

# Operating Systems

Youjip Won

**KAIST**



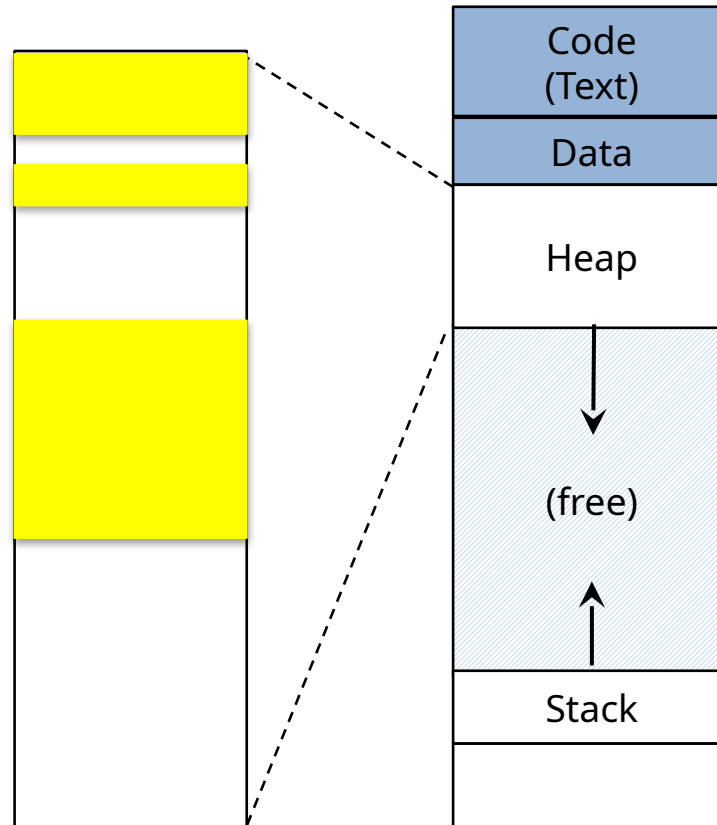
# 17. Free-Space Management

---

# Managing heap

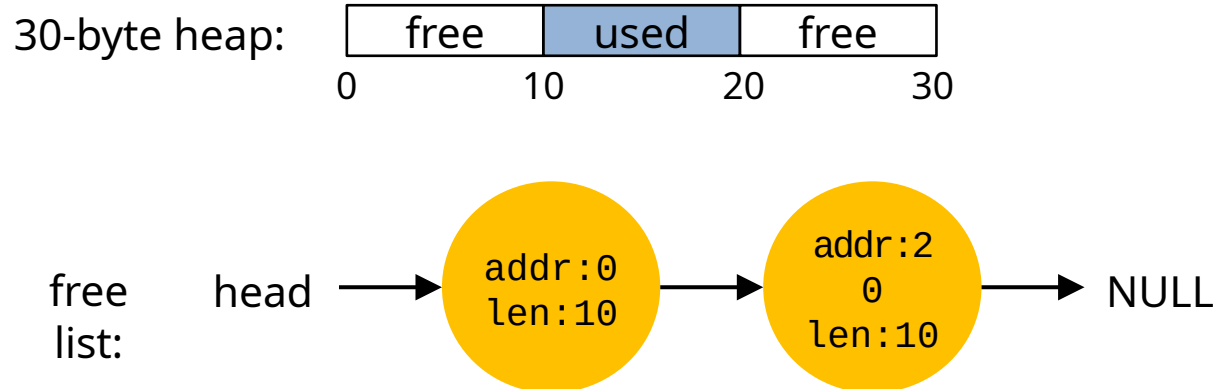
`malloc()`  
`free()`  
in **libc**

To manage the  
memory in a  
heap



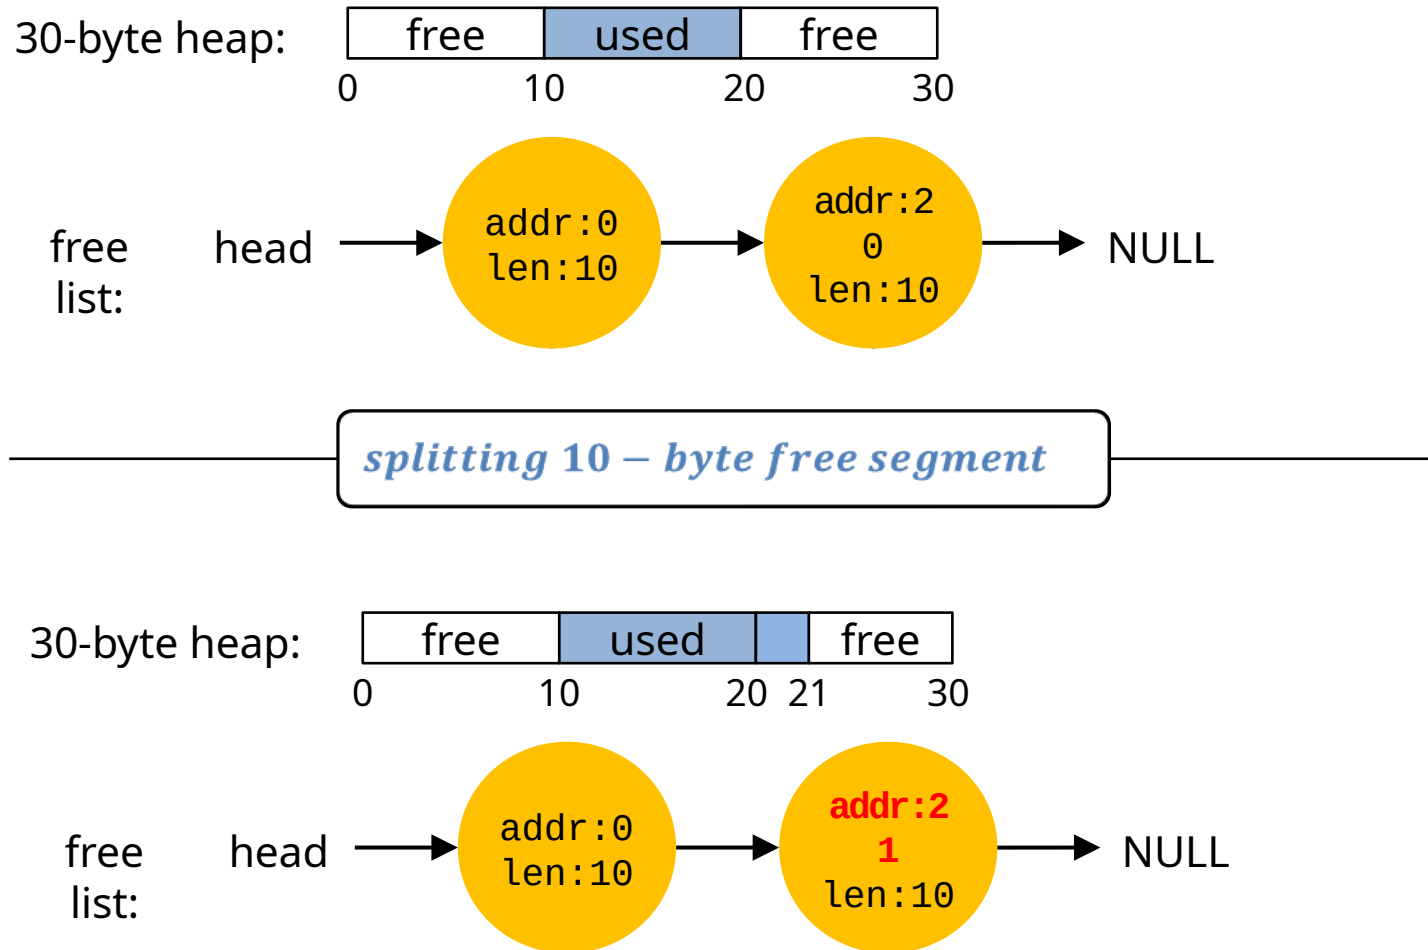
# Splitting

- Finding a free chunk of memory that can satisfy the request and splitting it into two.
  - ◆ When request for memory allocation is **smaller** than the size of free chunks.



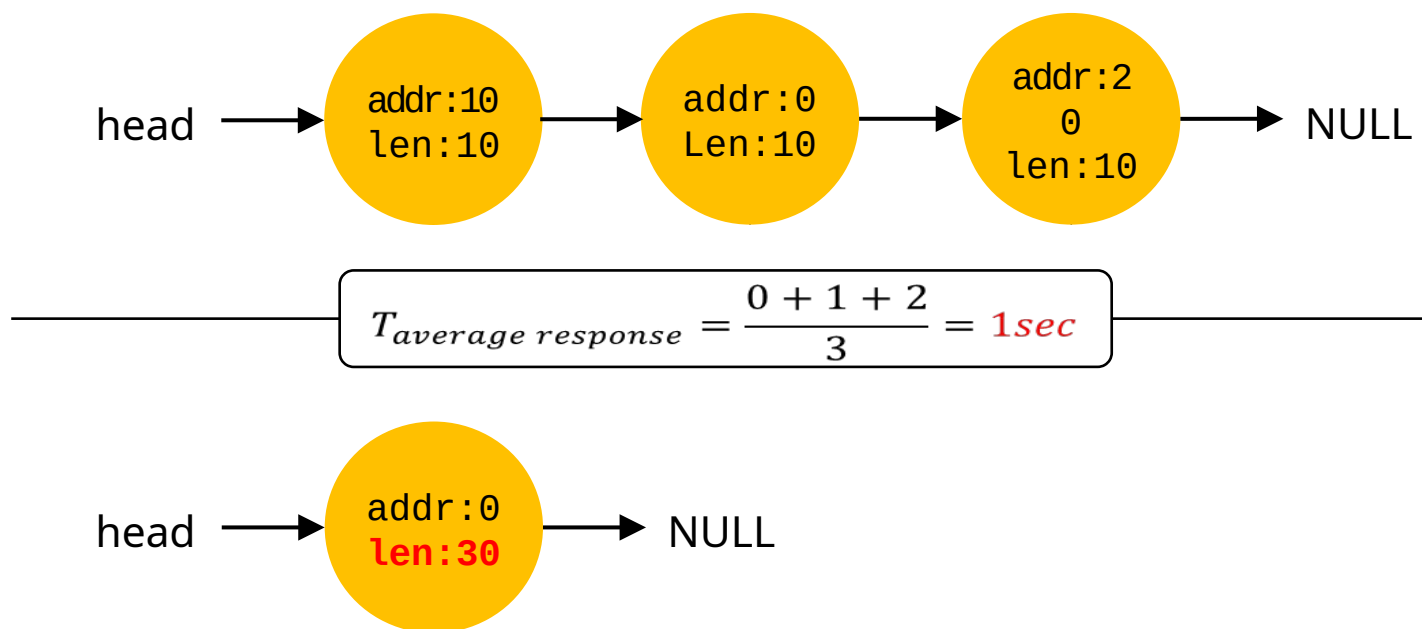
# Splitting(Cont.)

- Two 10-bytes free segment with **1-byte request**



# Coalescing

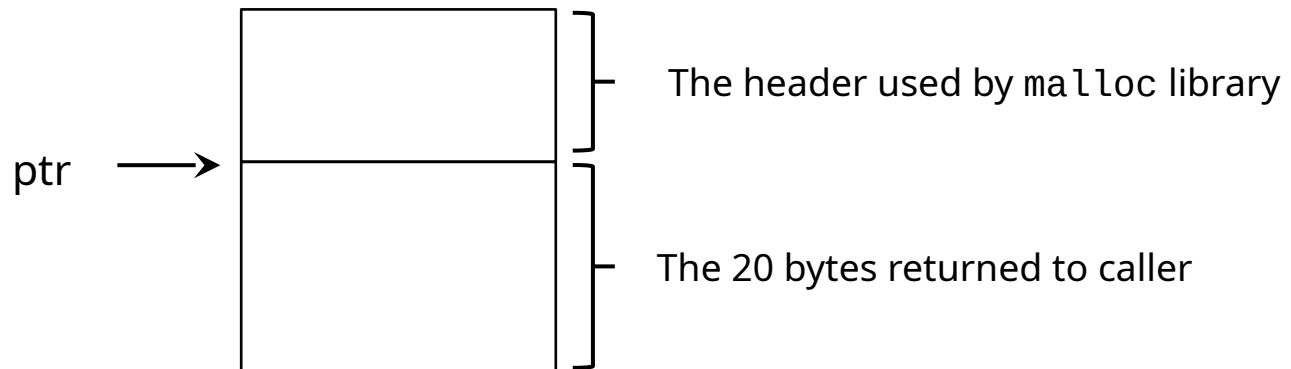
- If a user requests memory that is **bigger than free chunk size**, the list will **not find** such a free chunk.
- Coalescing: **Merge** returning a free chunk with existing chunks into a large single free chunk if **addresses** of them are **nearby**.



# Tracking The Size of Allocated Regions

- The interface to `free(void *ptr)` does **not take a size parameter**.
  - ◆ How does the library **know the size** of memory region that will be back into free list?

```
ptr =  
malloc(20);
```

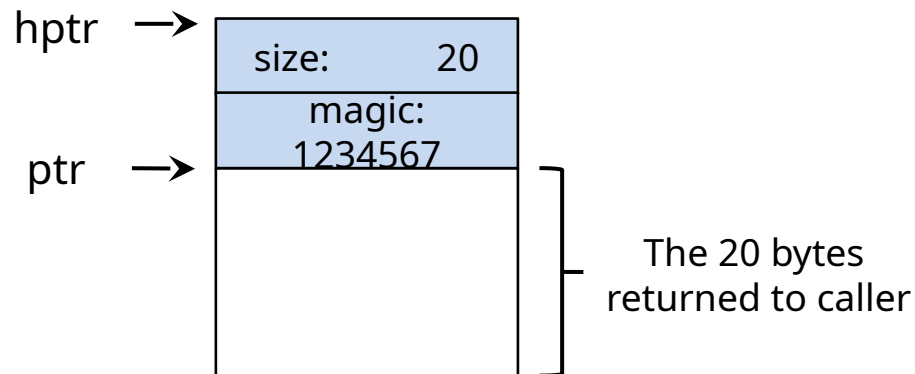


An Allocated Region Plus Header

# The Header of Allocated Memory Chunk

```
typedef struct __header_t {  
    int size;  
    int magic;  
} header_t;
```

Actual chunk size of malloc(N) = N + size of header  
Here, 28 Byte





## The Header of Allocated Memory Chunk(Cont.)

```
void free(void *ptr) {  
    header_t *hptr = (void *)ptr - sizeof(header_t);  
    ...  
    assert(hptr->magic==1234567) ;  
    ...  
}
```

# Embedding A Free List

```
typedef struct __node_t {  
    int size;  
    struct __node_t *next;  
} nodet_t;
```

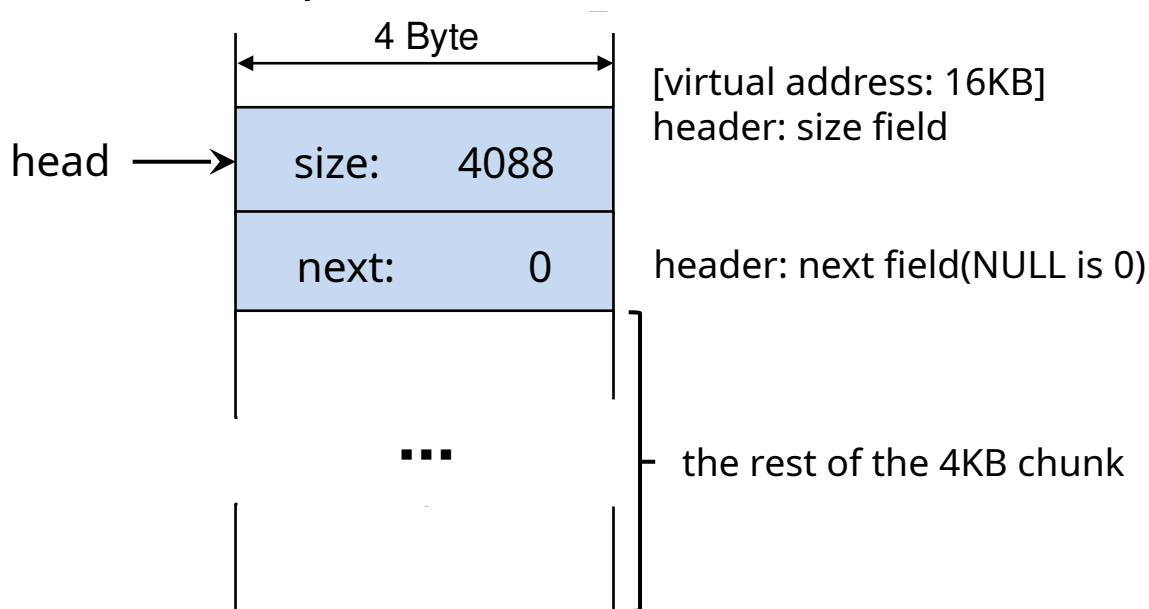
# Heap Initialization

// mmap() returns a pointer to a chunk of free space

```
node_t *head = mmap(NULL, 4096, PROT_READ|PROT_WRITE,  
                    MAP_ANON|MAP_PRIVATE, -1, 0);
```

```
head->size = 4096 - sizeof(node_t);
```

```
head->next = NULL;
```



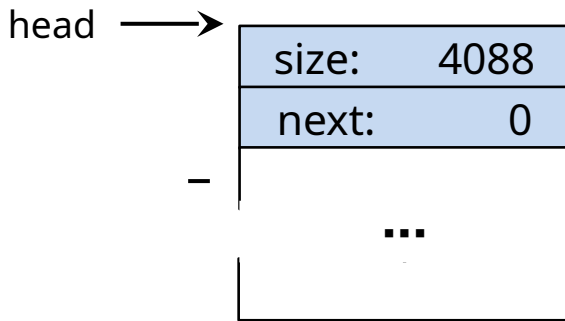
# Embedding A Free List: Allocation

- If a chunk of memory is requested, the library **will first find** a chunk that is **large enough** to accommodate the request.
- The library will
  - ◆ **Split** the large free chunk into two.
    - **One** for the **request** and the **remaining** free chunk
  - ◆ **Shrink** the size of free chunk in the list.

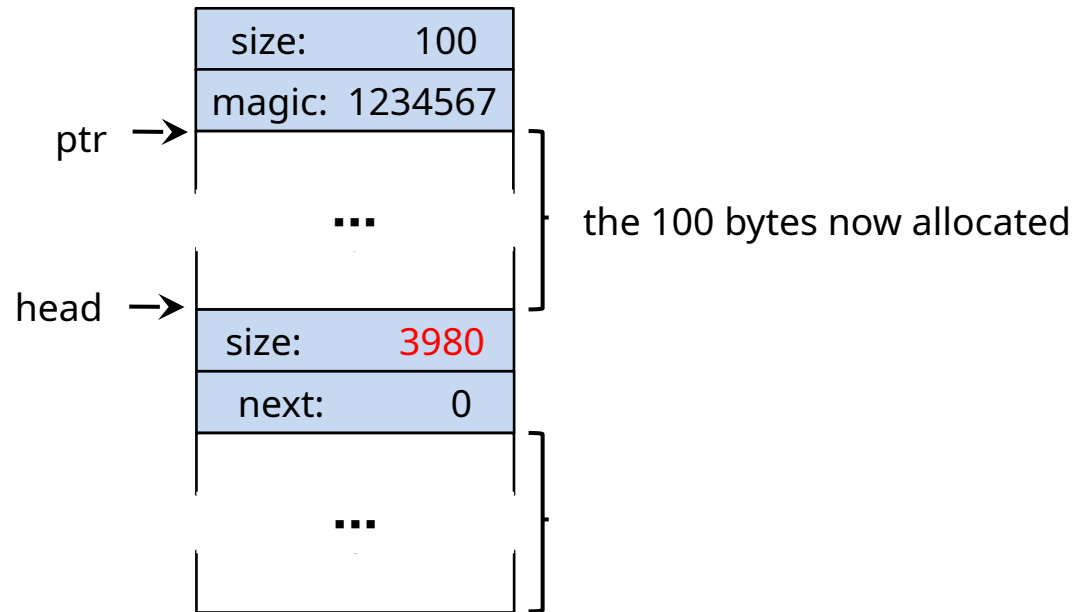
# Embedding A Free List: Allocation(Cont.)

- Example: a request for 100 bytes by `ptr = malloc(100)`  
→ 108 byte is returned.

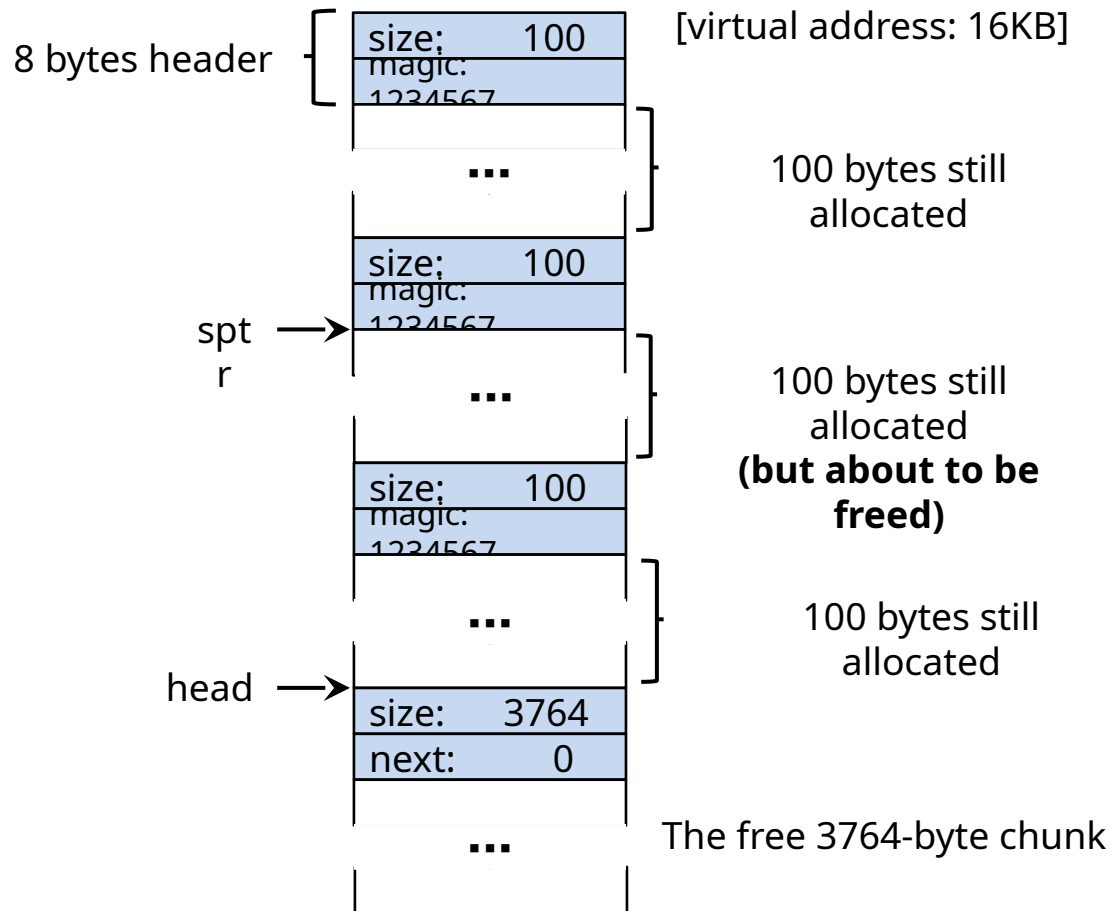
A 4KB Heap With One Free Chunk



A Heap : After One Allocation



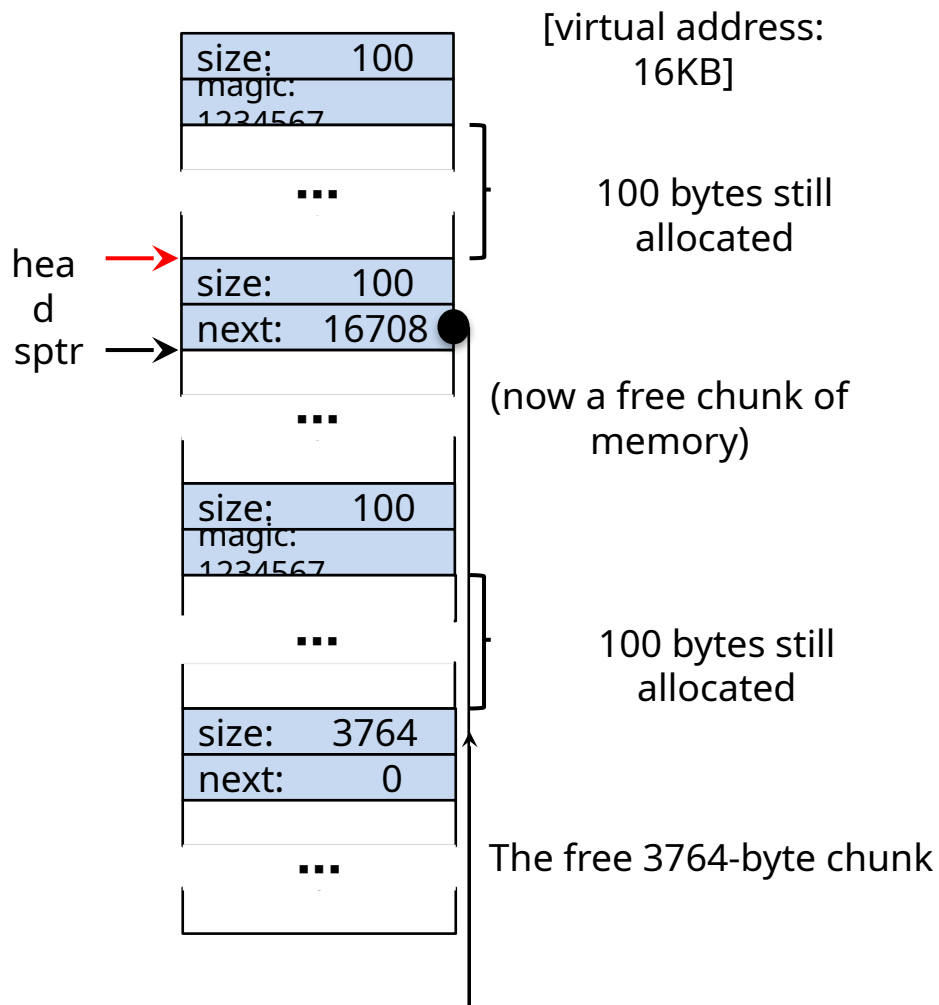
# Free Space With Chunks Allocated



Free Space With Three Chunks Allocated

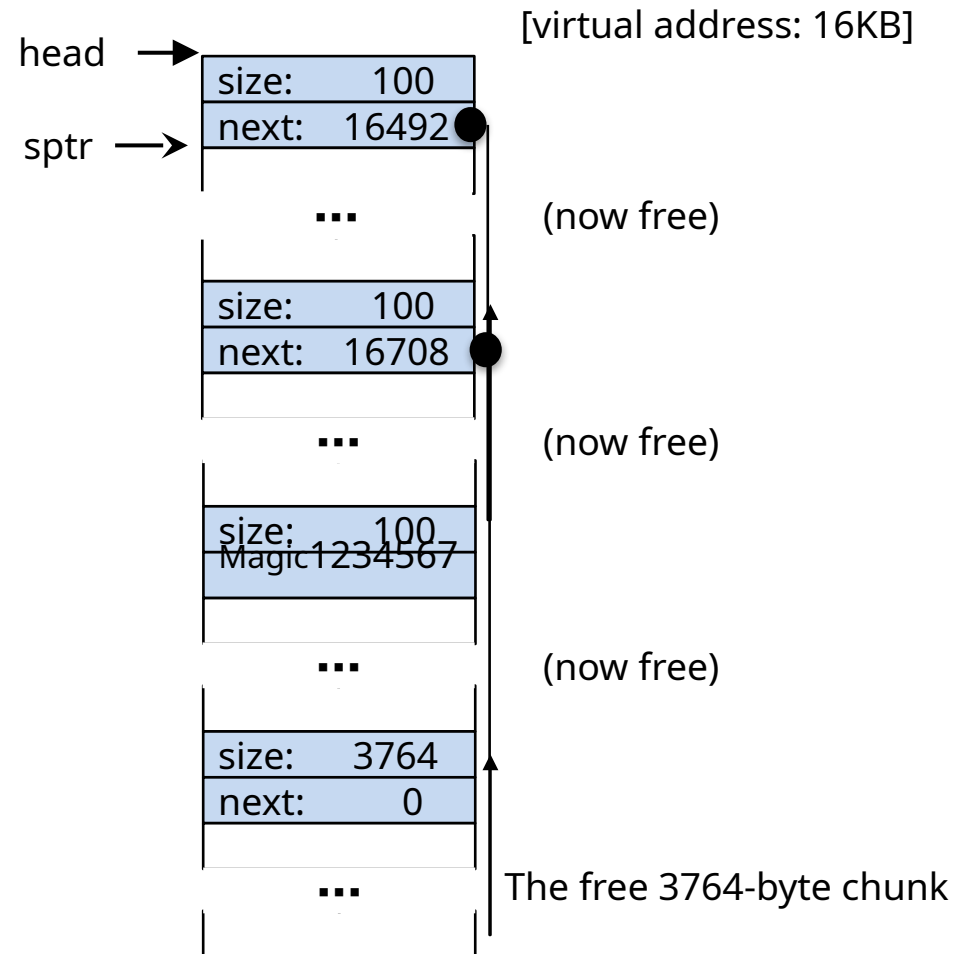
# Free Space With free()

```
□ free(sptr)
  void* tmp = head ;
  head = sptr ;
  head->next = tmp ;
```



# Free Space With Freed Chunks

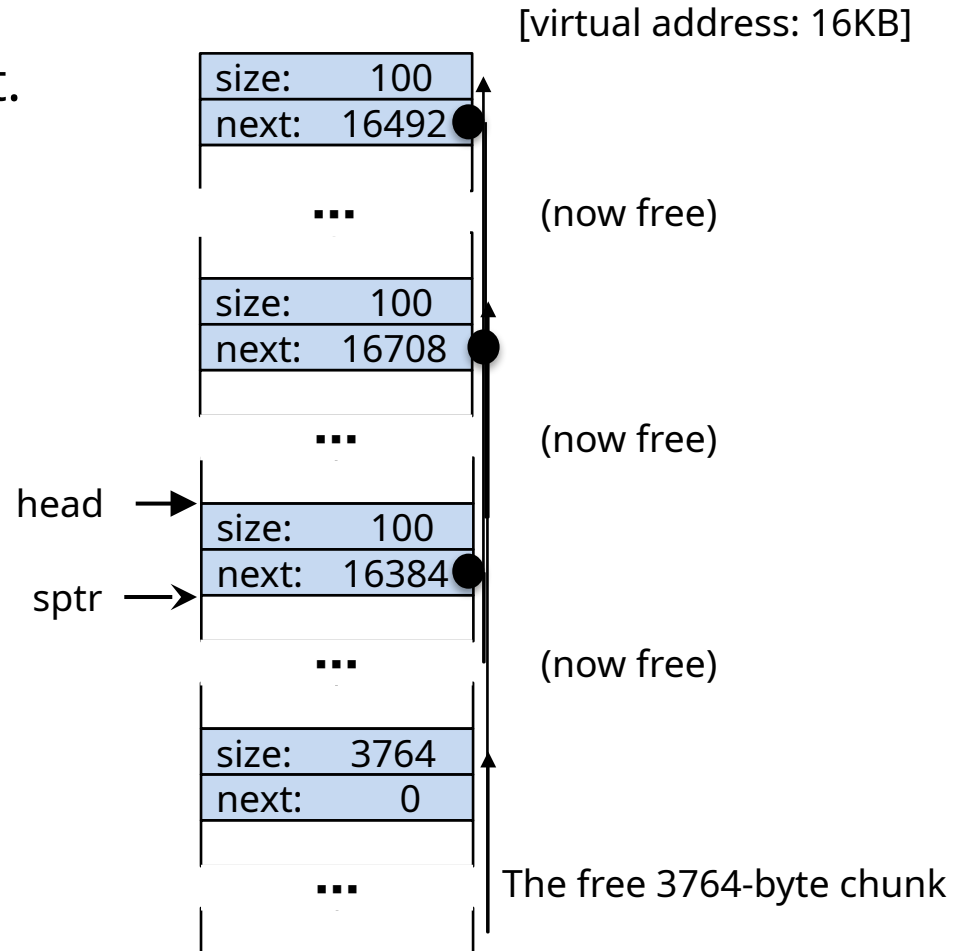
## ▣ free(sptr)





# Free Space With Freed **Chunks**

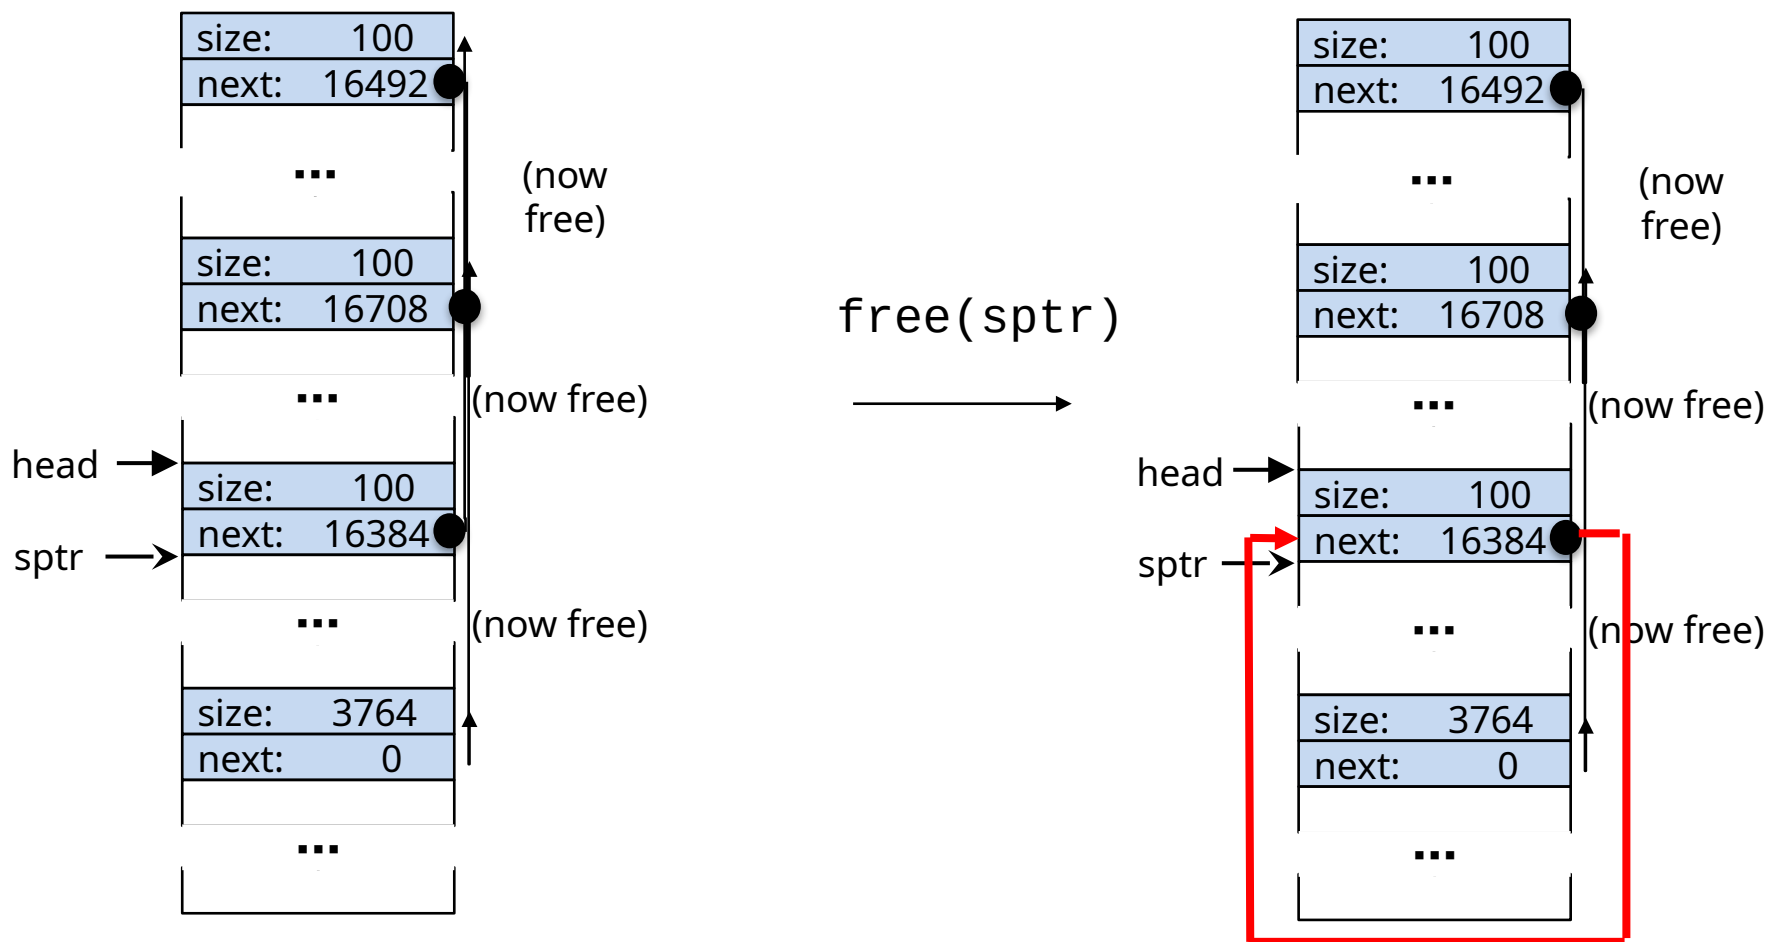
- ▣ `free(sptr)`
- ▣ **Coalescing** is needed in the list.



# Double free

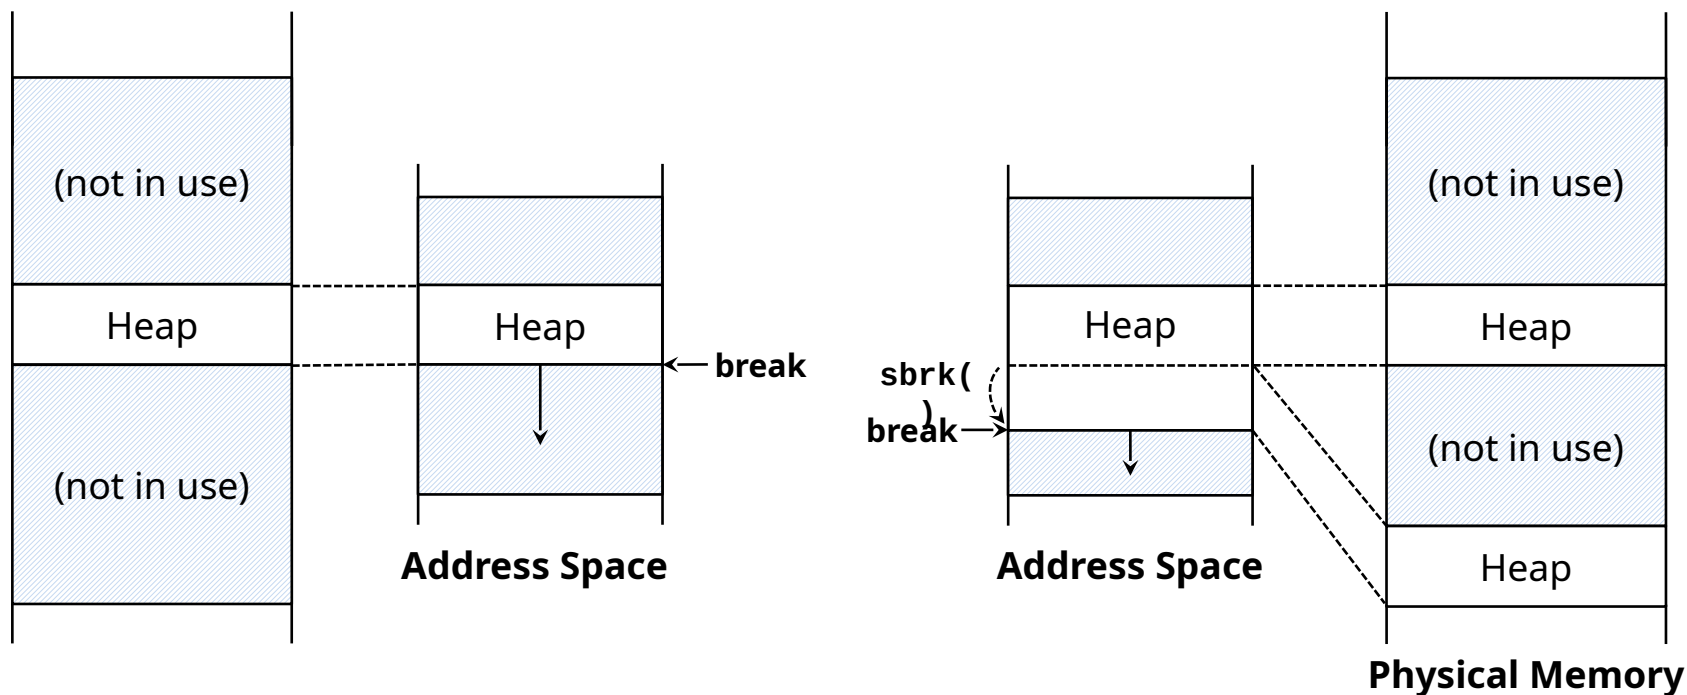
## Free(sptr): dangling reference

[virtual address: 16KB]



# Growing The Heap

- Most allocators **start** with a **small-sized heap** and then **request more** memory from the OS when they run out.
  - e.g., `sbrk()`, `brk()` in most UNIX systems.



# Managing Free Space: Basic Strategies

## □ Best Fit:

- ◆ Finding free chunks that are **big or bigger than the request**
- ◆ Returning the **one of smallest** in the chunks **in the group** of candidates

## □ Worst Fit:

- ◆ Finding the **largest free chunks** and allocation the amount of the request
- ◆ **Keeping the remaining chunk** on the free list.

# Managing Free Space: Basic Strategies(Cont.)

## □ First Fit:

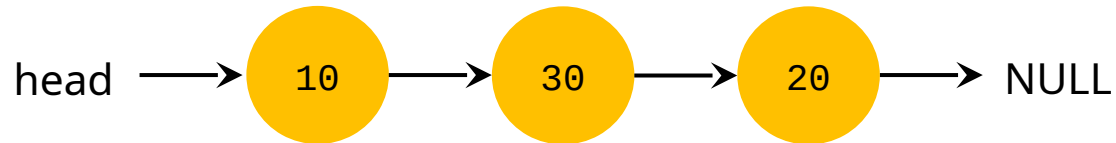
- ◆ Finding the **first chunk** that is **big enough** for the request
- ◆ Returning the requested amount and remaining the rest of the chunk.

## □ Next Fit:

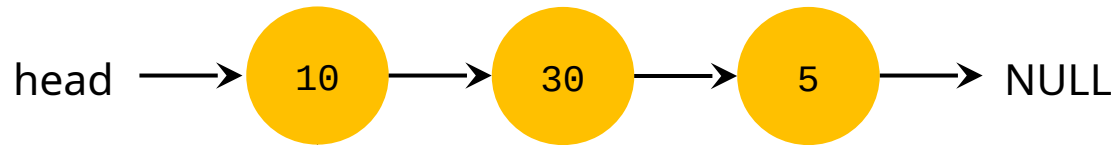
- ◆ Finding the first chunk that is big enough for the request.
- ◆ Searching at **where one was looking** at instead of the beginning of the list.

# Examples of Basic Strategies

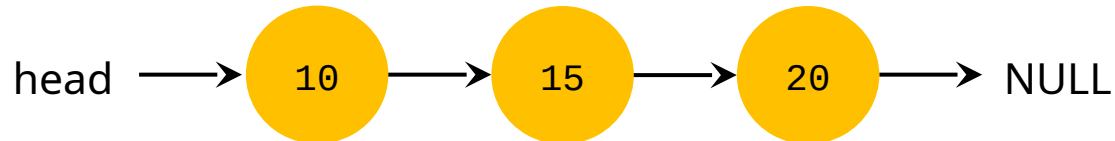
- Allocation Request Size 15



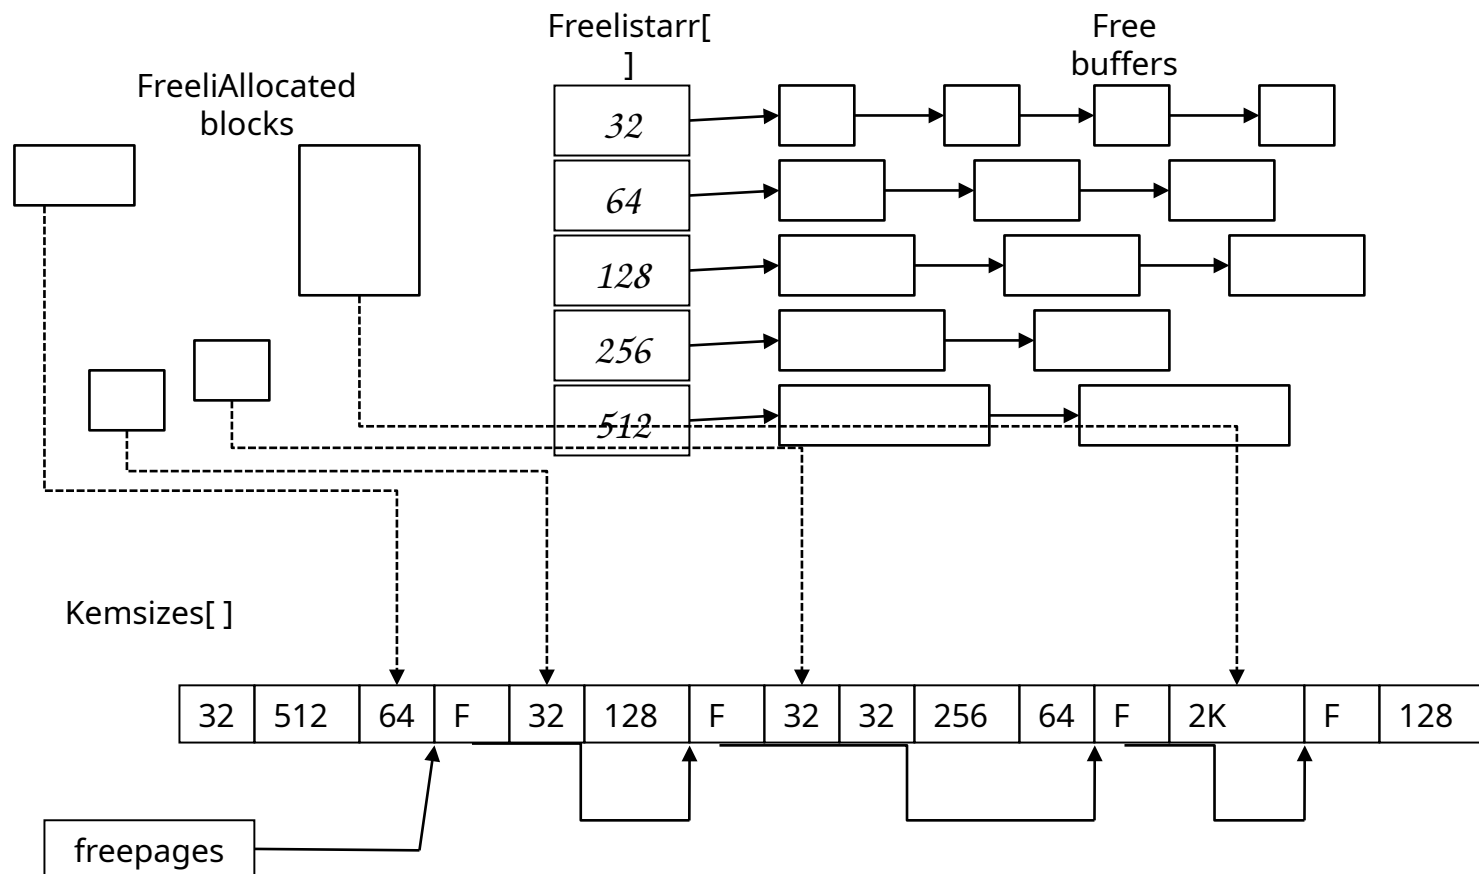
- Result of Best-fit



- Result of Worst-fit



# Segregated List: McKusick-Karels Allocator (4.3 BSD)



# McKusick-Karels Allocator

- Used in 4.3BSD and Digital UNIX.
- Management
  - ◆ Contiguous set of pages
  - ◆ Same buffer size within a page
- Manage pages by page usage array(kmemsizes[]).
- ◆ Element of kmemsizes: free | size of buffer
- `malloc()` allocates the requested sized buffer from the free list.
- `malloc()` invoke a routine that acquires free page and splits by a requested size if free list is empty.
- How much memory should dedicate to the pool of memory that serves specialized requests of a given size?



# Buddy System

- ❑ Create the small buffers by repeatedly halving a large buffer and coalesce the adjacent free buffers.
- ❑ When a buffer is split, each half is called the buddy of the other.
- ❑ Minimum size is 32byte.
- ❑ Use a bitmap to monitor each 32-byte chunk of the block.

# Detailed step

- Allocate 256 byte
  - ◆ Split the block into A and A'. Puts A' in the 512 byte free list.
  - ◆ Split A into B and B'. Put B' in the 256 byte free list.
  - ◆ Return B to the client
- Allocate 128 Byte
  - ◆ Split B' into C and C'. Put C' in the free list.
  - ◆ Return C to the client.
- Allocate 64 byte.
  - ◆ Split c' into D and D'. Put D' in the free list.
  - ◆ Return D to the client.
- Allocate 128
  - ◆ Split A' into E and E'. Put E' in the free list of 256 byte
  - ◆ Split E into F and F'. Put F' in the free list of 128 byte
  - ◆ Allocate F.

# Example

- Maintain freelist for every buffer size(32-512).

bitmap

1	1	1	1	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

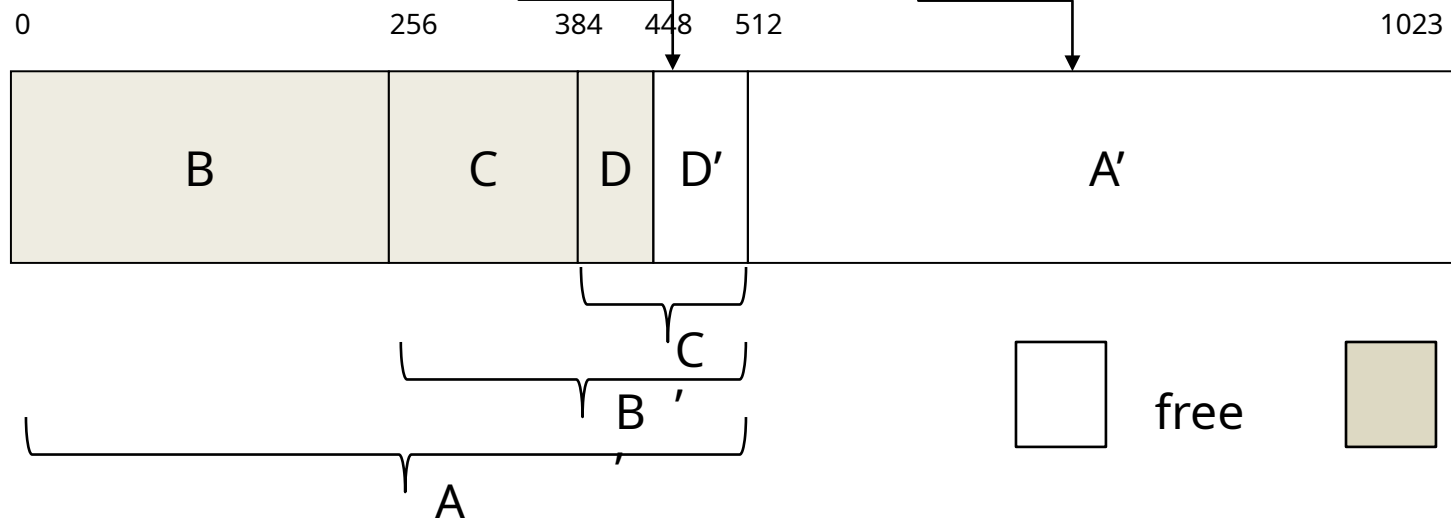
Free list headers

32	64	128	256	512
----	----	-----	-----	-----

*Allocate(256)*

*Allocate(128)*

*Allocate(64)*

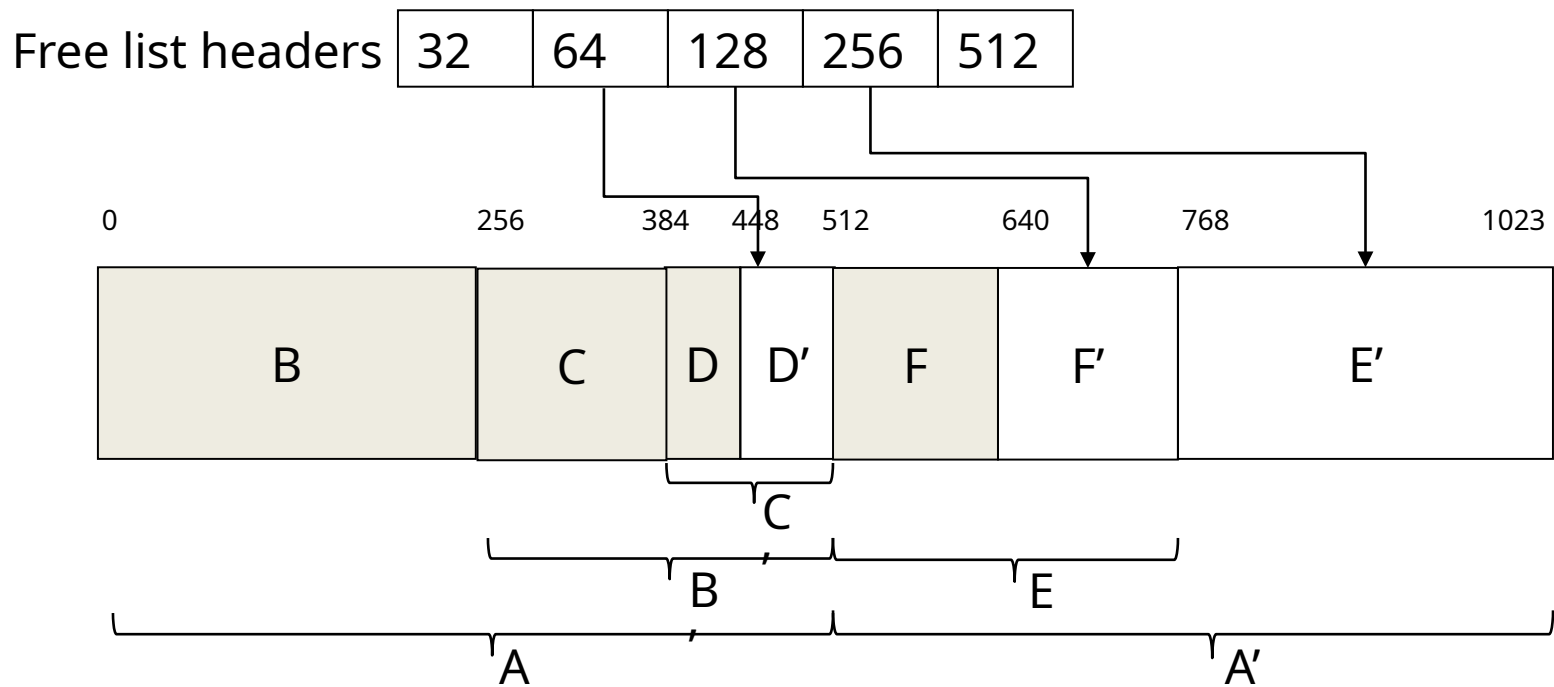


# Example

bitmap

1	1	1	1	1	1	1	1	1	0	0	0	0	1	1	0	0	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

*Allocate(128)*

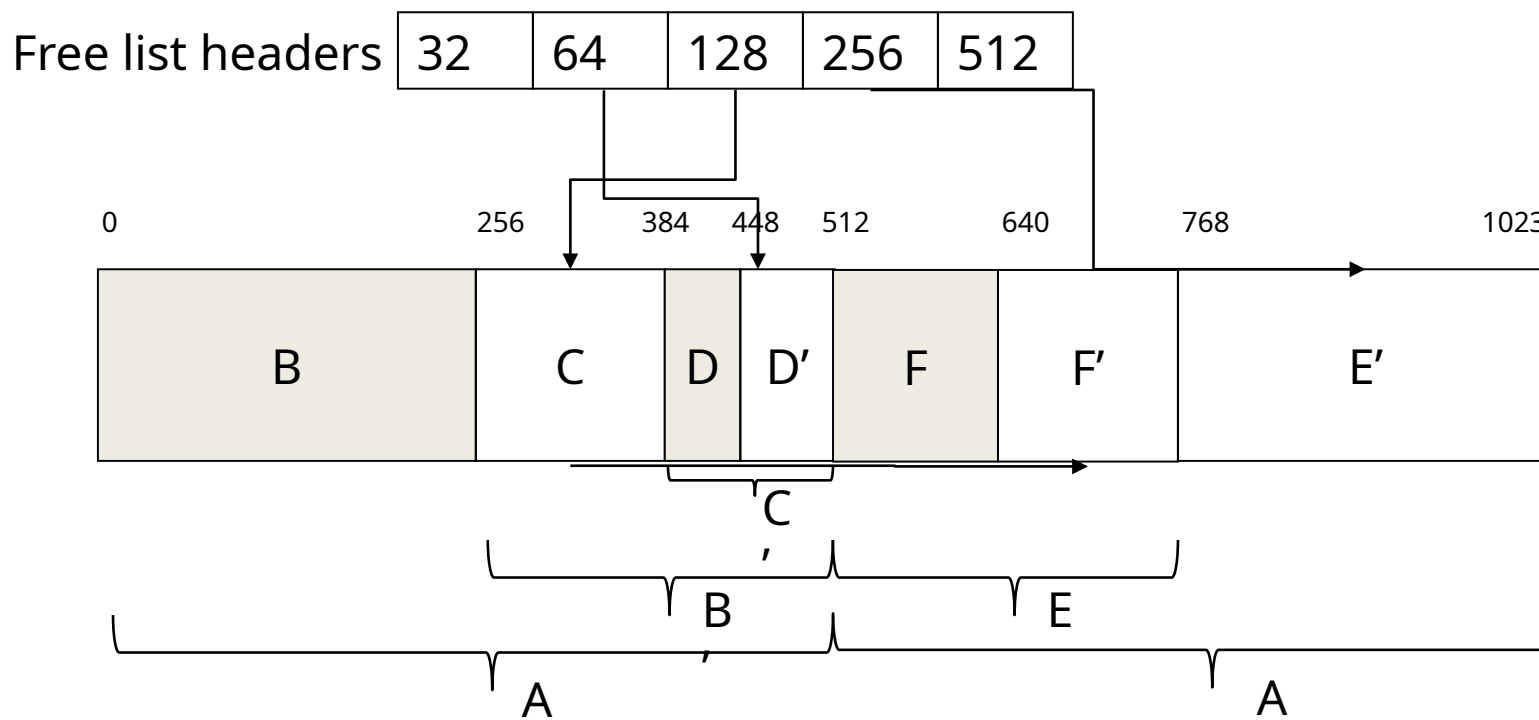


# Example

bitmap

1	1	1	1	1	1	1	1	1	0	0	0	0	1	1	0	0	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

*Release(C, 128)*

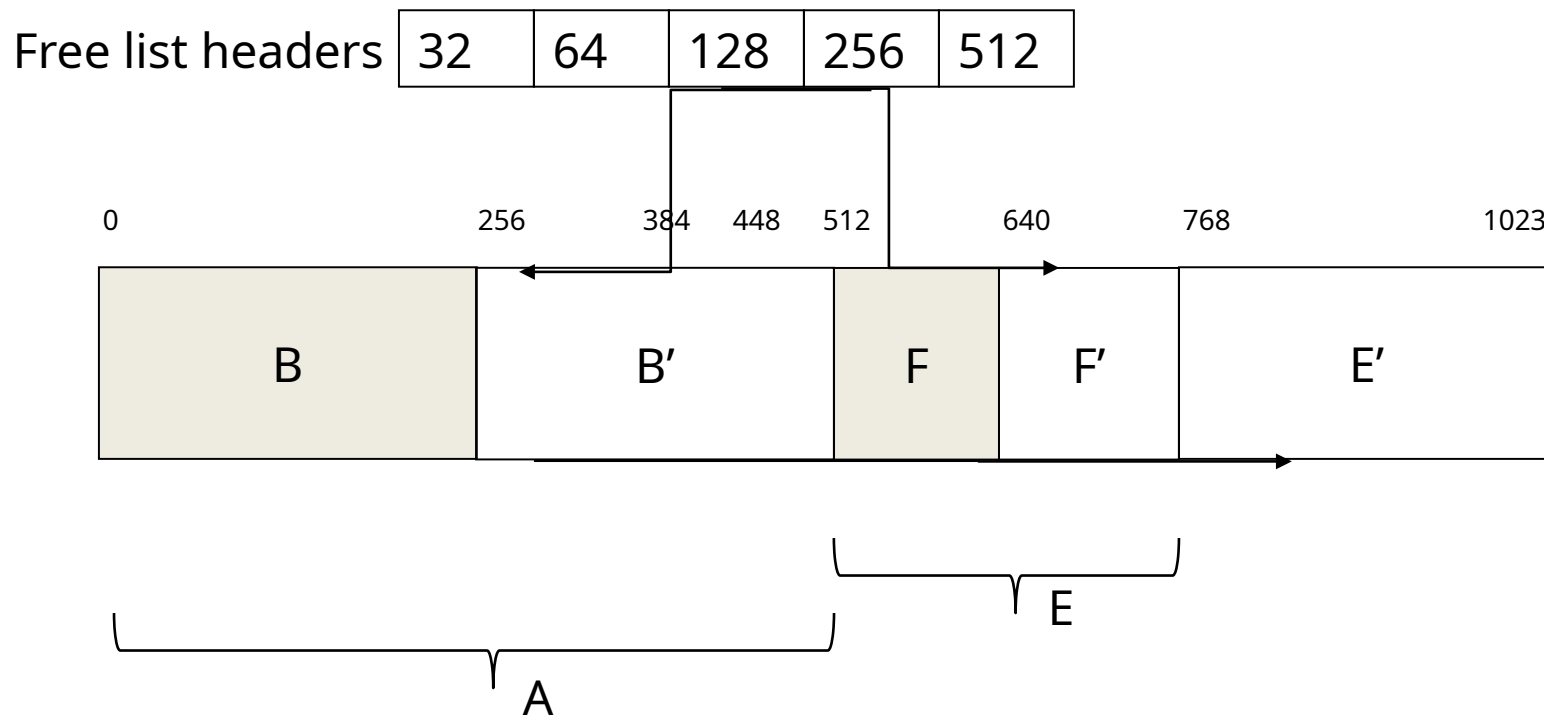


# Example

bitmap

1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

*Release(D, 64)*



# Analysis

- Characteristic
  - Internal fragmentation by power-of-two allocation
  - Easy to find a buddy of a buffer by address and size
  - Use bitmap for coalescing.
- Advantage
  - Does a good job of coalescing adjacent free buffers.
  - Easy exchange of memory between the allocator and the paging system
- Disadvantage
  - Performance degrade: every time a buffer is released, the allocator tries to coalesce as much as possible.
  - Release routine needs both the address and size of the buffer.
  - Partial release is insufficient.

# Slab allocator

- ❑ Advanced form of segregated list
- ❑ Slab: a set of kernel pages
  - ◆ Consists of same objects, e.g. inodes, locks, sockets
  - ◆ They are all initialized before allocation.
- ❑ Using object
  - ◆ Allocate memory and Construct the object
  - ◆ Use the object.
  - ◆ Deconstruct it, Free the memory.
- ❑ Construction
  - ◆ Initialization of various fields
- ❑ Object reuse
  - ◆ It is better to reuse possible data structures rather than release



# Memory allocation

- ▣ libc

- ◆ Based on linked list

- ▣ kernel

- ◆ Buddy: allocating memory to process
- ◆ Slab: allocating memory for small kernel objects (proc structure, inode, socket and etc.)

