

Operating Systems

Youjip Won

KAIST



18. Paging: Introduction

Concept of Paging


- Paging **splits up** address space into **fixed-zed** unit called a **page**.
 - ◆ Segmentation: variable size of logical segments(code, stack, heap, etc.)
- With paging, **physical memory** is also **split** into some number of pages called a **page frame**.
- **Page table** per process is needed **to translate** the virtual address to physical address.

Advantages Of Paging

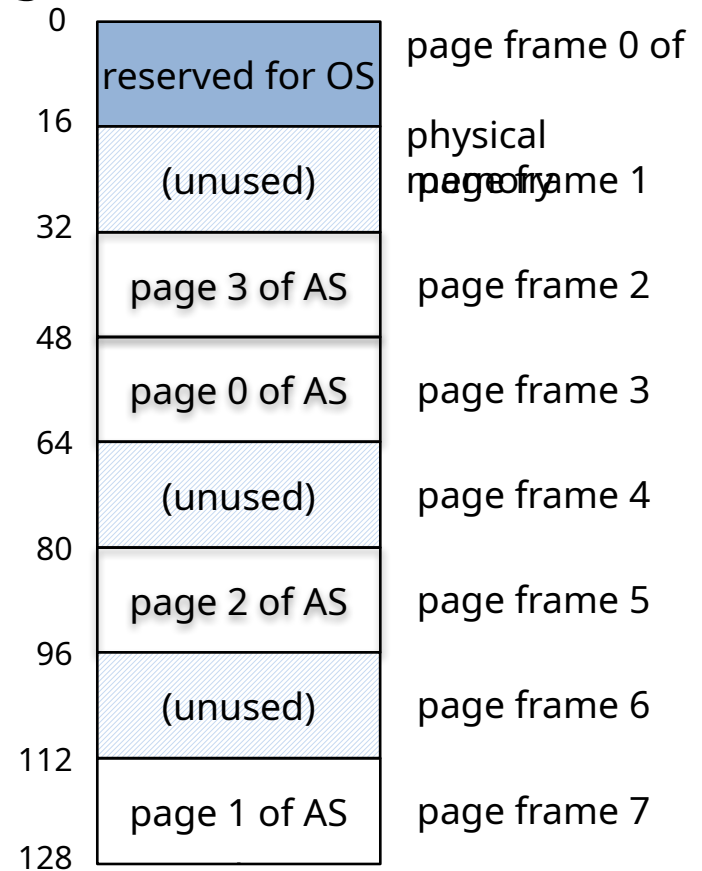
- ▣ **Flexibility:** Supporting the abstraction of address space effectively
 - ◆ Don't need assumption how heap and stack grow and are used.
- ▣ **Simplicity:** ease of free-space management
 - ◆ The page in address space and the page frame are the same size.
 - ◆ Easy to allocate and keep a free list

Example: A Simple Paging

- 128-byte physical memory with 16 bytes page frames
- 64-byte address space with 16 bytes pages




A Simple 64-byte Address Space




64-Byte Address Space Placed In Physical Memory

Address Translation

- Two components in the virtual address
 - VPN: virtual page number
 - Offset: offset within the page




- Example: virtual address 21 in 64-byte address space



Example: Address Translation

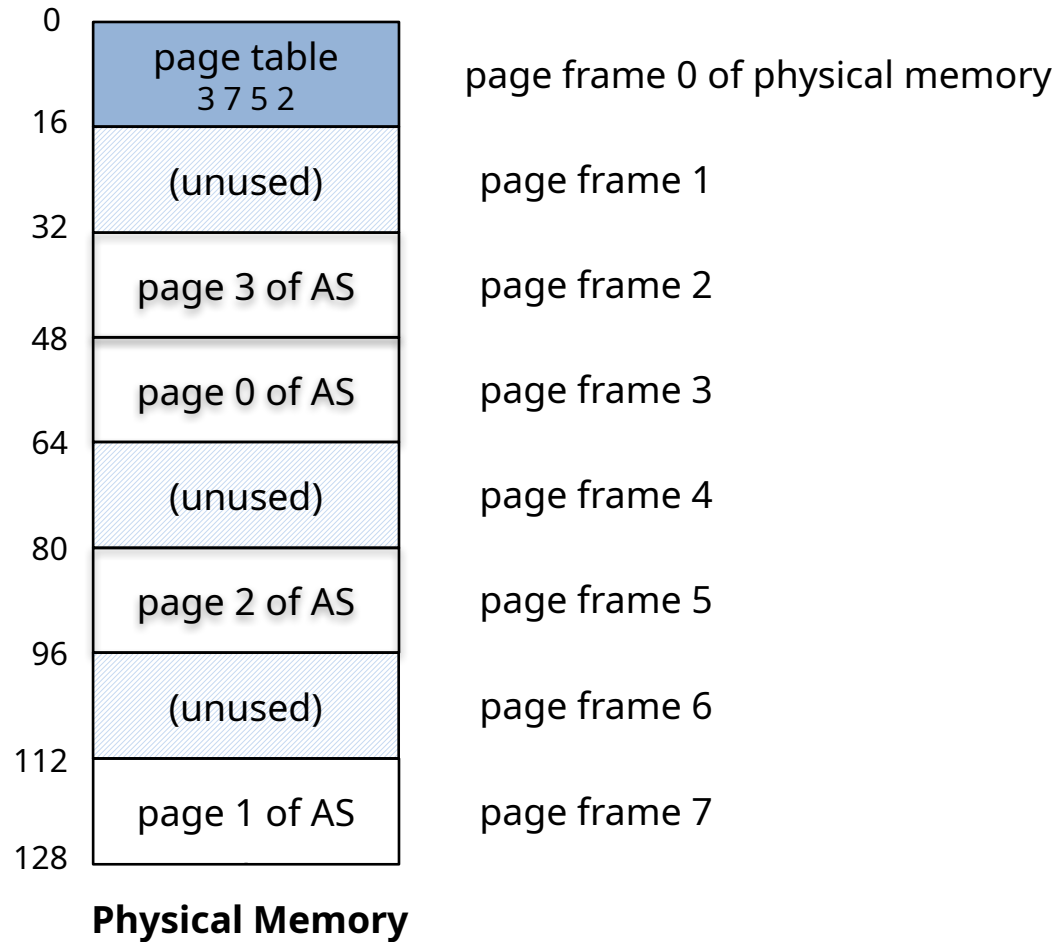
- ▣ The virtual address 21 in 64-byte address space



Where Are Page Tables Stored?

- ❑ Page tables can get awfully large.
 - ◆ 32-bit address space with 4-KB pages, 20 bits for VPN
 - Page offset for 4 Kbyte page: 20 bit
 - $4MB = 2^{20} \text{ entries} * 4 \text{ Bytes per page table entry}$
- ❑ Page tables for each process are stored in memory.

Example: Page Table in Kernel Physical Memory




What Is In The Page Table?

- ▣ The page table is a **data structure** that is used to map the virtual address to physical address.
 - ◆ Simplest form: a linear page table, an array
- ▣ The OS **indexes** the array by VPN, and looks up the page-table entry.

Common Flags Of Page Table Entry

- ▣ **Valid Bit:** Indicating whether the particular translation is valid.
- ▣ **Protection Bit:** Indicating whether the page could be read from, written to, or executed from
- ▣ **Present Bit:** Indicating whether this page is in physical memory or on disk(swapped out)
- ▣ **Dirty Bit:** Indicating whether the page has been modified since it was brought into memory
- ▣ **Reference Bit(Accessed Bit):** Indicating that a page has been accessed

Example: x86 Page Table Entry



An x86 Page Table Entry(PTE)

- P: present
- R/W: read/write bit
- U/S: supervisor
- A: accessed bit
- D: dirty bit
- PFN: the page frame number

Paging: Too Slow

- ▣ To find a location of the desired PTE, the **starting location** of the page table is **needed**.
- ▣ For every memory reference, paging requires the OS to perform one **extra memory reference**.

Accessing Memory With Paging

```
1 // Extract the VPN from the virtual address
2 VPN = (VirtualAddress & VPN_MASK) >> SHIFT
3
4 // Form the address of the page-table entry (PTE)
5 PTEAddr = PTBR + (VPN * sizeof(PTE))
6
7 // Fetch the PTE
8 PTE = AccessMemory(PTEAddr)
9
```

Accessing Memory With Paging

```
10 // Check if process can access the page
11 if (PTE.Valid == False)
12     RaiseException(SEGMENTATION_FAULT)
13 else if (CanAccess(PTE.ProtectBits) == False)
14     RaiseException(PROTECTION_FAULT)
15 else
16     // Access is OK: form physical address and fetch
17     it
18     offset = VirtualAddress & OFFSET_MASK
19     PhysAddr = (PTE.PFN << PFN_SHIFT) | offset
20     Register = AccessMemory(PhysAddr)
```

A Memory Trace

```
int array[1000];  
...  
for (i = 0; i < 1000; i++)  
    array[i] = 0;
```

```
prompt> gcc -o array array.c -Wall -o  
prompt> ./array
```

Memory access

```
0x1024 movl $0x0, (%edi,%eax,4) //[edi+eax*4]= 0
```

```
0x1028 incl %eax
```

```
0x102c cmpl $0x03e8,%eax //0000 0011 1110 10002 = 100010
```

```
0x1030 jne 0x1024
```


A Virtual(And Physical) Memory Trace

