

Operating Systems

Youjip Won

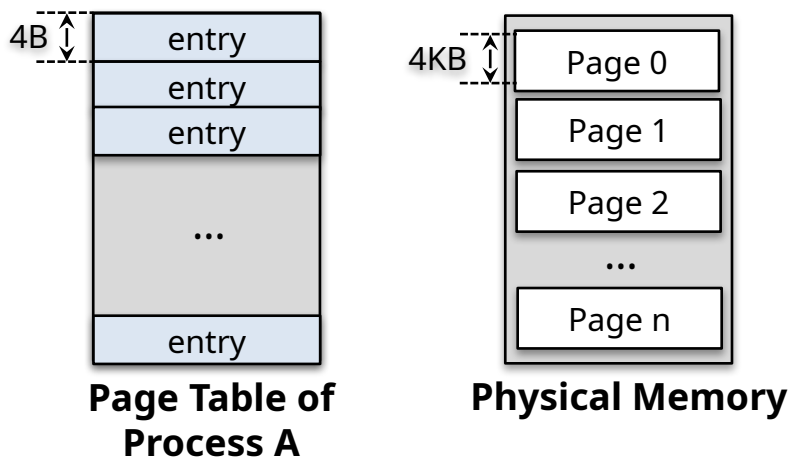
KAIST



20. Advanced Page Tables

Paging: Linear Tables

- We usually have one page table for every process in the system.
 - ◆ Assume that 32-bit address space with 4KB pages and 4-byte page-table entry.

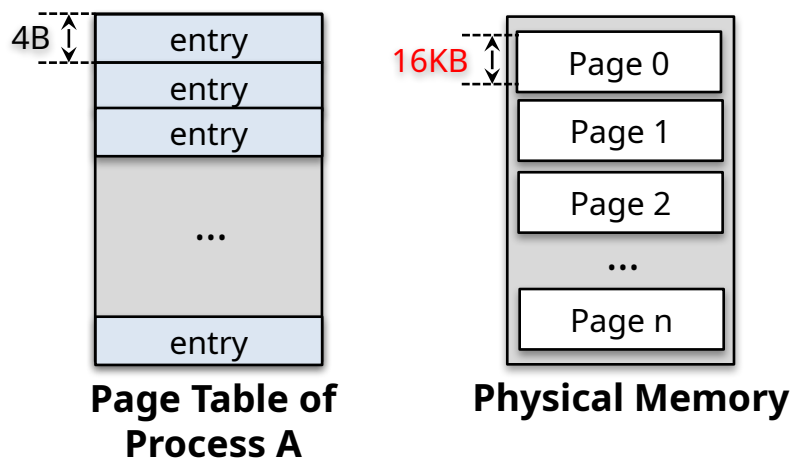


$$\text{Average turnaround time} = \frac{10 + 20 + 30}{3} = 20 \text{ sec}$$

Page table are **too big and thus consume too much memory.**

Paging: Smaller Tables

- Page tables are too big and thus consume too much memory.
 - ◆ Assume that 32-bit address space with **16KB** pages and 4-byte page-table entry.

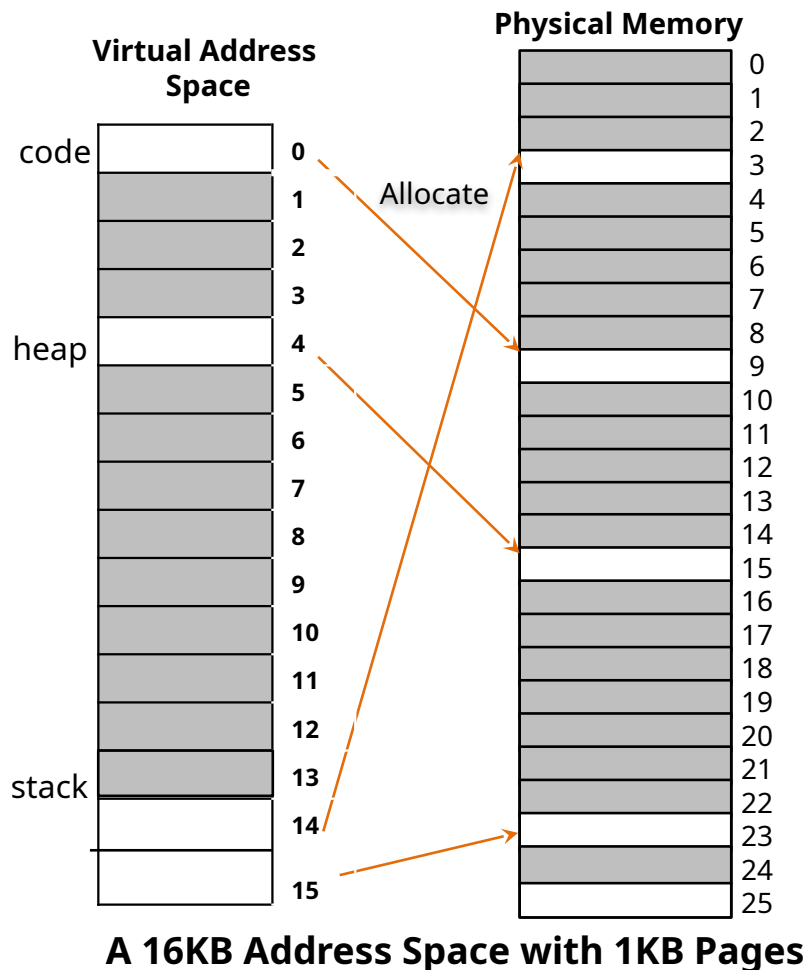


$$\text{Average turnaround time} = \frac{100 + 110 + 120}{3} = \mathbf{110 \text{ sec}}$$

Big pages lead to **internal fragmentation.**

Problem

- Single page table for the address space of process.

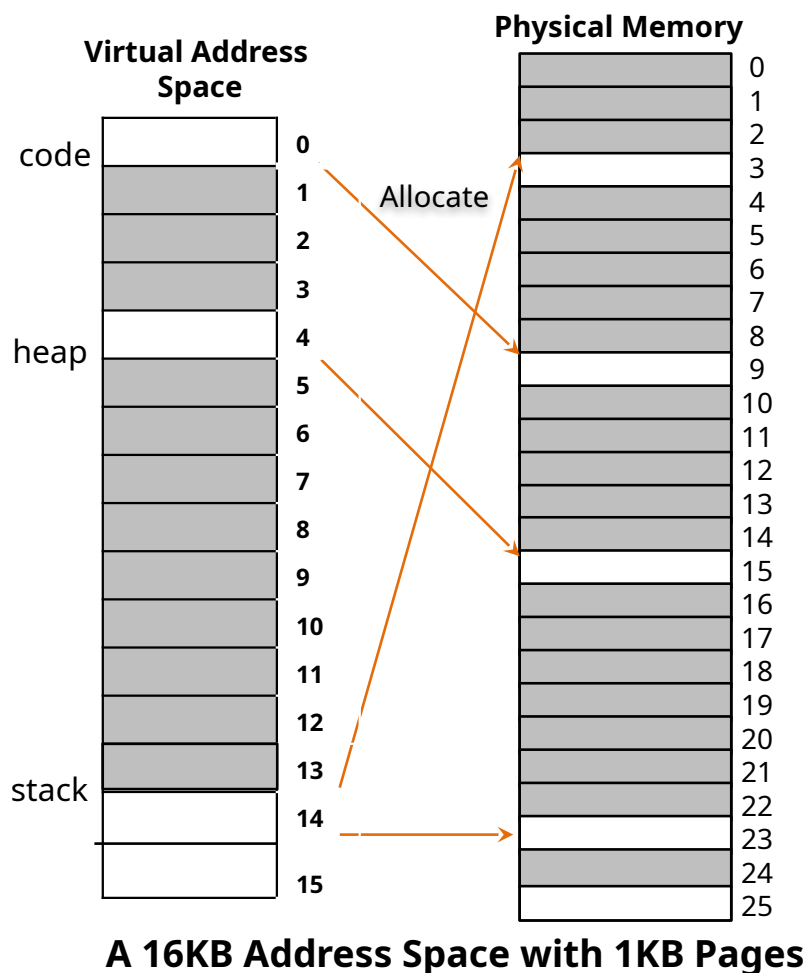


PFN	valid	prot	present	dirty
10	1	r-x	1	0
-	0	-	-	-
-	0	-	-	-
-	0	-	-	-
15	1	rw-	1	1
...
-	0	-	-	-
3	1	rw-	1	1
23	1	rw-	1	1

A Page Table For 16KB Address Space

Problem

- Most of the page table is **unused**, full of invalid entries.

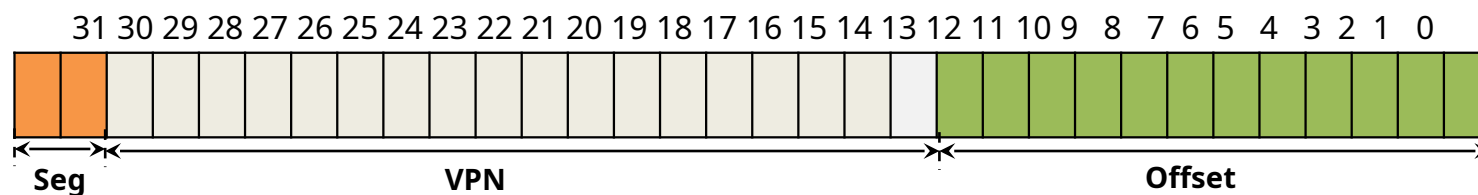


PFN	valid	prot	present	dirty
9	1	r-x	1	0
-	0	-	-	-
-	0	-	-	-
-	0	-	-	-
15	1	rw-	1	1
...
-	0	-	-	-
3	1	rw-	1	1
23	1	rw-	1	1

A Page Table For 16KB Address Space

Hybrid Approach: Paging and Segments

- Page table for each segment
 - ◆ the base register for each of these segments contains the physical address of a linear page table for that segment.
 - ◆ The bound register: indicate the end of the page table.
- Example: Each process has **three** page tables associated with it.



32-bit Virtual address space with 4KB pages

Seg value	Content
00	unused segment
01	code
10	heap
11	stack

TLB miss on Hybrid Approach

- ▣ The hardware gets **physical address** from **page table**.
 - ◆ The hardware uses the segment bits(SN) to determine which base and bounds pair to use.
 - ◆ The hardware then takes the **physical address** therein and **combines** it with the VPN as follows to form the address of the page table entry(PTE) .

```
01:    SN = (VirtualAddress & SEG_MASK) >> SN_SHIFT
02:    VPN = (VirtualAddress & VPN_MASK) >> VPN_SHIFT
03:    AddressOfPTE = Base[SN] + (VPN * sizeof(PTE))
```


Problem of Hybrid Approach

- Hybrid Approach is not without problems.
 - ◆ If we have a large but sparsely-used heap, we can still end up with a lot of page table waste.
 - ◆ Causing external fragmentation to arise again.

Multi-level Page Tables

- Turn the linear page table into something like a tree.
 - ◆ Chop up the page table into page-sized units.
 - ◆ If an entire page of page-table entries is invalid, don't allocate that page of the page table at all.
 - ◆ To track whether a page of the page table is valid, use a new structure, called **page directory**.

Multi-level Page Tables: Page directory

Linear Page Table

PBTR 201

valid	prot		
1	rx	12	PFN201
1	rx	13	
0	-	-	
1	rw	100	
<hr/>			
0	-	-	PFN202
0	-	-	
0	-	-	
0	-	-	
<hr/>			
0	-	-	PFN203
0	-	-	
0	-	-	
0	-	-	
<hr/>			
0	-	-	PFN204
0	-	-	
1	rw	86	
1	rw	15	

Multi-level Page Table

PBTR 200

	valid	PFN	
PFN200	1	201	PFN201
	0	-	
	0	-	
	1	203	

The Page Directory

[Page 1 of PT: Not Allocated]

[Page 2 of PT: Not Allocated]

	valid	prot	PFN	
PFN204	0	-	-	PFN204
	0	-	-	
	1	rw	86	
	1	rw	15	

Multi-level Page Tables

▣ Page Directory

- ◆ The page directory contains one entry per page of the page table.
- ◆ It consists of a number of **page directory entries(PDE)**.
- ◆ PDE has a valid bit and page frame number(PFN).

▣ Advantage

- ◆ Only allocates page-table space in proportion to the amount of address space you are using.
- ◆ The OS can grab the next free page when it needs to allocate or grow a page table.

▣ Disadvantage

- ◆ Multi-level table is a small example of a **time-space trade-off**.
- ◆ **Complexity**.

Example

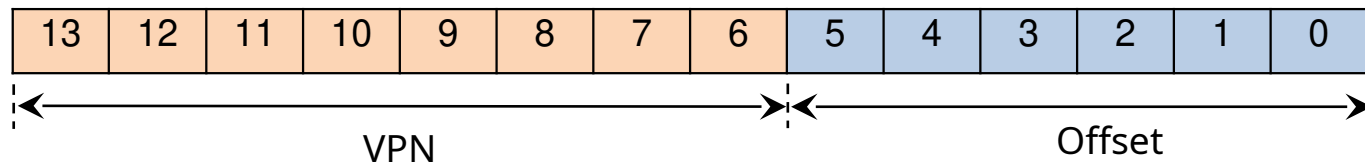
0000 0000	code
0000 0001	code
...	
	(free)
	(free)
0000 0100	
0000 0101	heap
.....	
	heap
	(free)
	(free)
1111 1110	
1111 1111	stack
	stack

- ▣ Page 0,1: code
- ▣ Page 4,5: heap
- ▣ Page 254, 255: stack

Example


Flag	Detail
Address space	16 KB (2^{14} Byte)
Page size	64 byte
Virtual address	14 bit
VPN	8 bit
Offset	6 bit
Page table entry	4 Byte

A 16-KB Address Space With 64-byte Pages



- Single level paging
 - ♦ 256 page table entries: 2^8 entries
 - ♦ Page table size: $256 * 4 \text{ Byte} = 1 \text{ Kbyte}$
 - ♦ Sixteen pages (64 byte each): $1024/64 = 16$

Example: single level page table

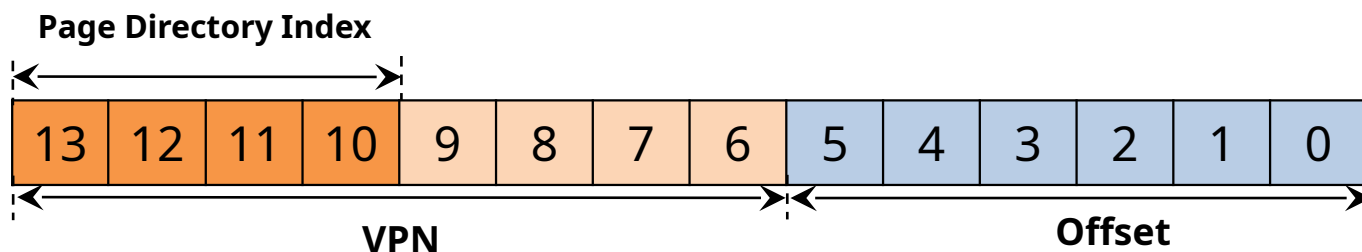


Example: two level paging

□ Page directory index

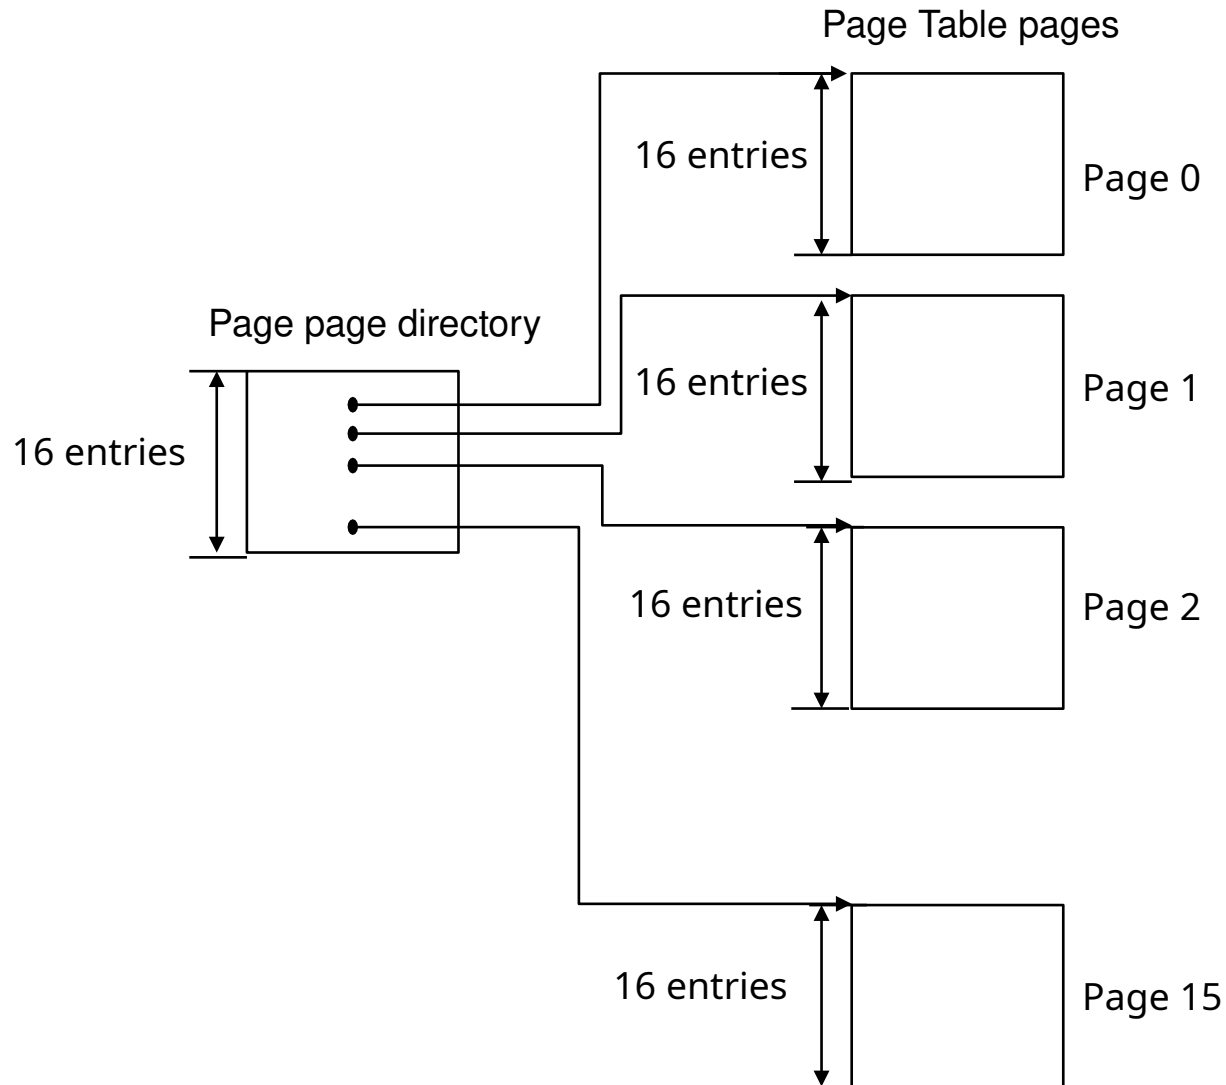
- ◆ A page table consists of 16 pages.
- ◆ 16 entries for page directory: one entry per page of the page table.
- ◆ $16 * 4 \text{ byte} = 64 \text{ byte}$ is required for page directory. → it can fit in a page.
- ◆ 4 bits for page directory index.

$$\text{PDEAddr} = \text{PageDirBase} + (\text{PDIndex} * \text{sizeof(PDE)})$$



- If the page-directory entry is **invalid**, raise an exception (The access is invalid).

Example: two level paging




Example: Page Table index

□ Page table index

- ◆ It is used to find the address of the page table entry.

$$\text{PTEAddr} = (\text{PDE.PFN} \ll \text{SHIFT}) + (\text{PTIndex} * \text{sizeof}(\text{PTE}))$$




Examples

Page Directory		Page of PT (@PFN:100)			Page of PT (@PFN:101)		
PFN	valid?	PFN	valid	prot	PFN	valid	prot
100	1	10	1	r-x	—	0	—
—	0	23	1	r-x	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	80	1	rw-	—	0	—
—	0	59	1	rw-	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	55	1	rw-
101	1	—	0	—	45	1	rw-


Single level paging: 16 pages

→ Two level paging: 3 pages



More than Two Levels

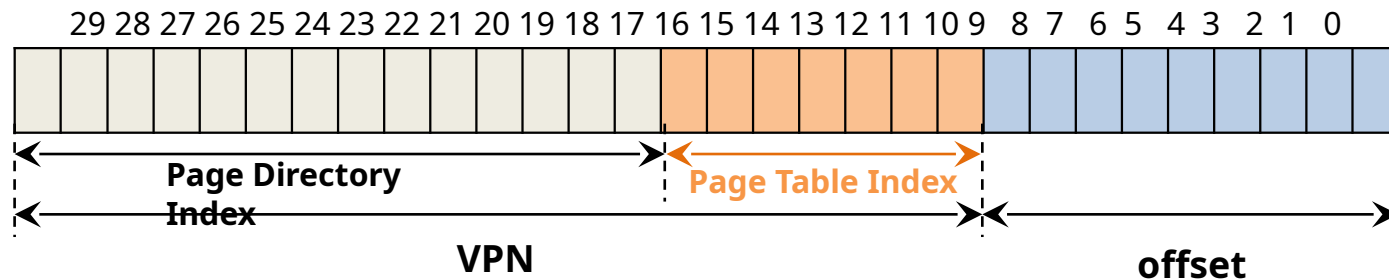
- In some cases, a deeper tree is required.
- Consider the following address space.



Flag	Detail
Virtual address	30 bit
Page size	512 byte
VPN	21 bit
Offset	9 bit

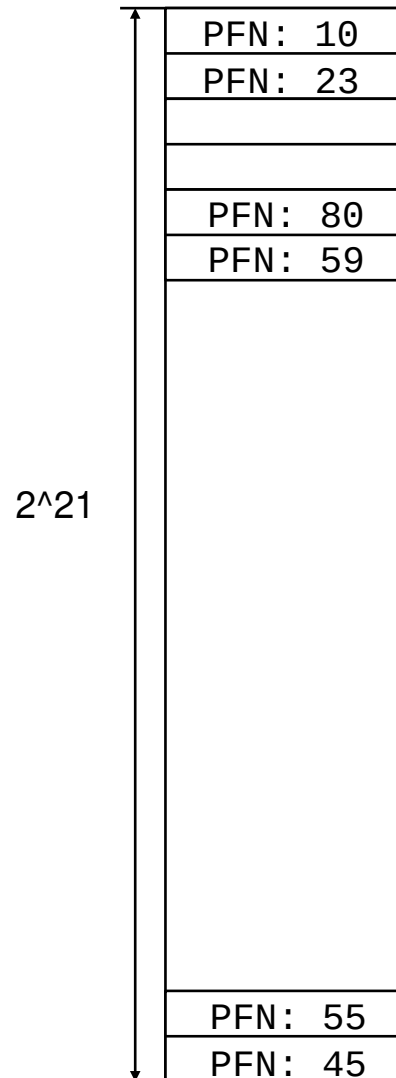
More than Two Level : Page Table Index

- Number of pages in a page table
 - ♦ 2^{21} page table entries
 - ♦ 2^7 (128) PTE's in a single page (512/4, 512 page size, 4 byte pte size)
 - ♦ 2^{14} pages in a page table
- Number of page directory entries: 2^{14}

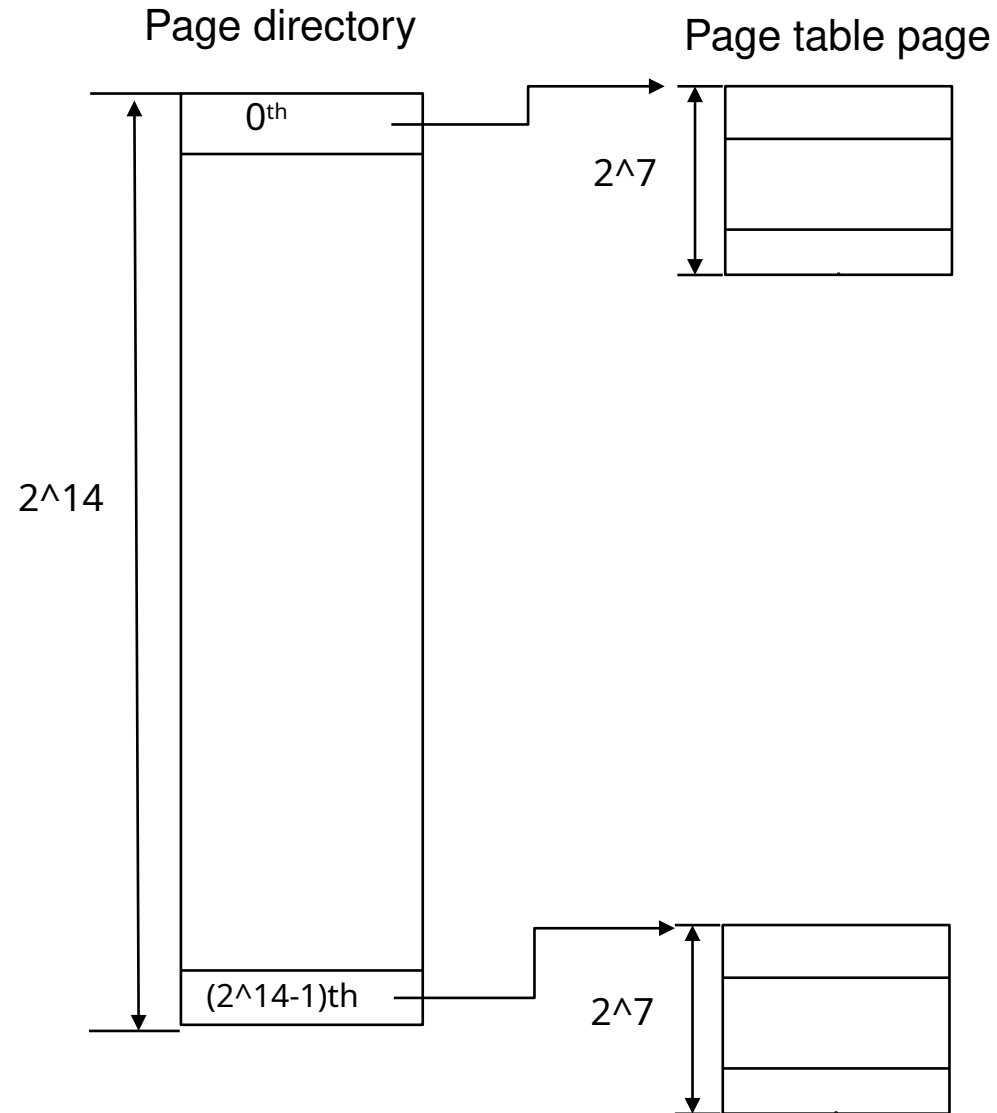


Flag	Detail
Virtual address	30 bit
Page size	512 byte
VPN	21 bit
Offset	9 bit
Page entry per page	128 PTEs (512/4)

Single level page table

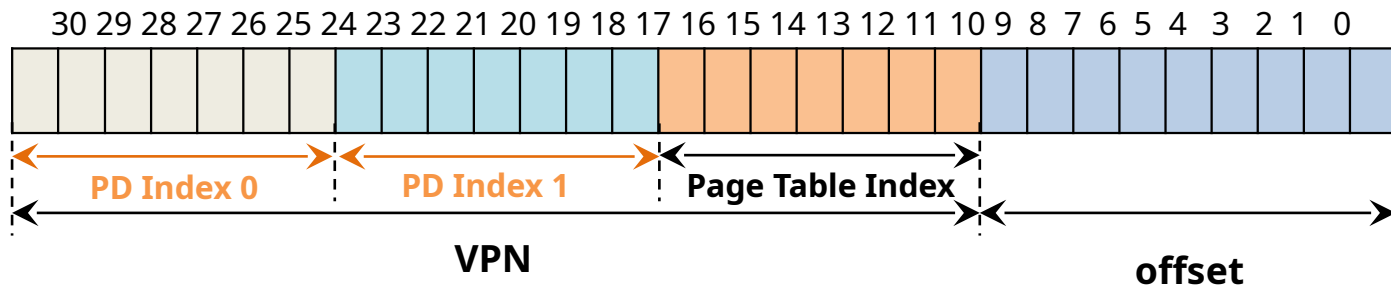


two level page table

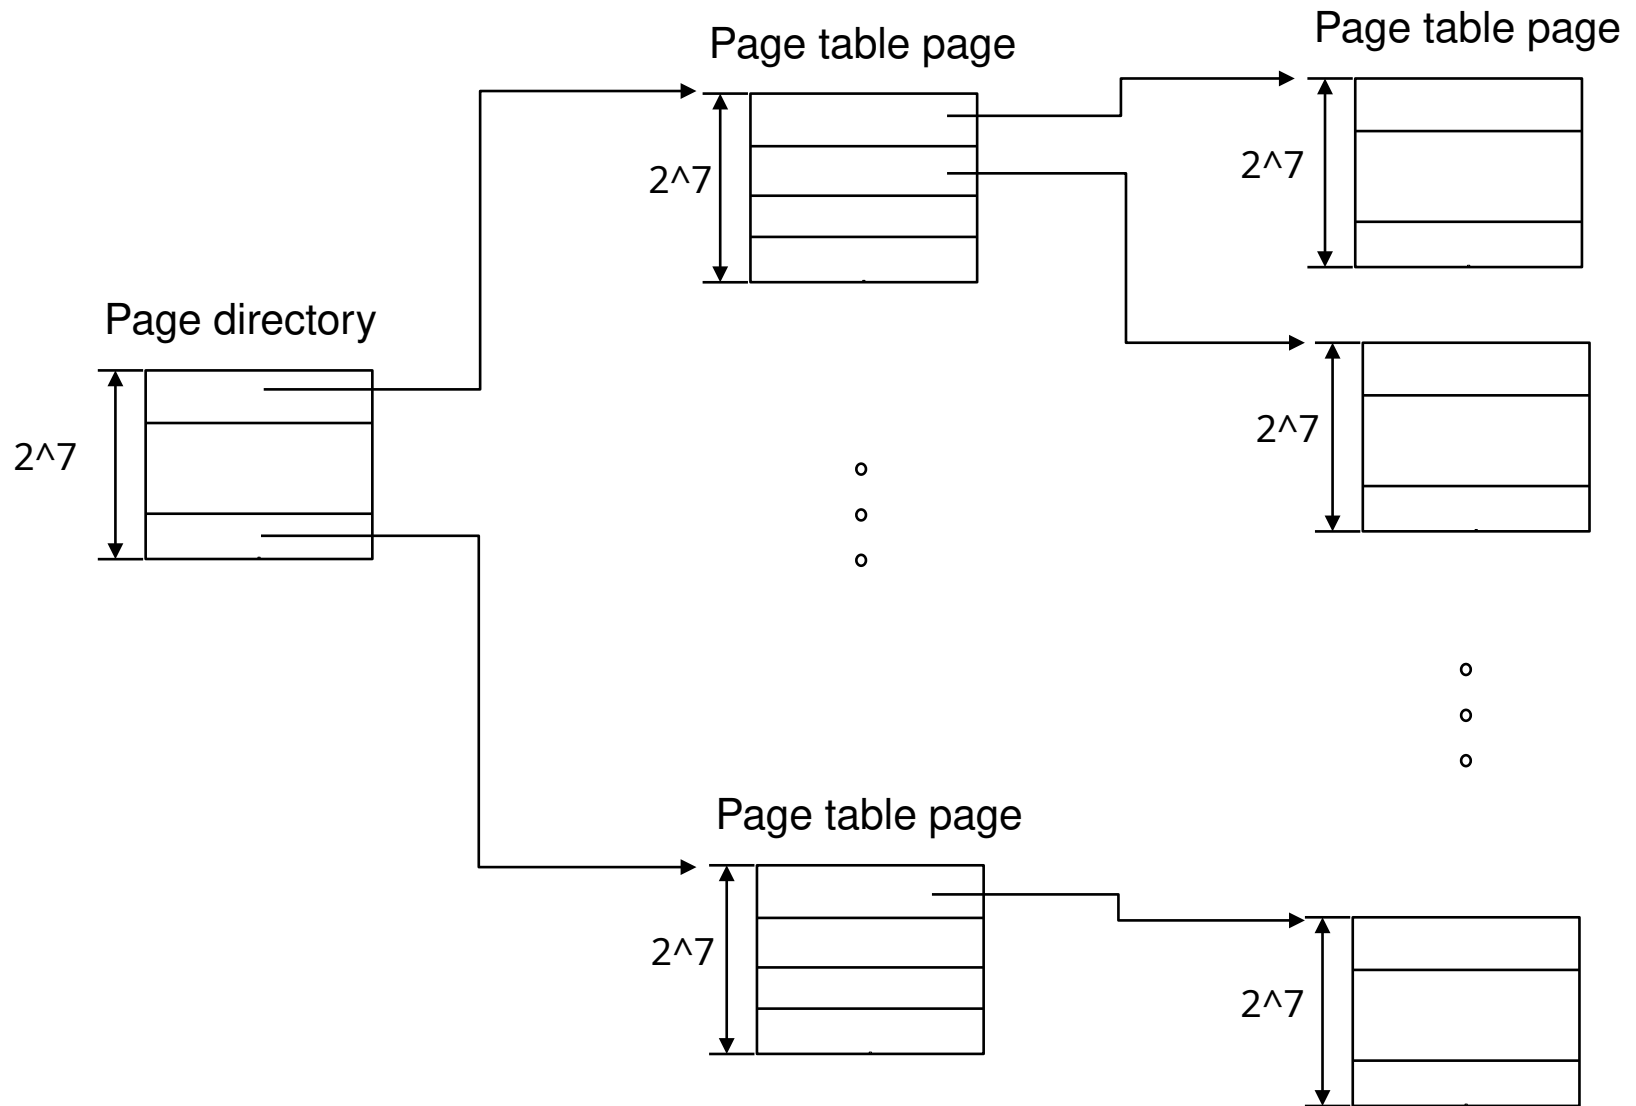


More than Two Level : Page Directory

- ▣ The number of pages in a page directory: $2^{14}/2^7 = 2^7$ (128)
- ▣ Page the page directory



three level page table



Multi-level Page Table Control Flow

```
01:     VPN = (VirtualAddress & VPN_MASK) >> SHIFT
02:     (Success, TlbEntry) = TLB_Lookup(VPN)
03:     if(Success == True)           //TLB Hit
04:         if(CanAccess(TlbEntry.ProtectBits) == True)
05:             Offset = VirtualAddress & OFFSET_MASK
06:             PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
07:             Register = AccessMemory(PhysAddr)
08:         else RaiseException(PROTECTION_FAULT);
09:     else // perform the full multi-level lookup
```

- ◆ (1 lines) extract the virtual page number(VPN)
- ◆ (2 lines) check if the TLB holds the translation for this VPN
- ◆ (5-8 lines) extract the page frame number from the relevant TLB entry, and form the desired physical address and access memory

Multi-level Page Table Control Flow

```
11:         else
12:             PDIndex = (VPN & PD_MASK) >> PD_SHIFT
13:             PDEAddr = PDBR + (PDIndex * sizeof(PDE))
14:             PDE = AccessMemory(PDEAddr)
15:             if(PDE.Valid == False)
16:                 RaiseException(SEGMENTATION_FAULT)
17:             else // PDE is Valid: now fetch PTE from PT
```

- ♦ (11 lines) extract the Page Directory Index(PDIndex)
- ♦ (13 lines) get Page Directory Entry(PDE)
- ♦ (15-17 lines) Check PDE valid flag. If valid flag is true, fetch Page Table entry from Page Table

The Translation Process: Remember the TLB

```
18:     PTIndex = (VPN & PT_MASK) >> PT_SHIFT
19:     PTEAddr = (PDE.PFN << SHIFT) + (PTIndex * sizeof(PTE))
20:     PTE = AccessMemory(PTEAddr)
21:     if(PTE.Valid == False)
22:         RaiseException(SEGMENTATION_FAULT)
23:     else if(CanAccess(PTE.ProtectBits) == False)
24:         RaiseException(PROTECTION_FAULT);
25:     else
26:         TLB_Insert(VPN, PTE.PFN , PTE.ProtectBits)
27:         RetryInstruction()
```

Inverted Page Tables

- ▣ Keeping a single page table that has an entry for each physical page of the system.
- ▣ The entry tells us which process is using this page, and which virtual page of that process maps to this physical page.

Summary

- ▣ Reducing the page table size
- ▣ Per segment page table
- ▣ Multi-level paging
 - ◆ Efficient use of memory
 - ◆ Severe TLB miss