# Operating Systems

**Youjip Won**

KAIST

KAIST OSLab
Operating Systems Laboratory

# 40. Filesystem Implementation

# Overview

- In this chapter, we study very simple file system (vsfs)
  - Basic on-disk structures, access methods, and various policies of vsfs
- We will study…
  - How can we build a simple file system?
  - What structures are needed on the disk?
  - What do they need to track?
  - How are they accessed?

# File system Implementation

- What types of data structures are utilized by the file system?
- How file system organize its data and metadata?
- Understand access methods of a file system.
  - `open()`, `read()`, `write()`, etc.

# Overall Organization

- Let's develop the overall organization of the file system data structure.

- Divide the disk into blocks.
  - Block size is 4 KB.
  - The blocks are addressed from *0 to N -1*.

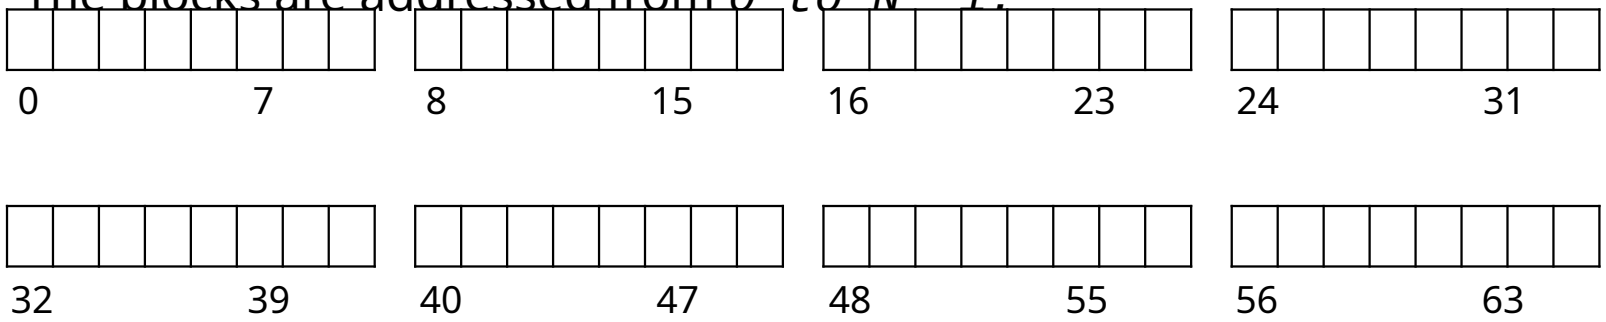

# Data region in file system

- Reserve data region to store user data

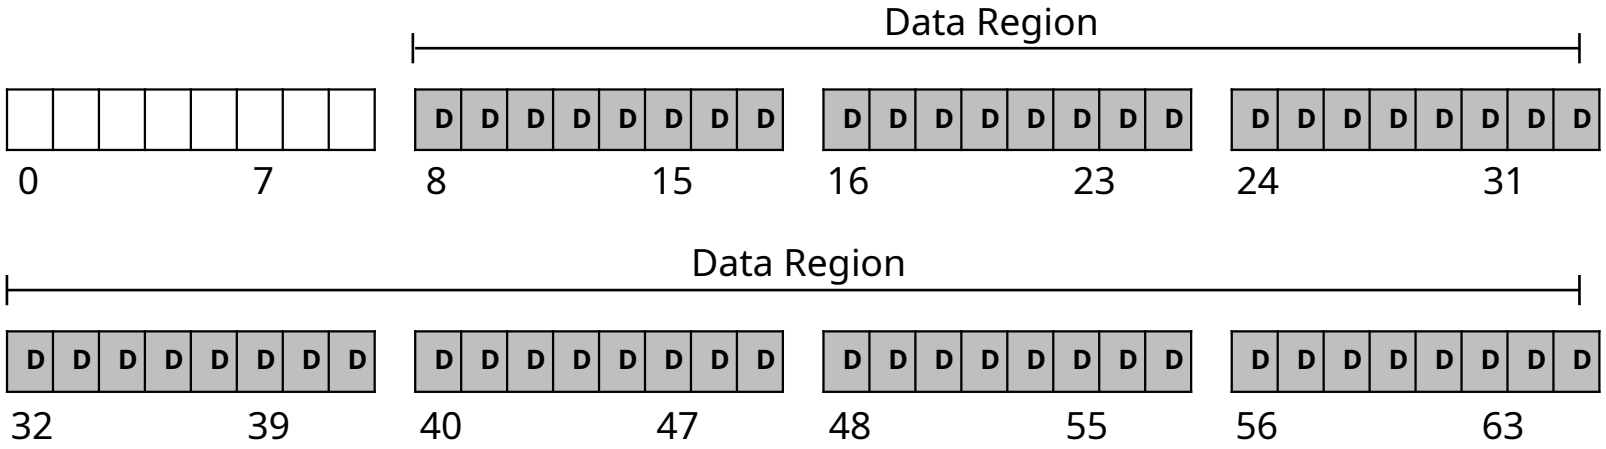


- File system has to track which data block comprise a file, the size of the file, its owner, etc.

**How we store these inodes in file system?**

# Inode table in file system

- Reserve some space for inode table
  - This holds an array of on-disk inodes.
  - Ex) inode tables : 3 ~ 7, inode size : 256 bytes
    - 4-KB block can hold 16 inodes.
    - The file system contains 80 inodes. (maximum number of files)

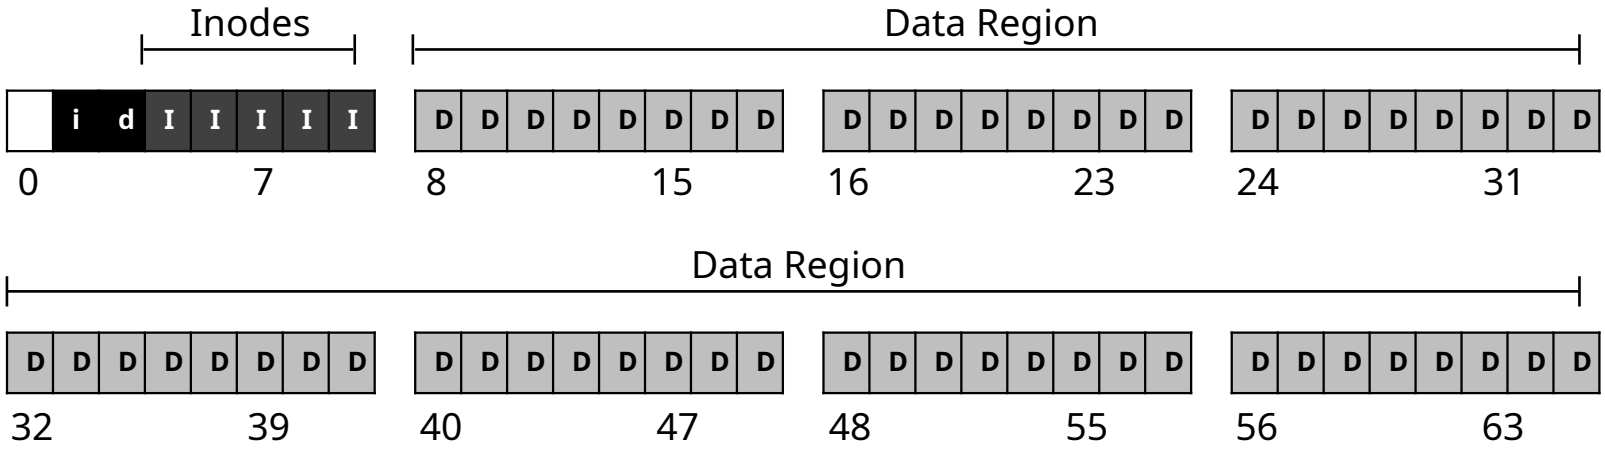

# allocation structures

- This is to track whether inodes or data blocks are free or allocated.
- Use bitmap, each bit indicates free(0) or in-use(1)
  - data bitmap: for data region for data region
  - inode bitmap: for inode table

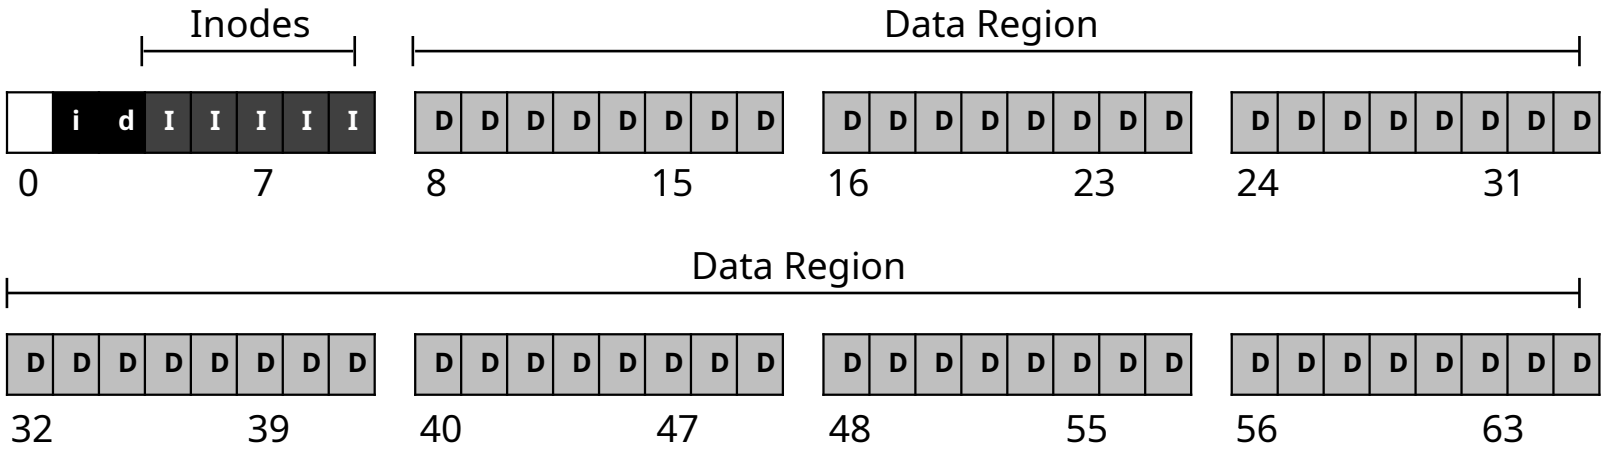

# super block

- Super block contains this information for particular file system
  - Ex) The number of inodes, begin location of inode table. etc

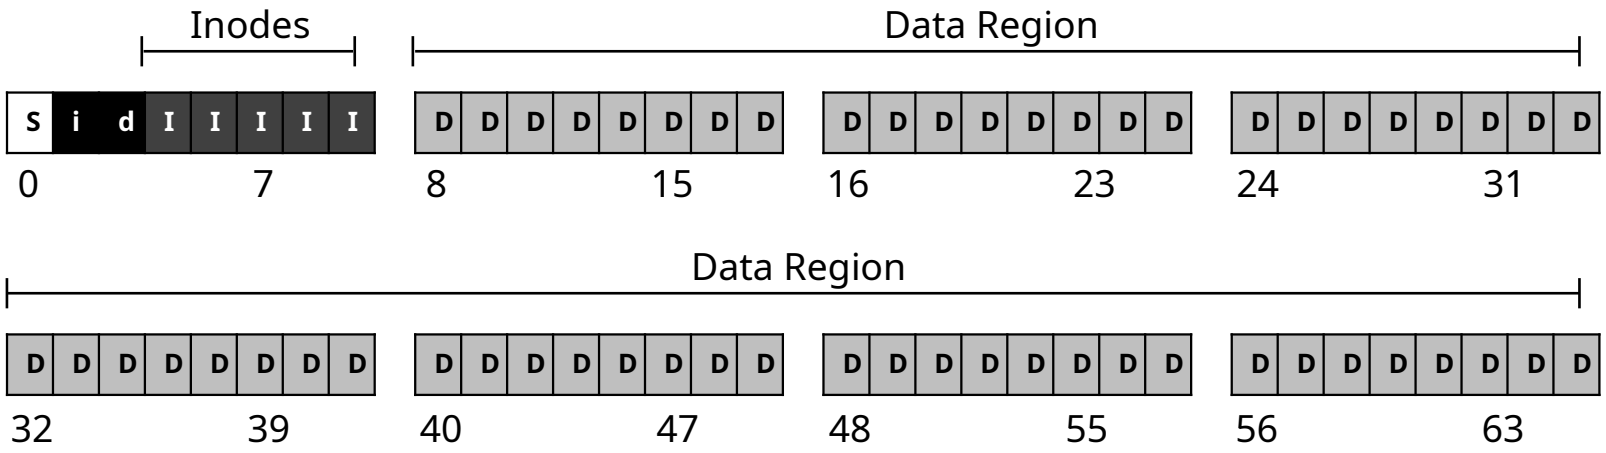


- Thus, when mounting a file system, OS will read the superblock first, to initialize various information.

# File Organization: The inode

- Each inode is referred to by inode number.
  - by inode number, File system calculate where the inode is on the disk.
  - Ex) inode number: 32
    - Calculate the offset into the inode region (32 x sizeof(inode) (256 bytes) = 8192
    - Add start address of the inode table(12 KB) + inode region(8 KB) = 20 KB

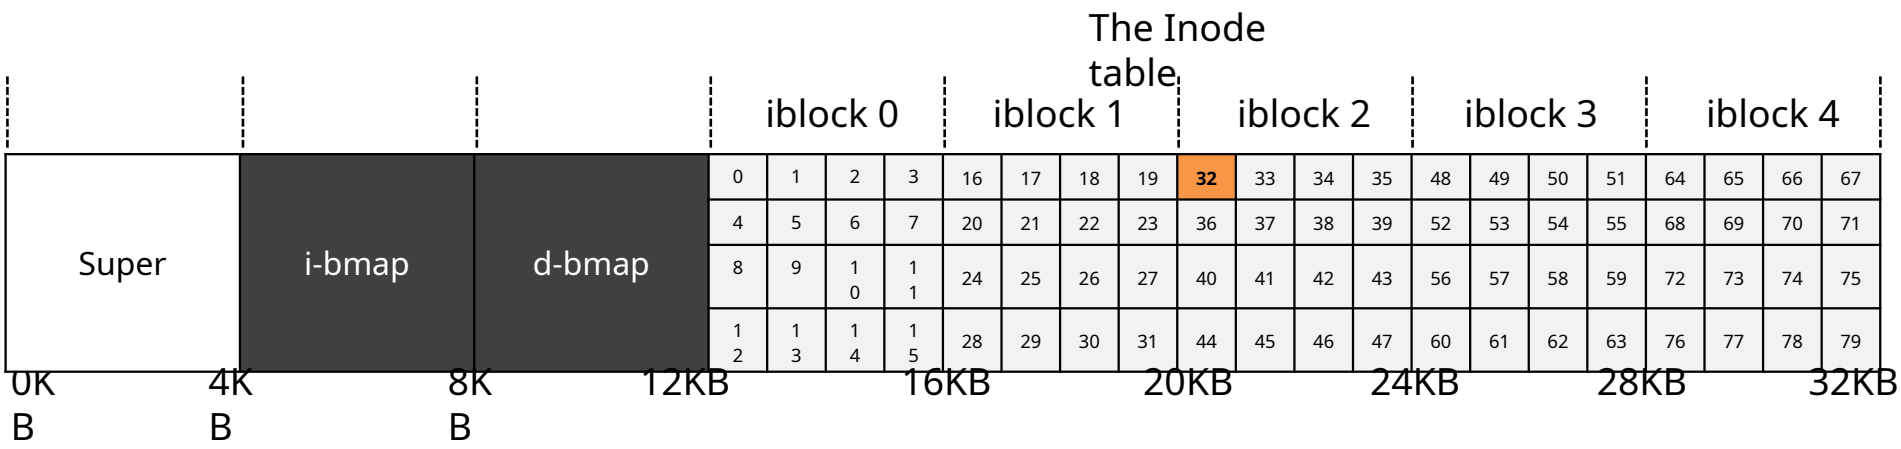

# File Structure: Indexed Allocation

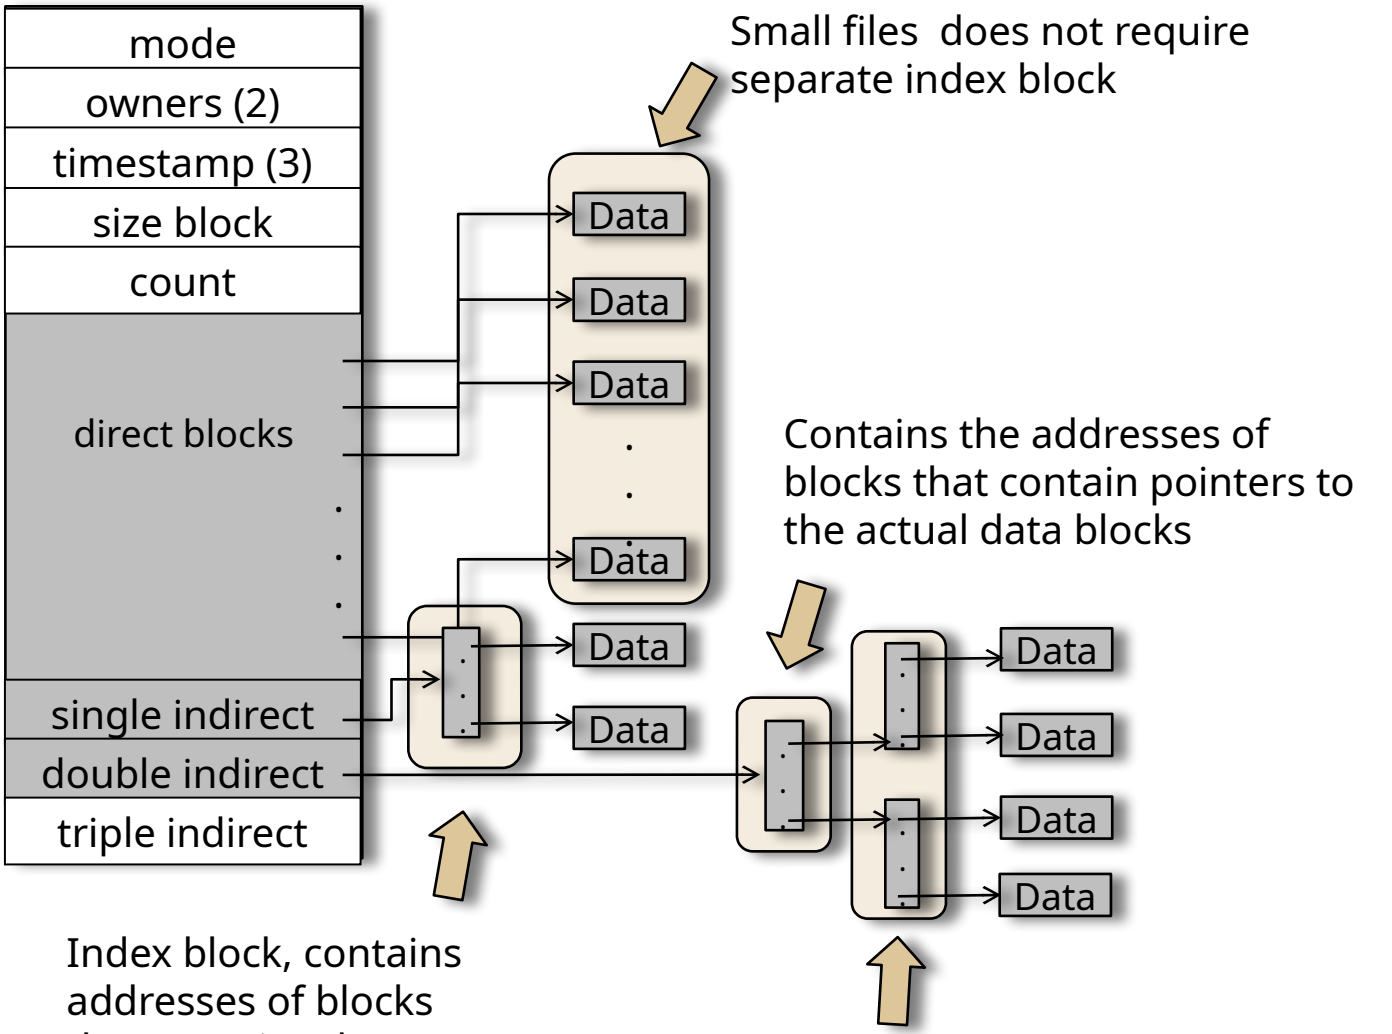

# Directory Structure

**VSFS**

| inum | reclen | strlen | name |
|---|---|---|---|
| 5 | 4 | 2 | . |
| 2 | 4 | 3 | .. |
| 12 | 4 | 4 | foo |
| 13 | 4 | 4 | bar |
| 24 | 8 | 7 | foobar |

**EXT4**

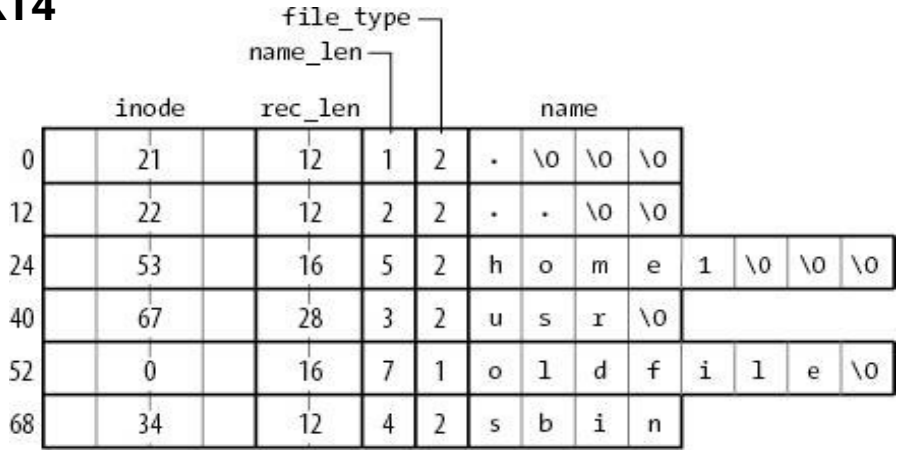

# File Read

| | data bitmap | inode bitmap | root inode | foo inode | bar inode | root data | foo data | bar data[0] | bar data[1] | bar data[2] |
|---|---|---|---|---|---|---|---|---|---|---|
| open(bar) | | | read | read | read | read | read | | | |
| read() | | | | | read<br>write | | | read | | |
| read() | | | | | read<br>write | | | | read | |
| read() | | | | | read<br>write | | | | | read |

# File Creation

| | data bitmap | inode bitmap | root inode | foo inode | bar inode | root data | foo data | bar data[0] | bar data[1] | bar data[2] |
|---|---|---|---|---|---|---|---|---|---|---|
| create (/foo/bar) | | read<br>write | read | read | read<br>write | read | read<br>write | | | |
| write() | read<br>write | | | | read<br>write | | | write | | |
| | | | | | ... | | | | | |

# File Creation (Cont.)

| | data bitmap | inode bitmap | root inode | foo inode | bar inode | root data | foo data | bar data[0] | bar data[1] | bar data[2] |
|---|---|---|---|---|---|---|---|---|---|---|
| | | | | | ... | | | | | |
| write() | read<br>write | | | | read<br>write | | | | write | |
| write() | read<br>write | | | | read<br>write | | | | write | |

# Caching and Buffering

- Reading and writing can very IO intensive.
  - File open: two IO for each directory component and one read for the data.
- Buffer Cache
  - cache the disk blocks to reduce the IO.
  - LRU replacement
  - Static partitioning: 10% of DRAM, inefficient usage

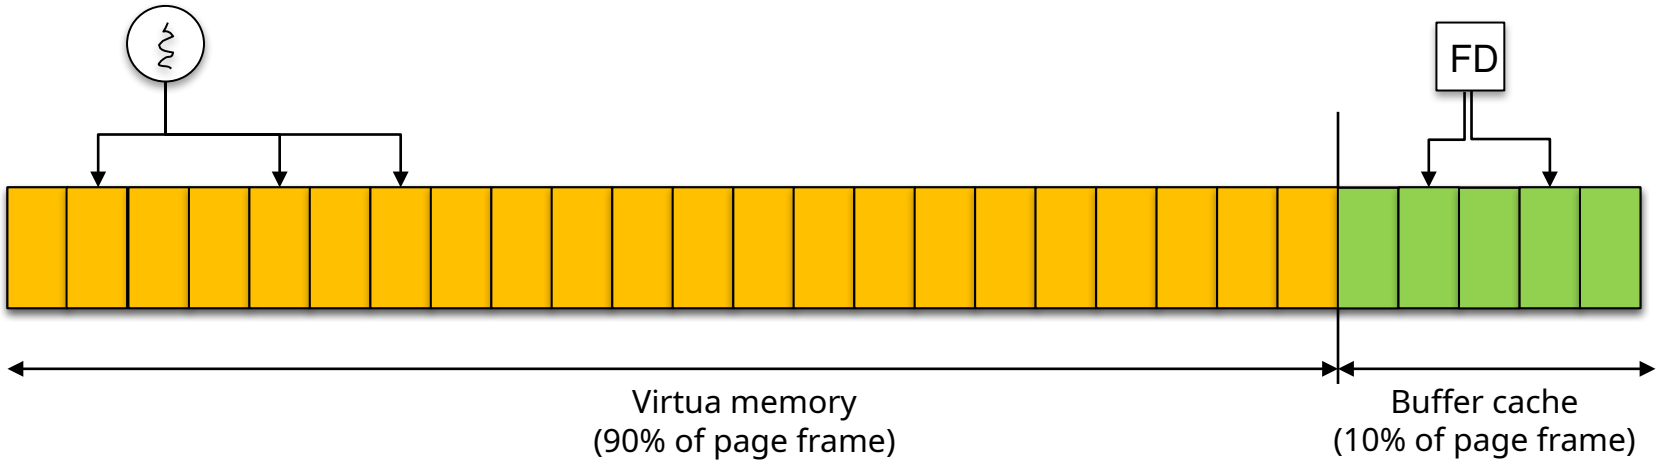

# Caching and Buffering

- Page Cache
  - Merge virtual memory and buffer cache
  - A physical page frame can host either a page in the process address space or a file block.
    - Process uses page table to map a virtual page to a page frame.
    - A file IO uses “`address_space`”(Linux) to map a file block to a physical page frame.
  - Dynamic partitioning

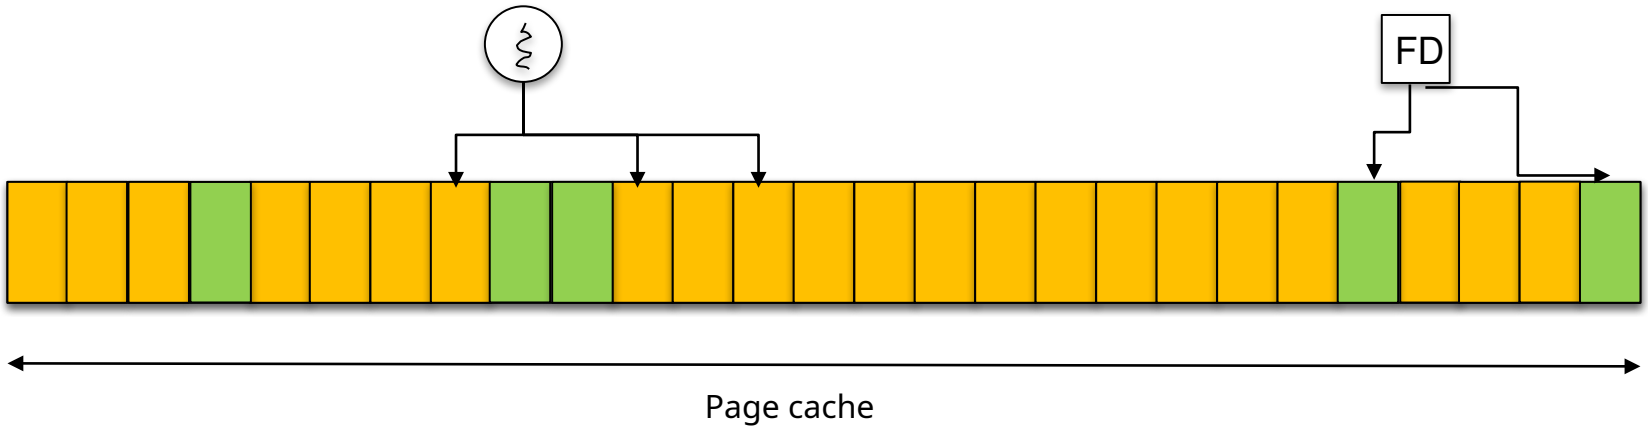

# Summary

- Requirements for building filesystem
  - File information: inode
  - File structure: indexed file
  - Directory (name→inode-number): array of <inode #, name>'s
  - Free block information: Bitmap
- All are flexible.