

# What is a Distributed System?

*A distributed system is one where a machine I've never heard of can cause my program to fail.*

— Leslie Lamport

Definition:

More than 1 machine working together to solve a problem

Examples:

- client/server: web server and web client
- cluster: page rank computation

Other courses:

- **CS 640**: Networking
- **CS 739**: Distributed Systems

# Why Go Distributed?

More computing power

More storage capacity

Fault tolerance

Data sharing

# New Challenges

**System failure:** need to worry about partial failure

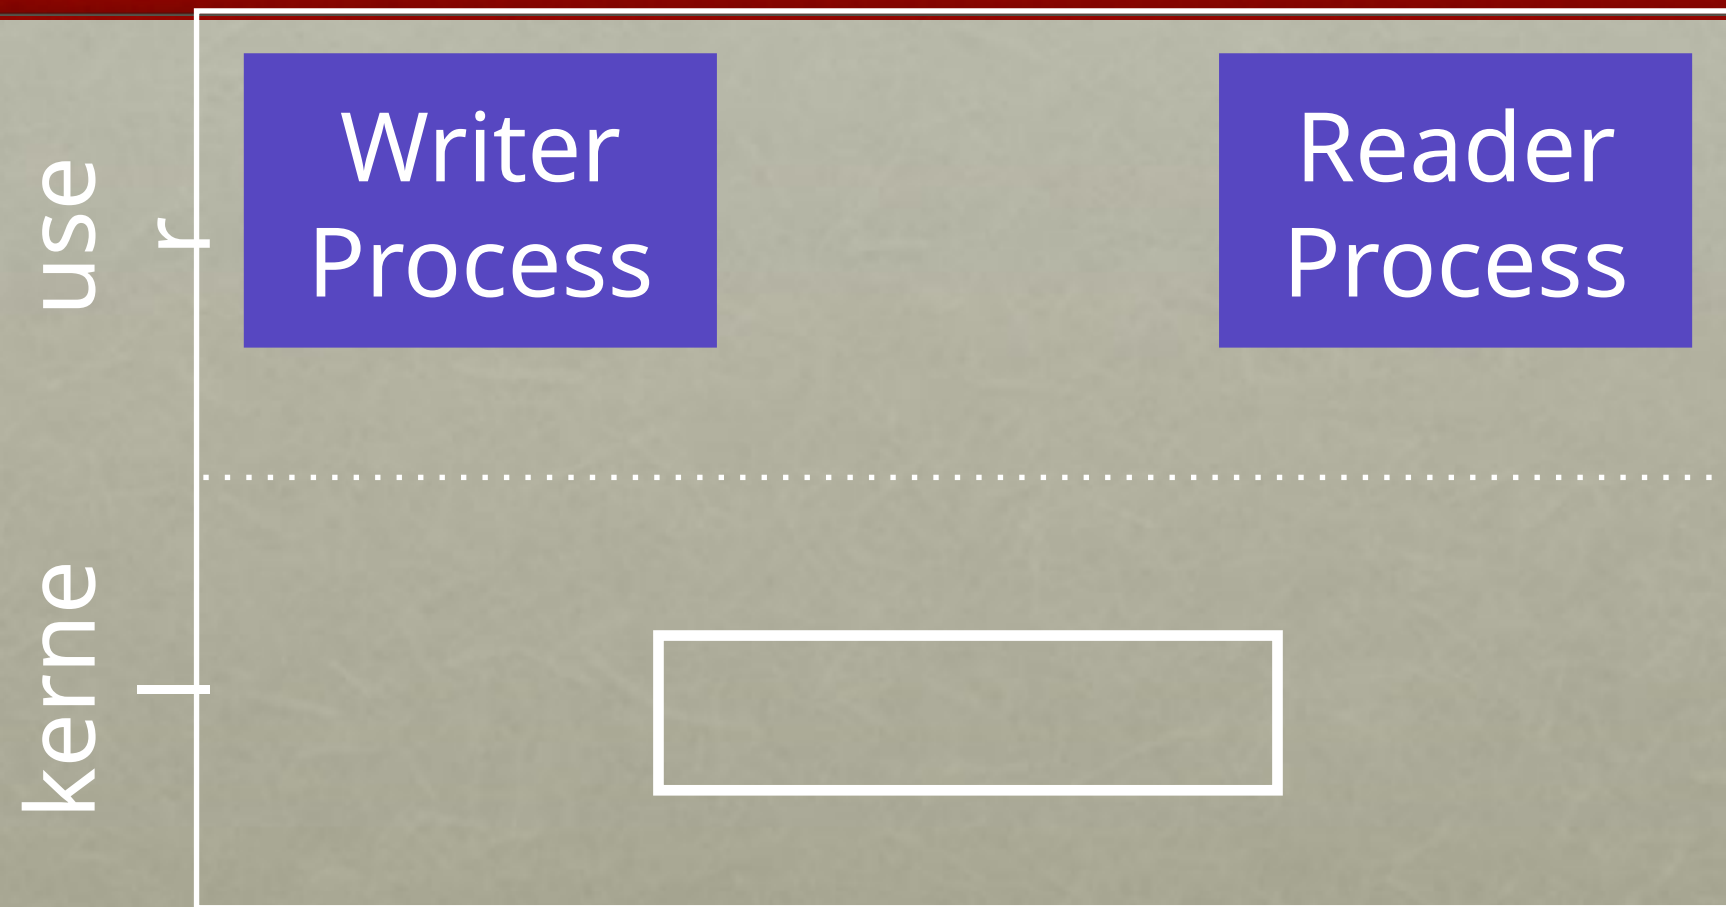
**Communication failure:** links unreliable

- bit errors
- packet loss
- node/link failure

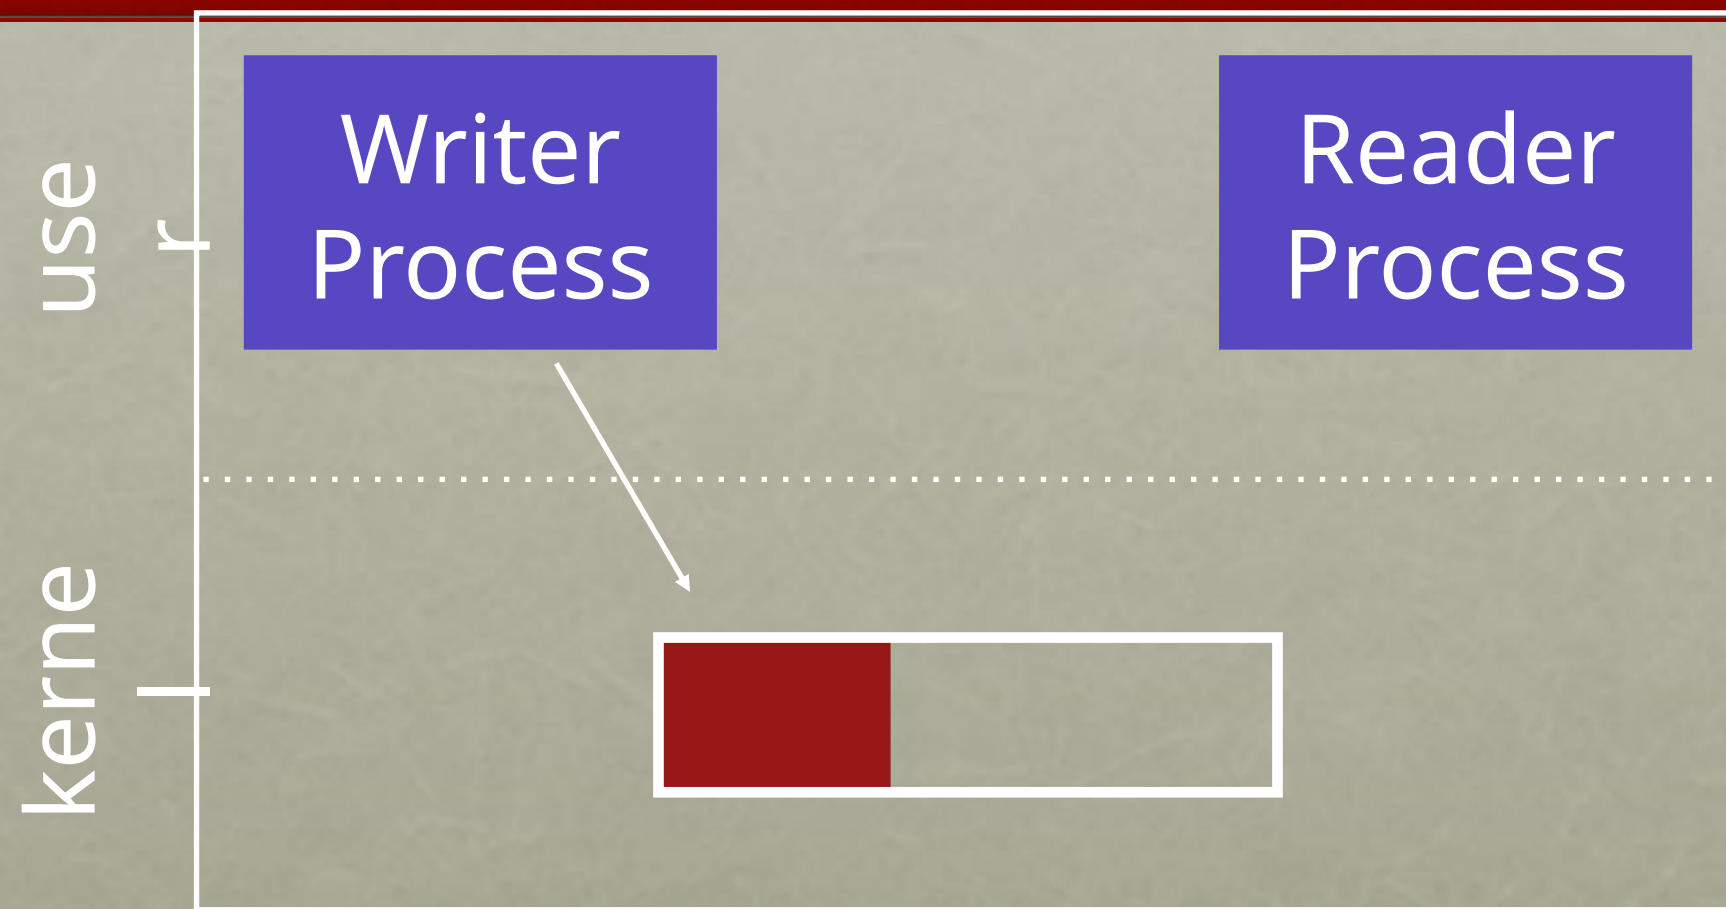
Motivation example:

Why are network sockets less reliable than pipes?

# Pipe

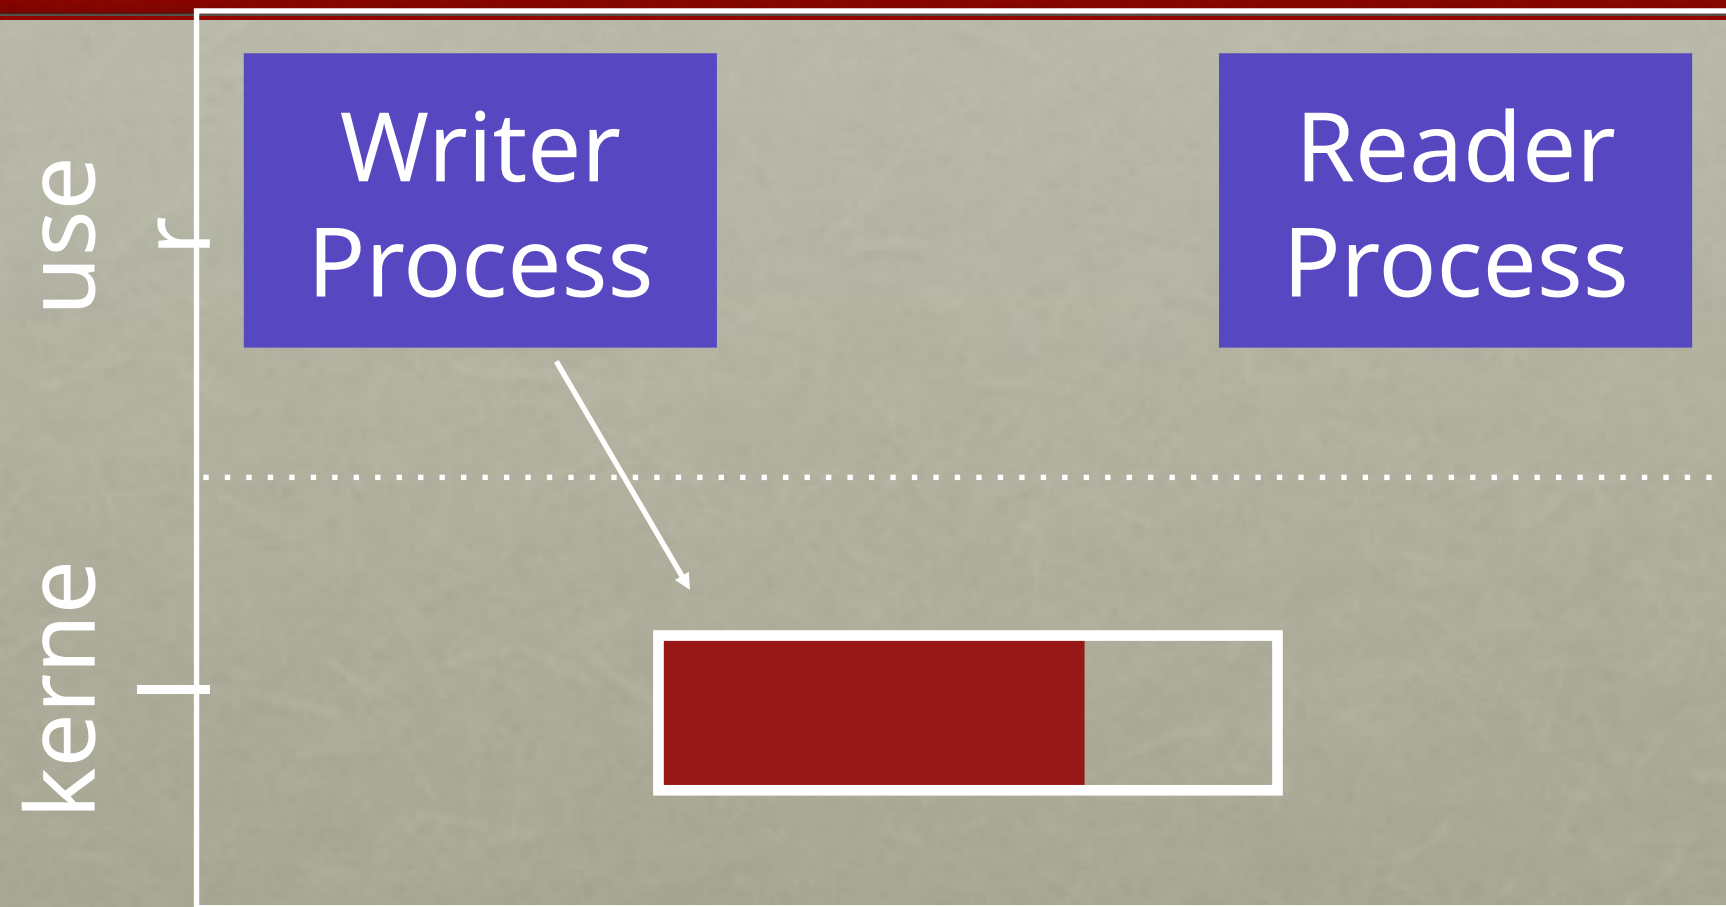


# Pipe

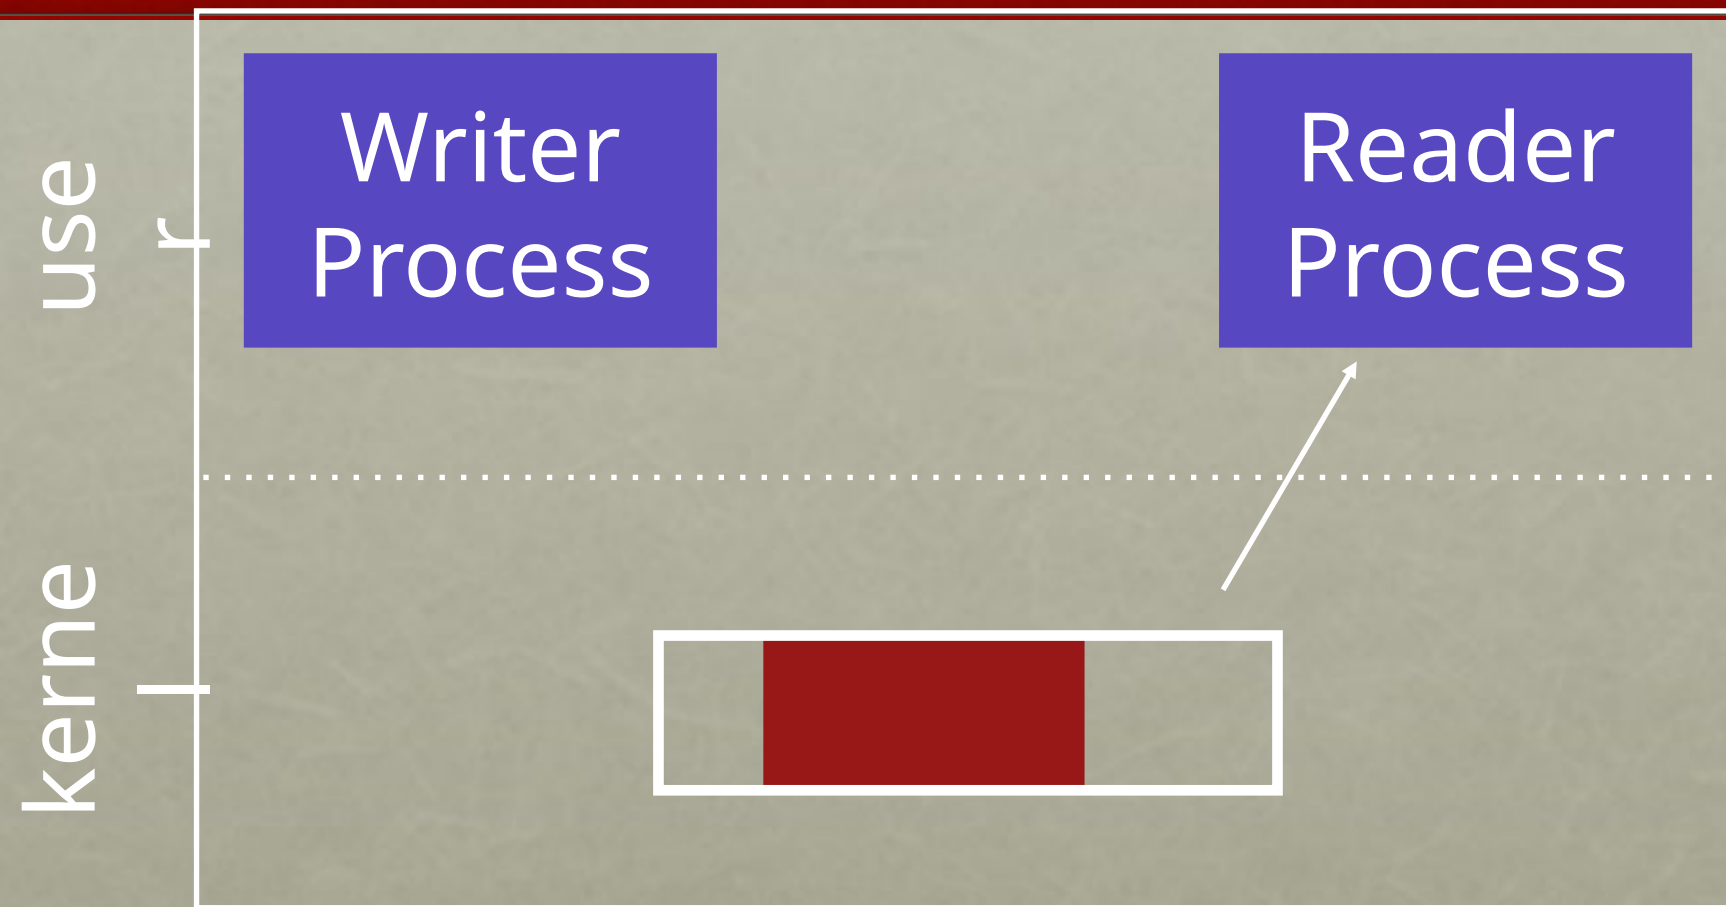




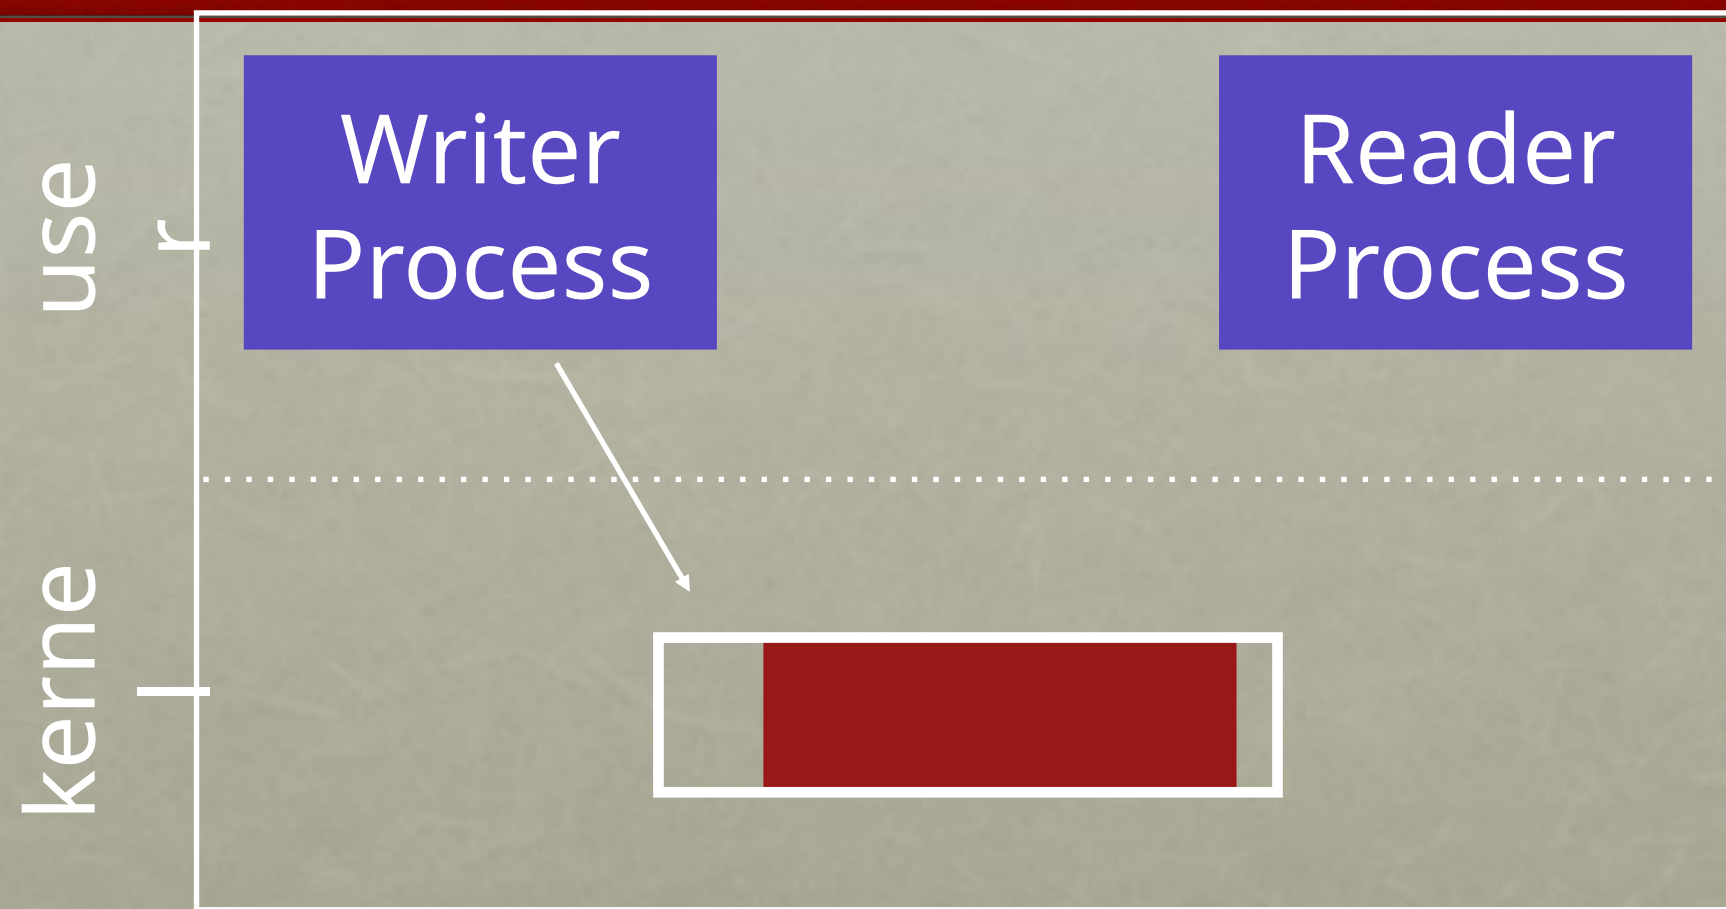
# Pipe



# Pipe

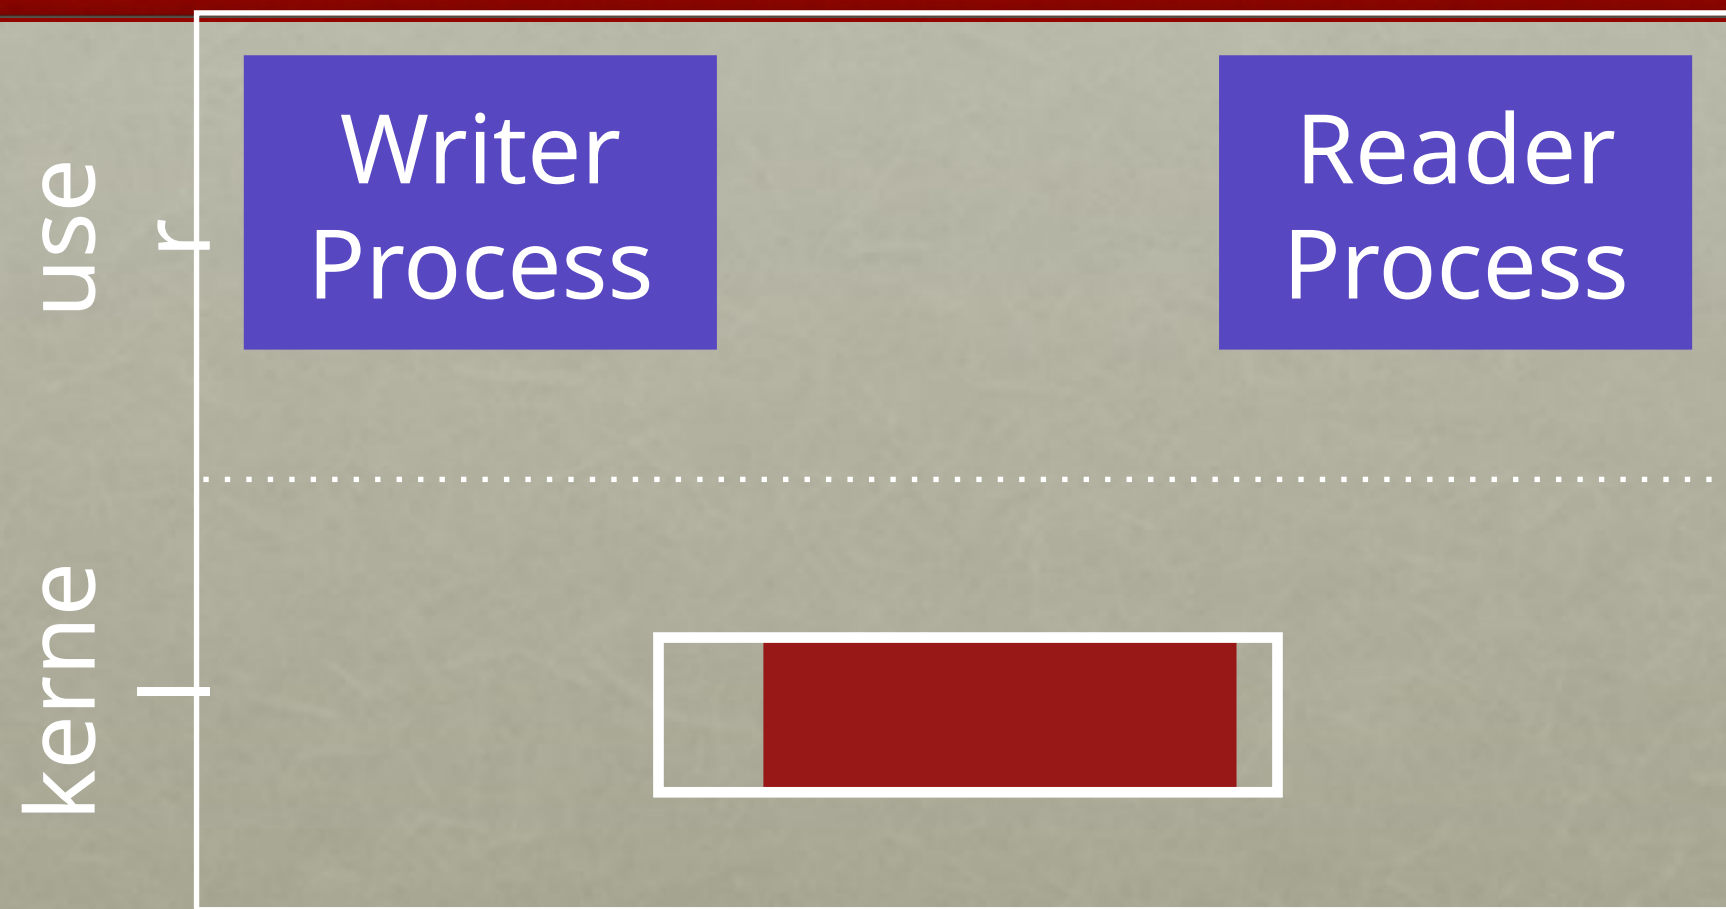


# Pipe

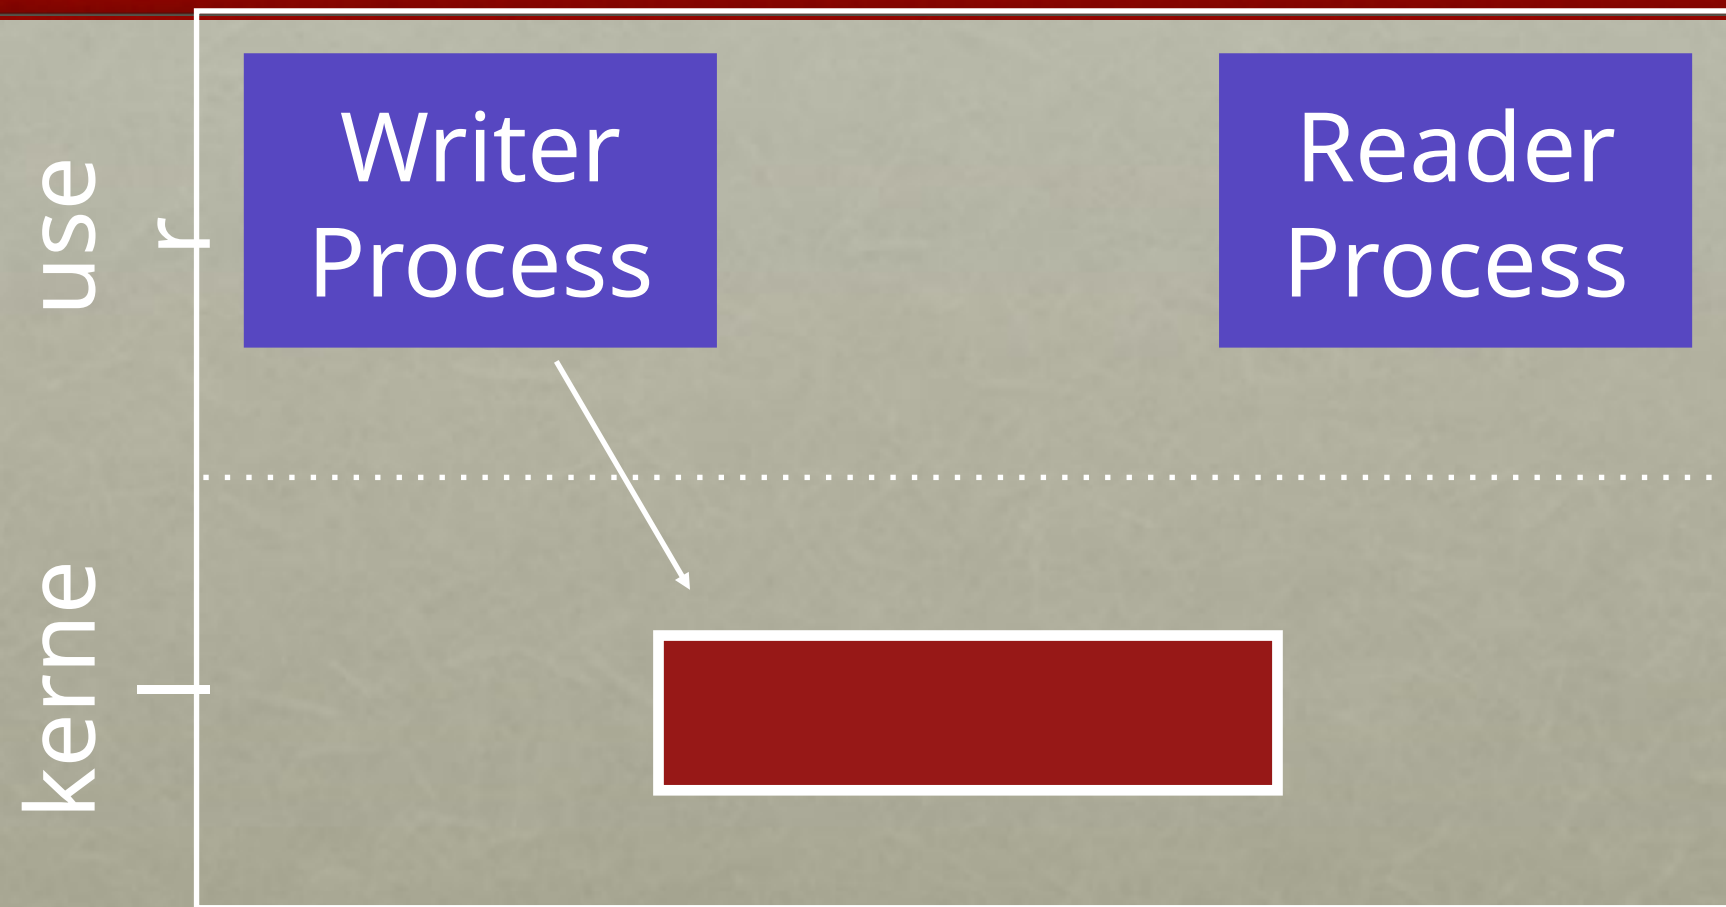




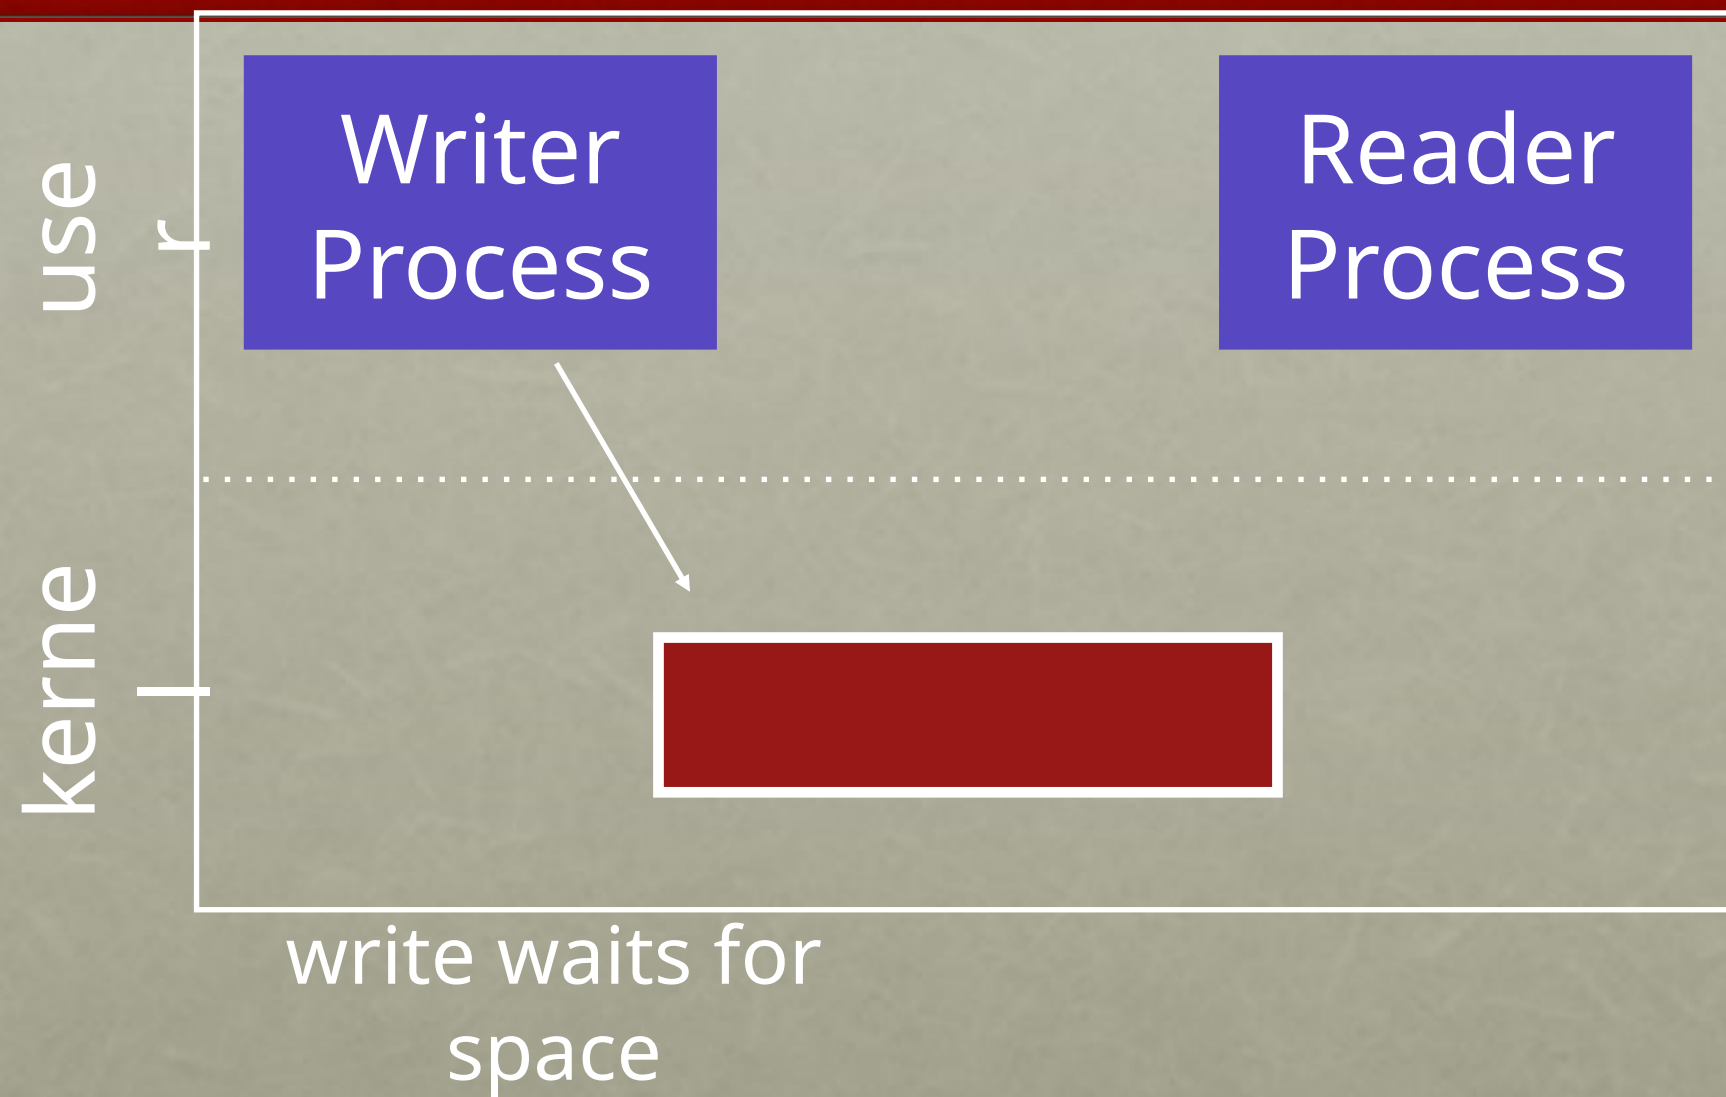
# Pipe



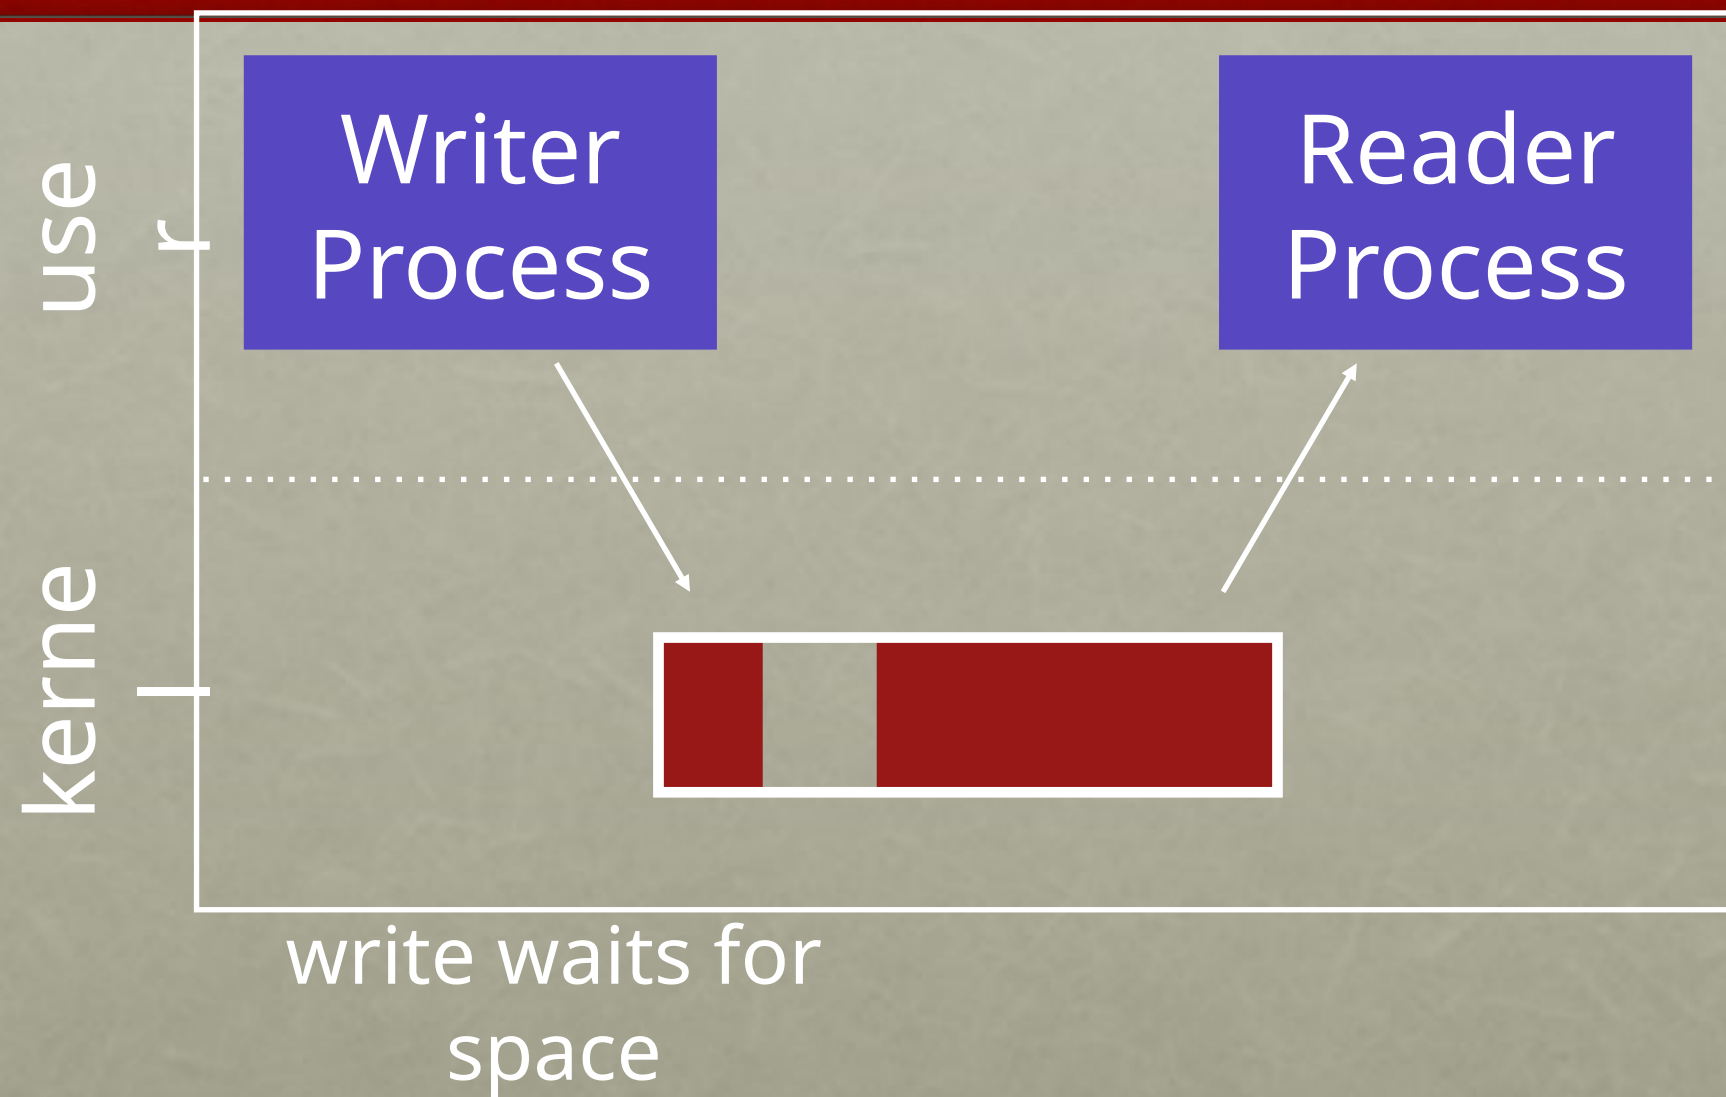
# Pipe



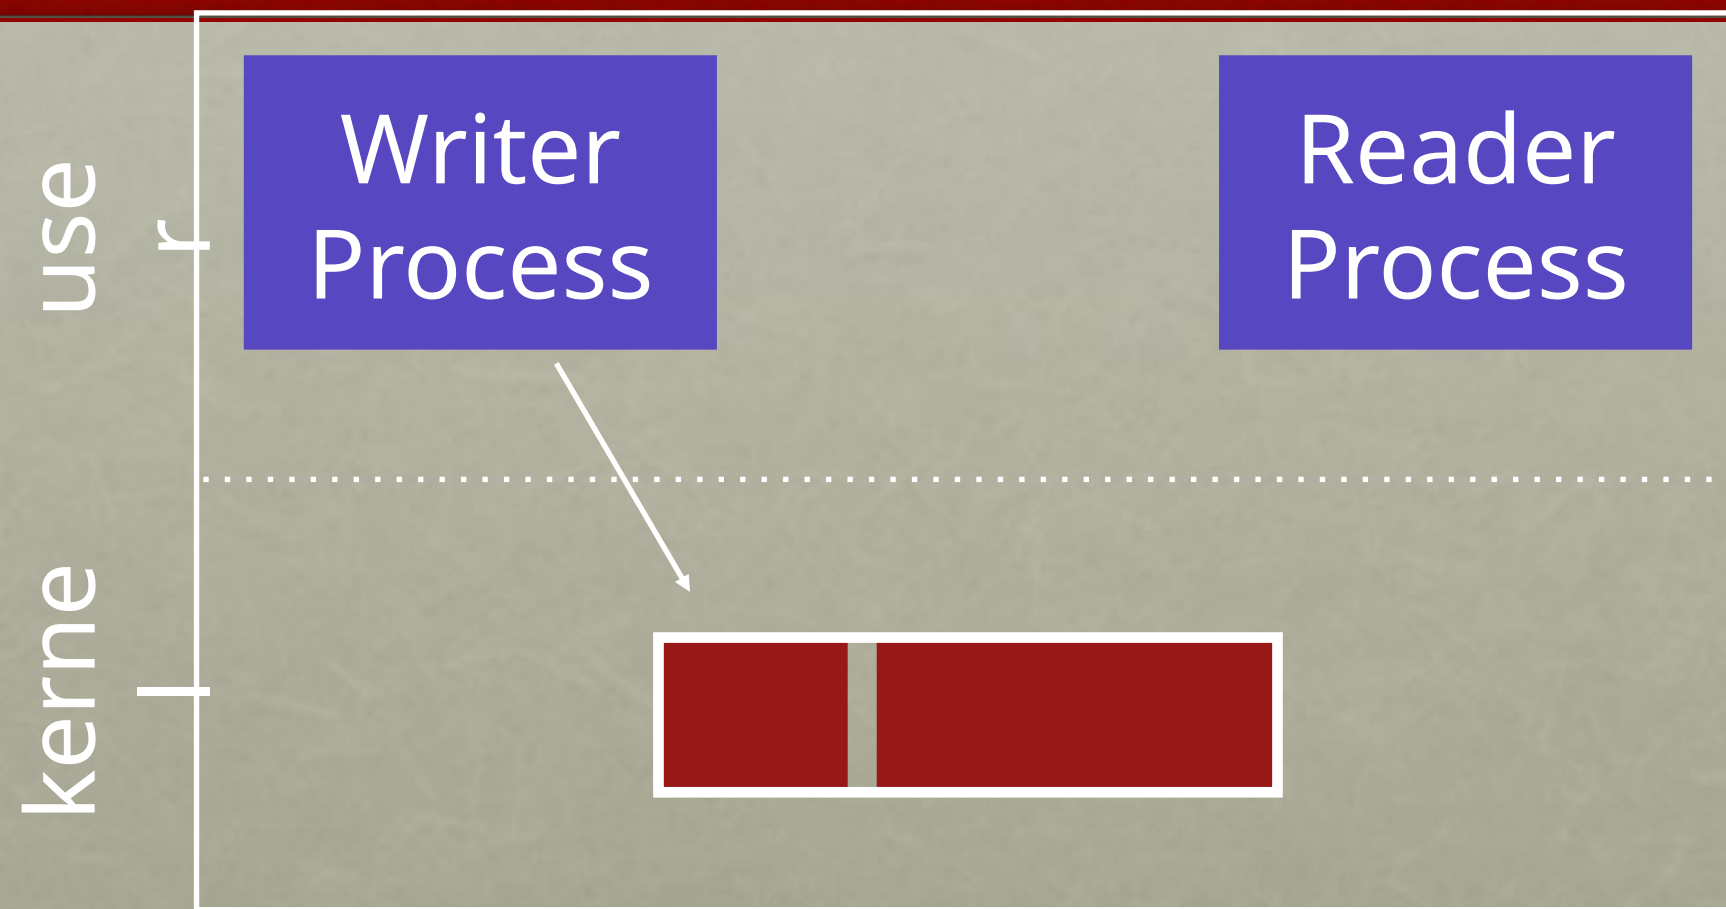
# Pipe



# Pipe




# Pipe

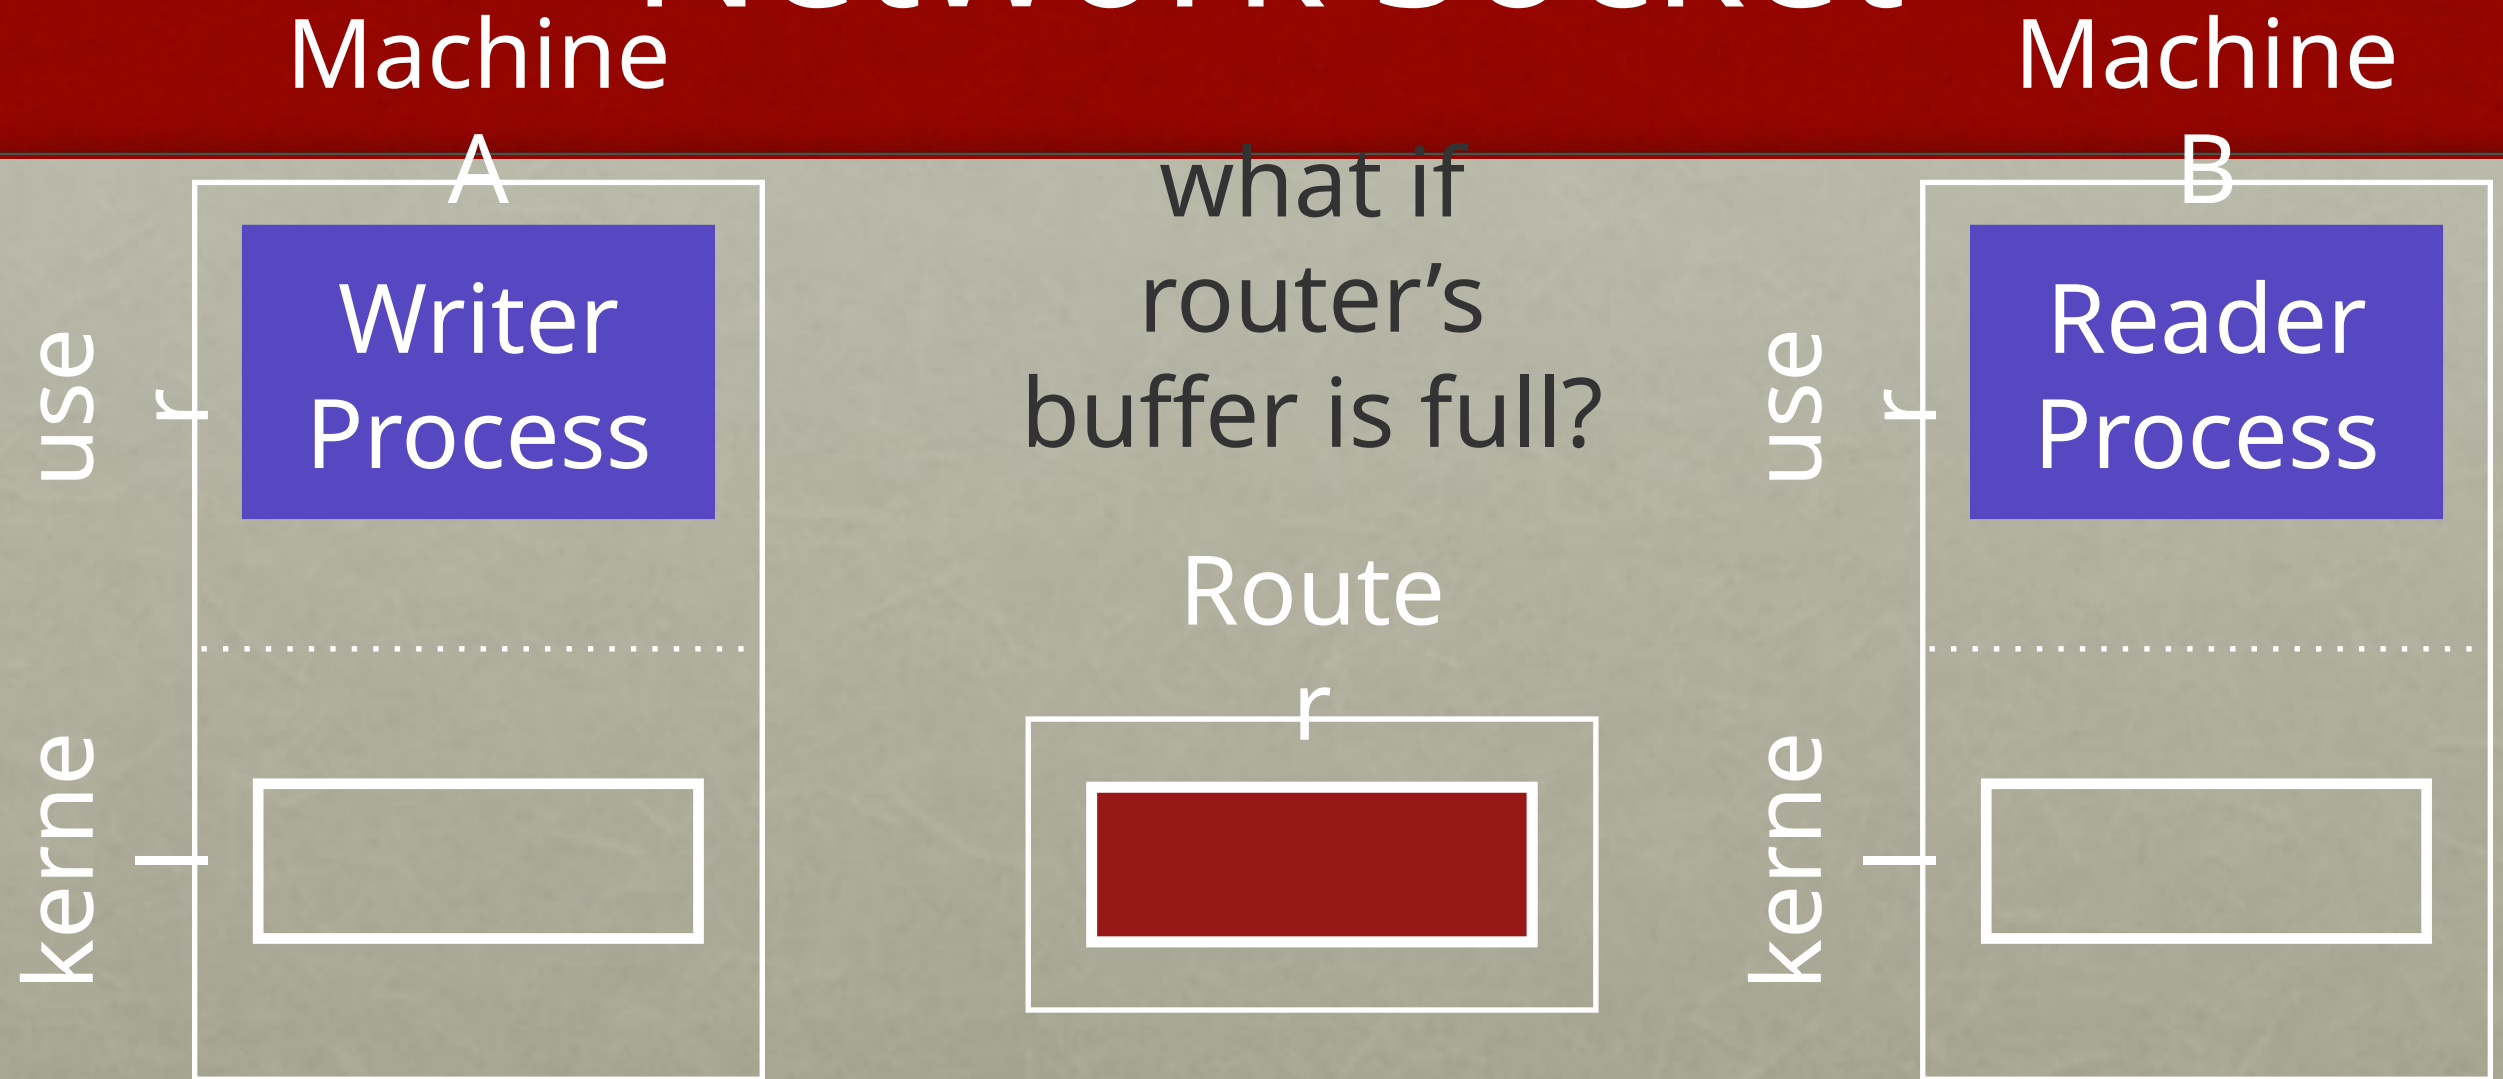




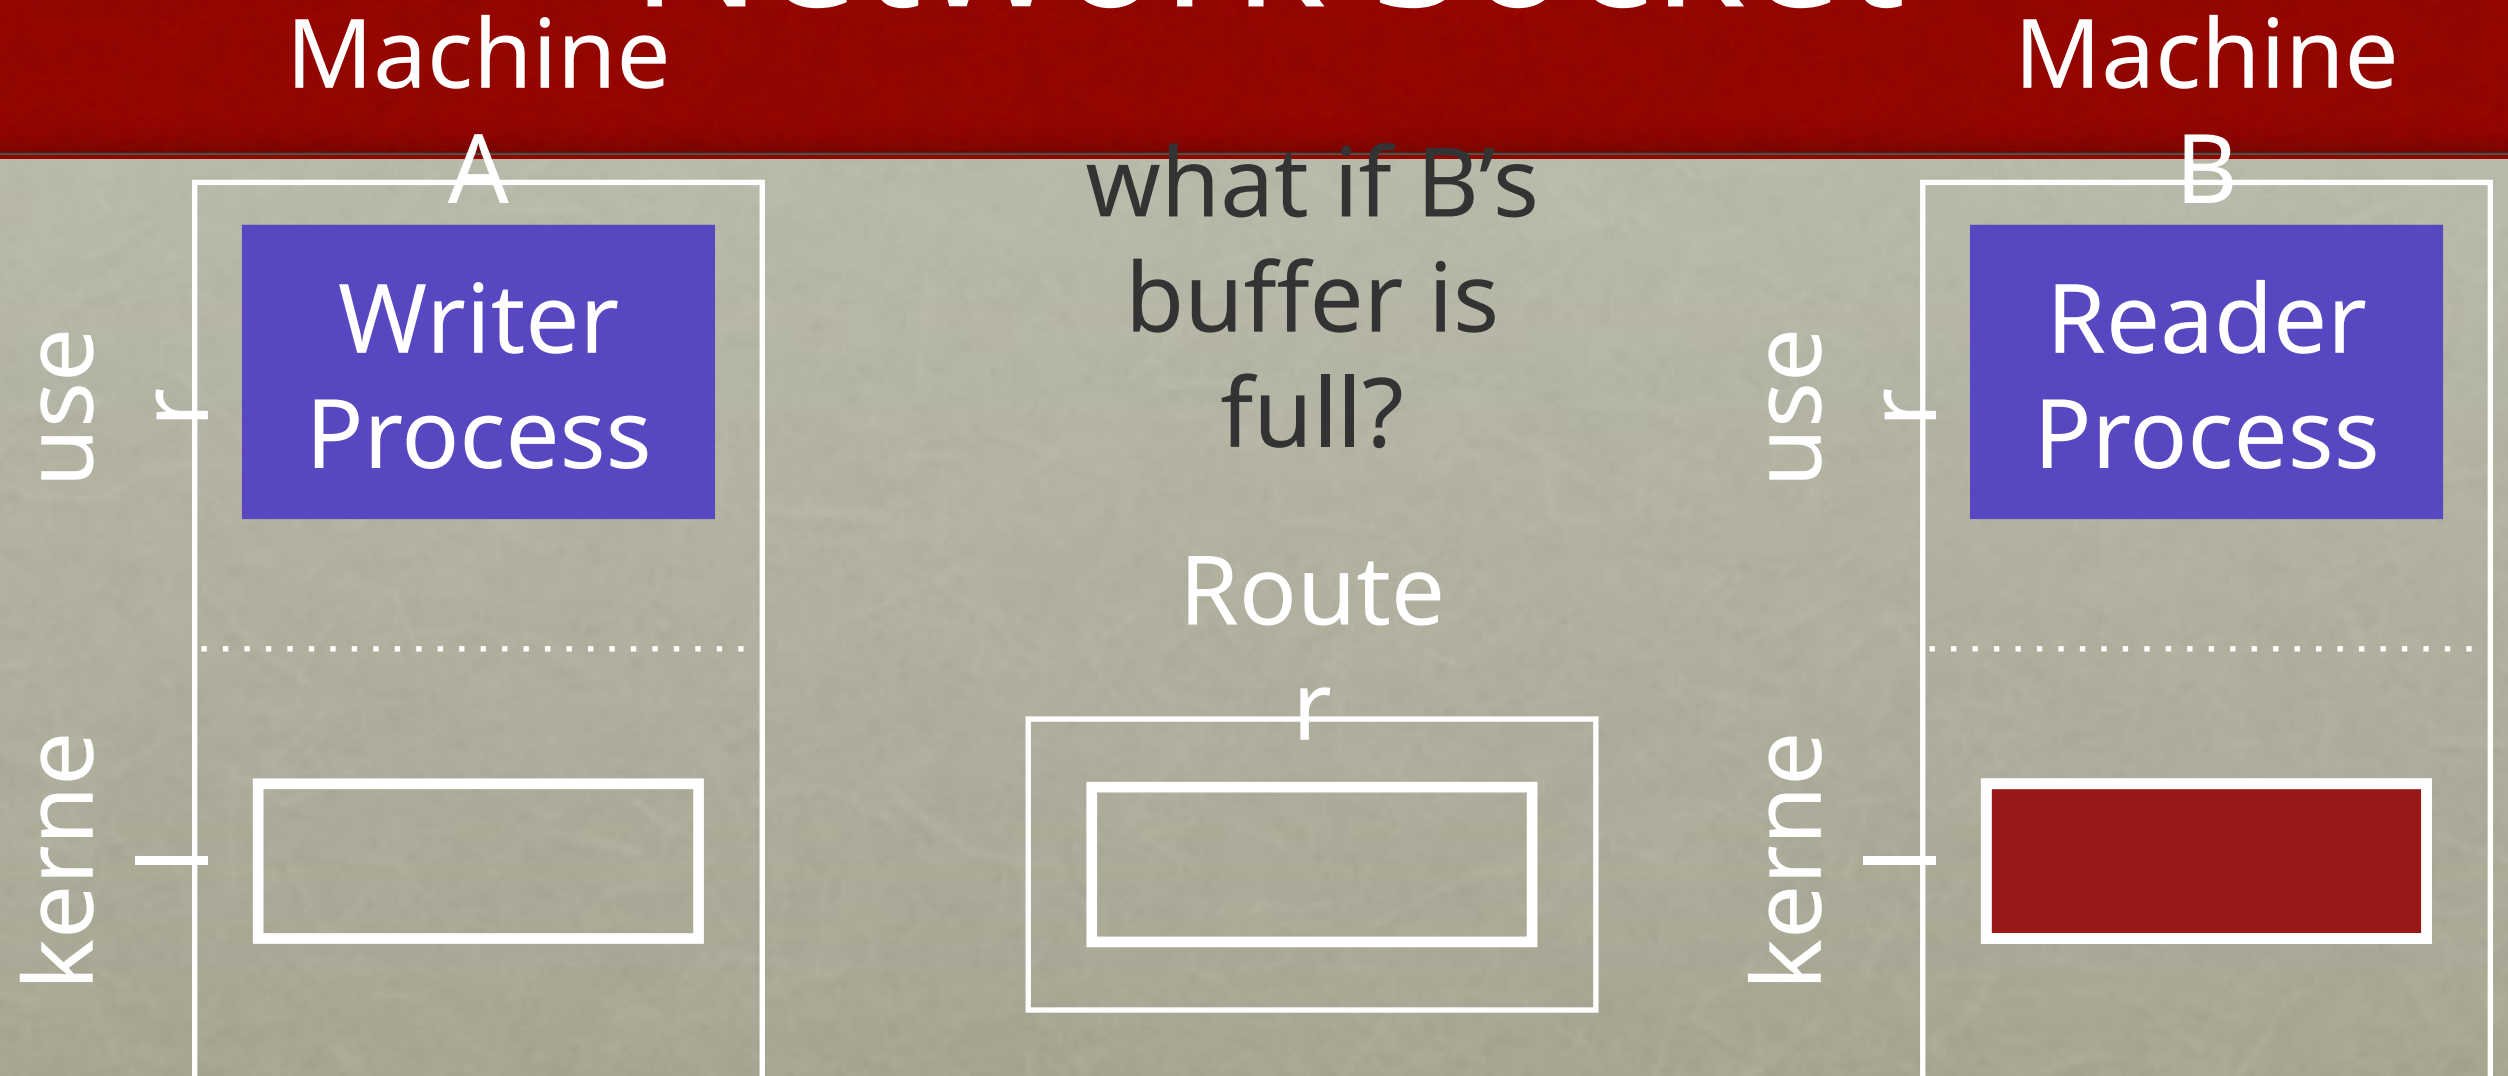
# Network Socket



# Network Socket




# Network Socket



# Network Socket

Machine



# Communication Overview

Raw messages: UDP

Reliable messages: TCP

Remote procedure call: RPC



# Raw Messages: UDP

UDP : User Datagram Protocol

API:

- reads and writes over socket file descriptors
- messages sent from/to ports to target a process on machine

Provide minimal reliability features:

- messages may be lost
- messages may be reordered
- messages may be duplicated
- only protection: checksums to ensure data not corrupted

# Raw Messages: UDP

## Advantages

- Lightweight
- Some applications make better reliability decisions themselves (e.g., video conferencing programs)

## Disadvantages

- More difficult to write applications correctly

# Reliable Messages: Layering strategy

TCP: Transmission Control Protocol

Using software, build reliable, logical connections over unreliable connections

Techniques:

- acknowledgment (ACK)


# Technique #1: ACK



Sender knows message was received



# ACK

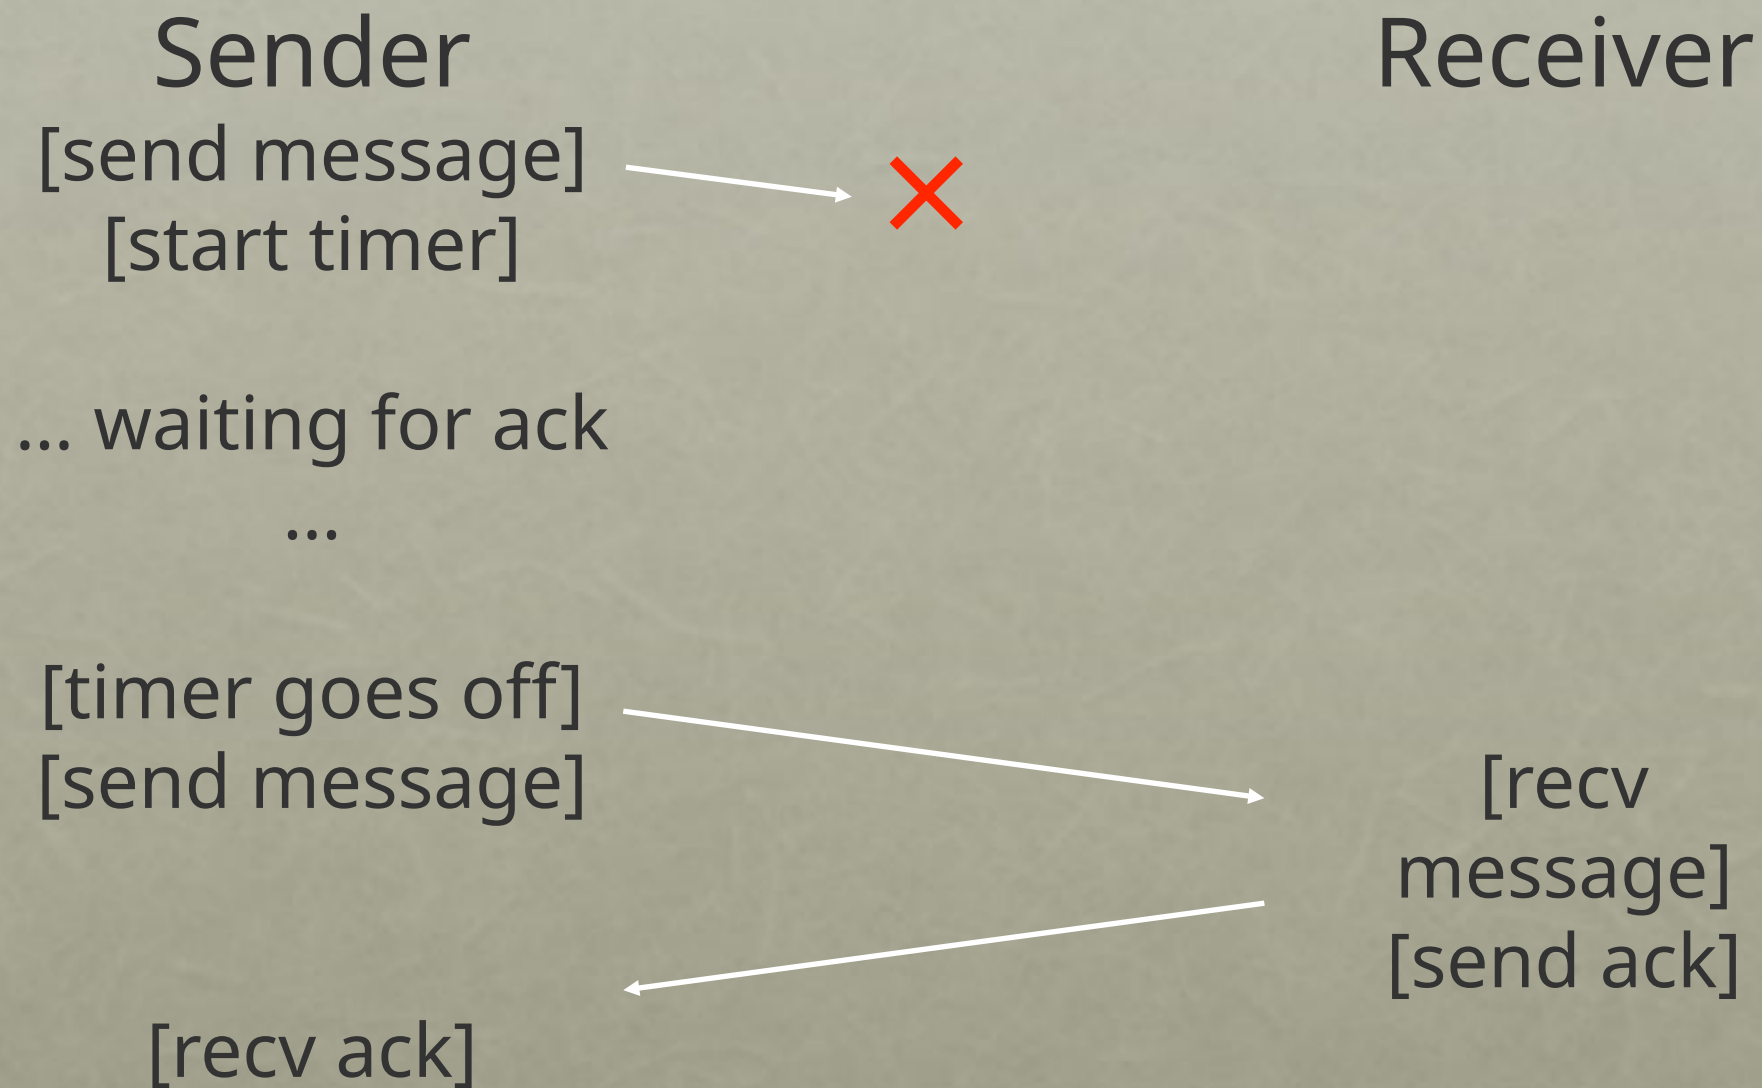


Sender doesn't receive ACK...

What to do?



# Technique #2: Timeout



# Lost ACK: Issue 1

How long to wait?

Too long?

- System feels unresponsive

Too short?

- Messages needlessly re-sent
- Messages may have been dropped due to overloaded server. Resending makes overload worse!

# Lost Ack: Issue 1

How long to wait?

One strategy: be adaptive

Adjust time based on how long acks usually take

For each missing ack, wait longer between retries

# Lost Ack: Issue 2

What does a lost ack really mean?

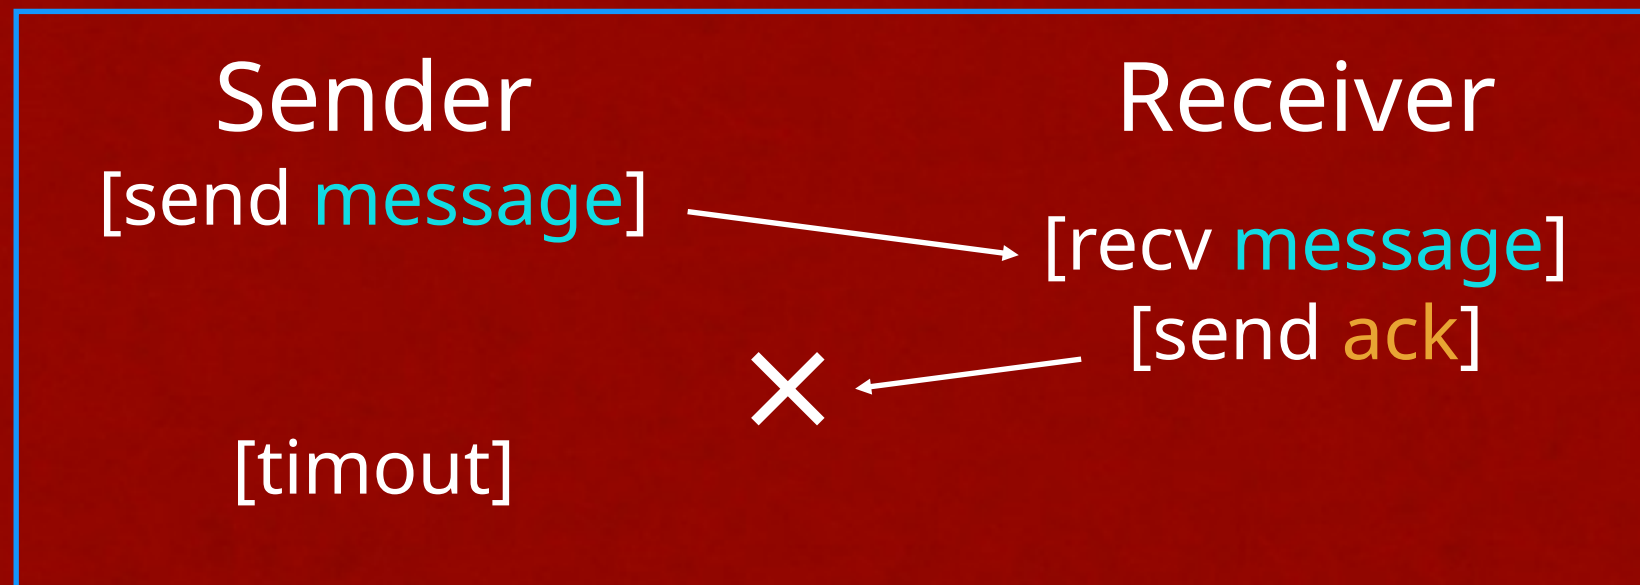


Case 1



Lost ACK:  
How can  
sender  
tell between  
these  
two cases?

Case 2



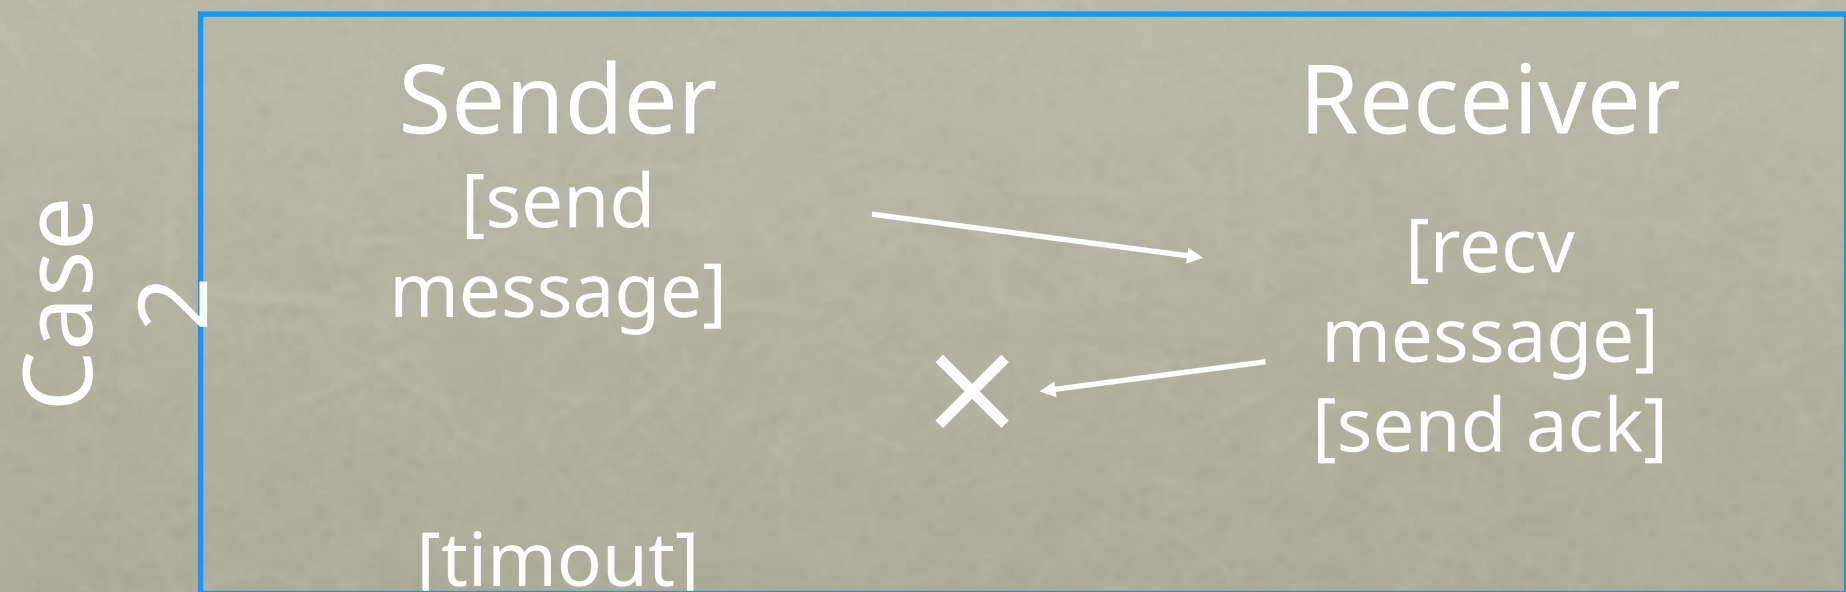
ACK: message received exactly once

No ACK: message may or may not have been received

What if message is command to increment counter?



# Proposed Solution




Proposal:

Sender could send an **AckAck** so receiver knows whether to retry sending an Ack

Sound good?

# Aside: Two Generals' Problem



Suppose generals agree after N messages

Did the arrival of the N'th message change decision?

- if yes: then what if the N'th message had been lost?
- if no: then why bother sending N messages?

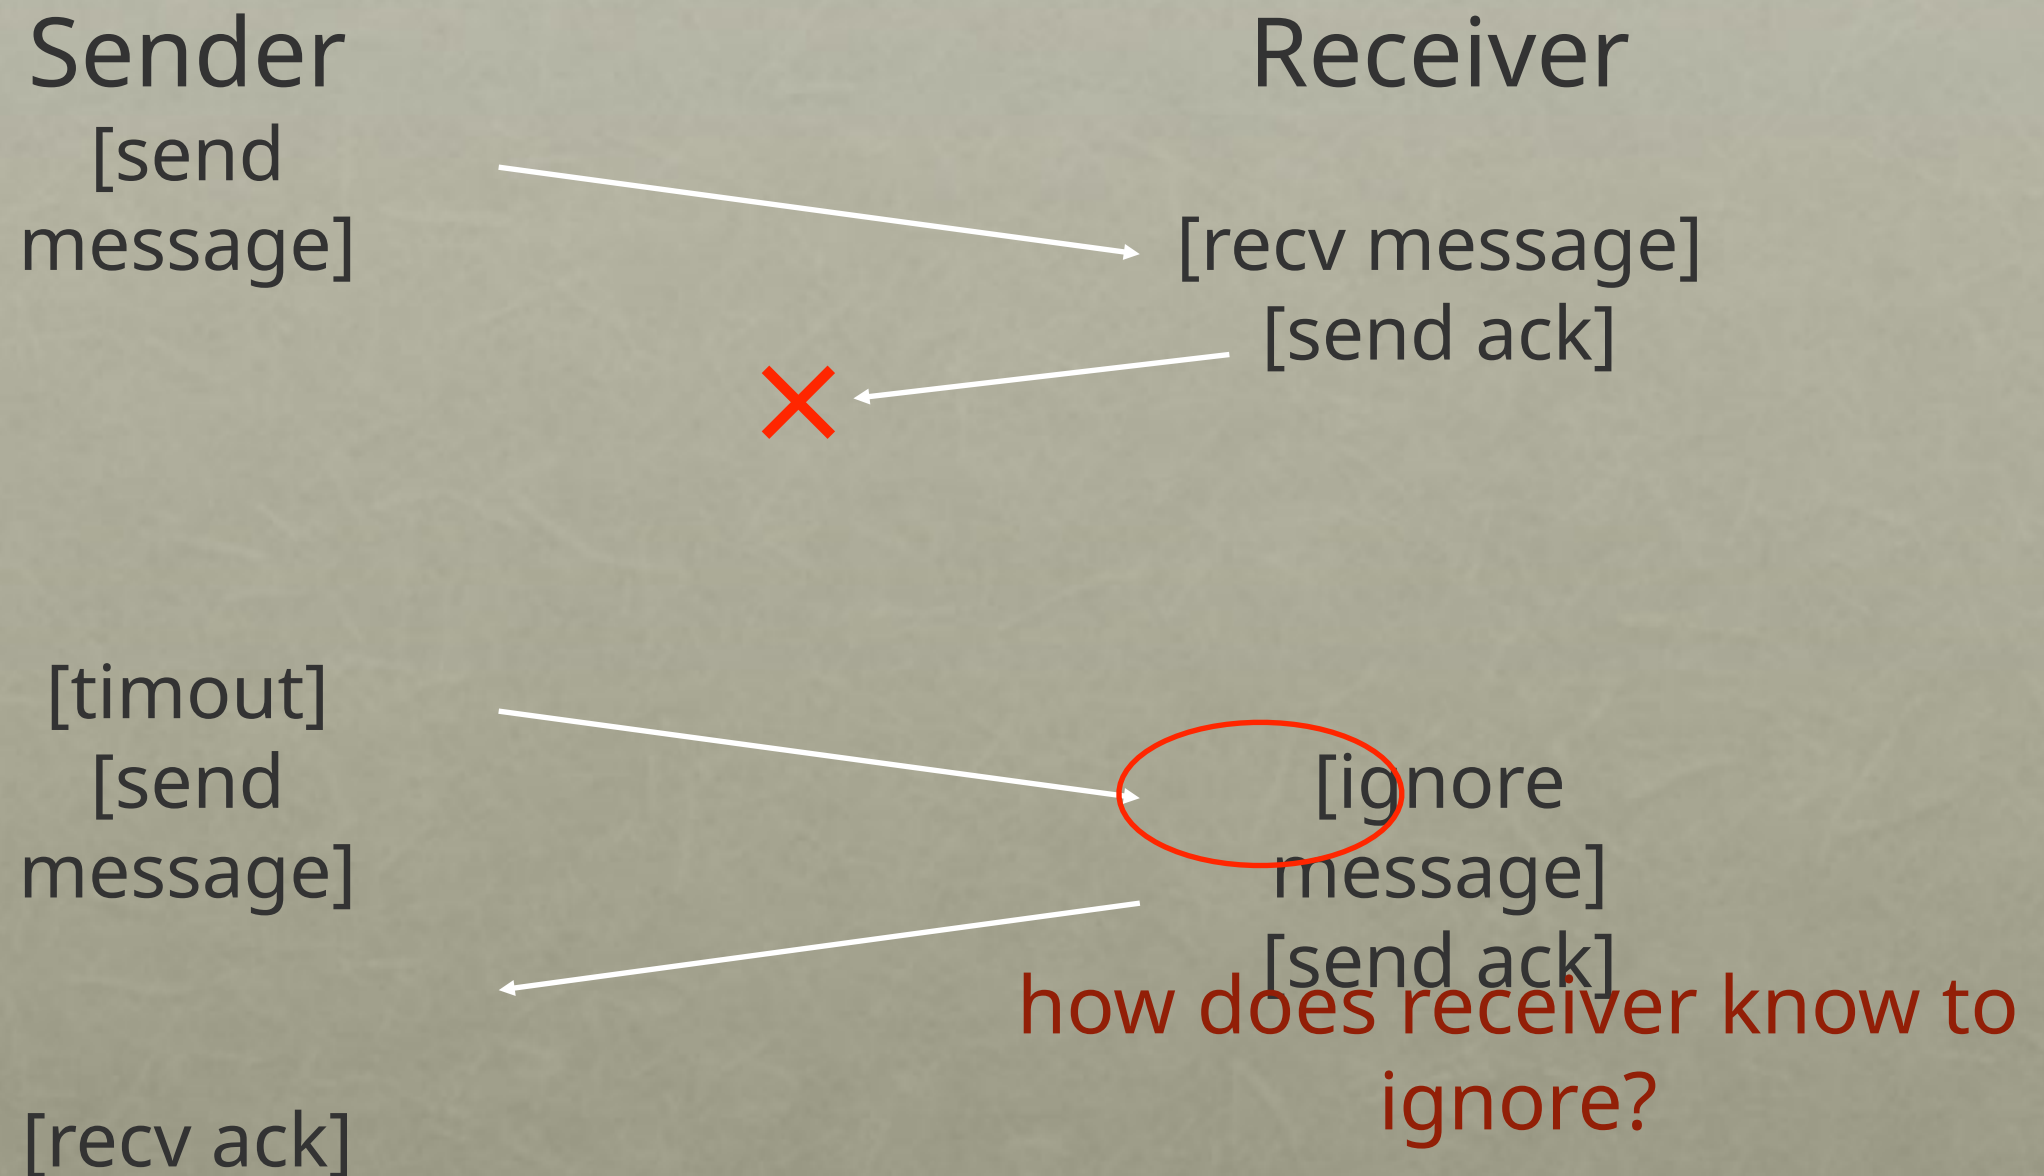
# Reliable Messages: Layering Strategy

Using software, build reliable, logical connections over unreliable connections

Techniques:

- acknowledgment
- timeout
- remember sent messages

# Technique #3: Receiver Remembers Messages





# Solutions

Solution 1: remember every message ever received

Solution 2: sequence numbers

- senders gives each message an increasing unique seq number
- receiver knows it has seen all messages before N
- receiver remembers messages received after N

Suppose message K is received. Suppress message if:

- $K < N$
- Msg K is already buffered



# TCP

TCP: Transmission Control Protocol

Most popular protocol based on seq nums

Buffers messages so arrive in order

Timeouts are adaptive

# Communications Overview

Raw messages: UDP

Reliable messages: TCP

Remote procedure call: RPC

# RPC

## Remote **P**rocedure **C**all

What could be easier than calling a function?

**Strategy:** create wrappers so calling a function on another machine feels just like calling a local function

Very common abstraction

# RPC

Machine

Machine

**A**

```
int main(...) {  
    int x =  
    foo("hello");  
}  
  
int foo(char *msg) {  
    send msg to B  
    recv msg from B  
}
```

**B**

```
int foo(char *msg) {  
    ...  
}  
  
void foo_listener() {  
    while(1) {  
        recv, call  
        foo  
    }  
}
```

What it feels like for  
programmer

# RPC

Machine

A

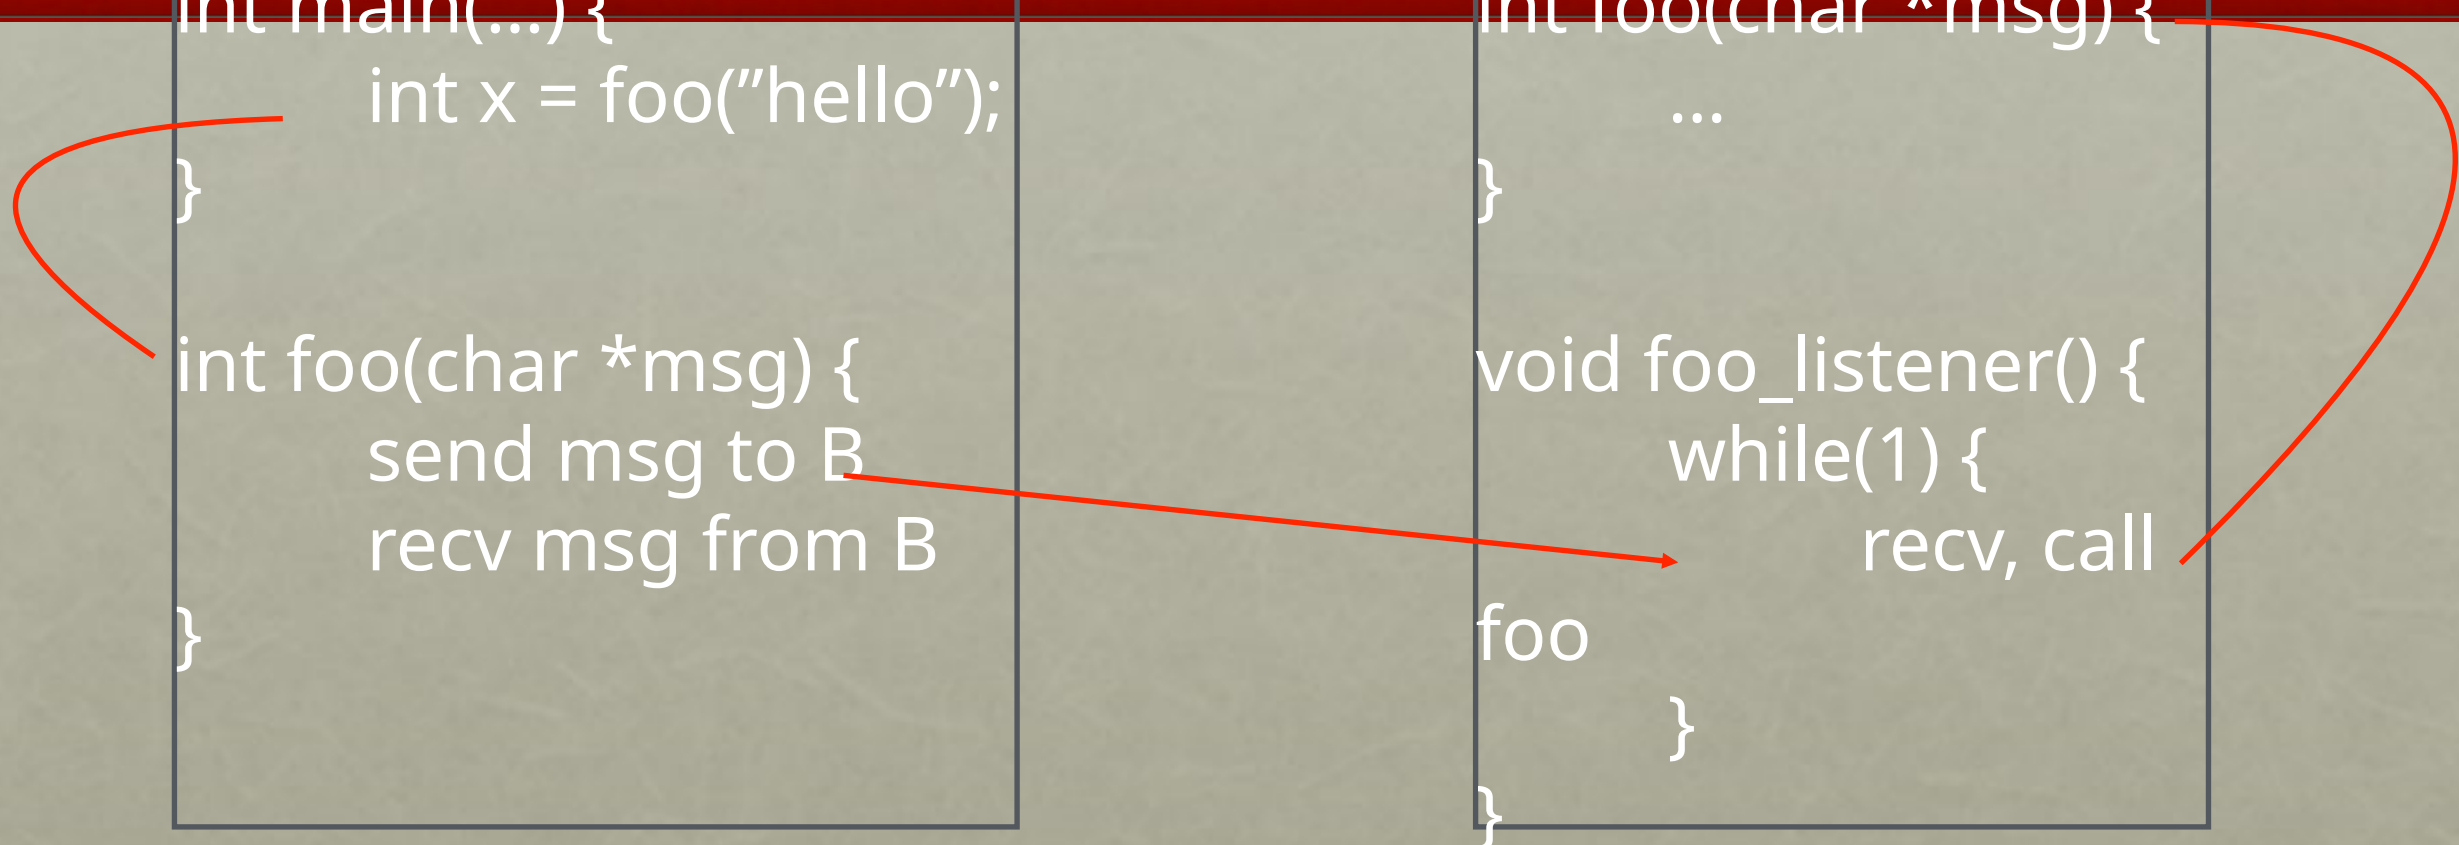
```
int main(...) {  
    int x = foo("hello");  
}  
  
int foo(char *msg) {  
    send msg to B  
    recv msg from B  
}
```

Machine

B

```
int foo(char *msg) {  
    ...  
}  
  
void foo_listener() {  
    while(1) {  
        recv, call  
        foo  
    }  
}
```

Actual  
calls





# RPC

Machine

A

```
int main(...) {  
    int x =  
    foo("hello");  
}
```

client  
wrapper

```
int foo(char *msg) {  
    send msg to B  
    recv msg from B  
}
```

Machine

B

```
int foo(char *msg) {  
    ...  
}
```

server  
wrapper

```
void foo_listener() {  
    while(1) {  
        recv, call  
        foo  
    }  
}
```

Wrapper  
s

# RPC Tools

RPC packages help with two components

## (1) Runtime library

- Thread pool
- Socket listeners call functions on server

## (2) Stub generation

- Create wrappers automatically
- Many tools available (rpcgen, thrift, protobufs)

# Wrapper Generation

Wrappers must do conversions:

- client arguments to message
- message to server arguments
- convert server return value to message
- convert message to client return value

Need uniform endianness (wrappers do this)

Conversion is called marshaling/unmarshaling, or serializing/deserializing

# Wrapper Generation: Pointers

Why are pointers problematic?

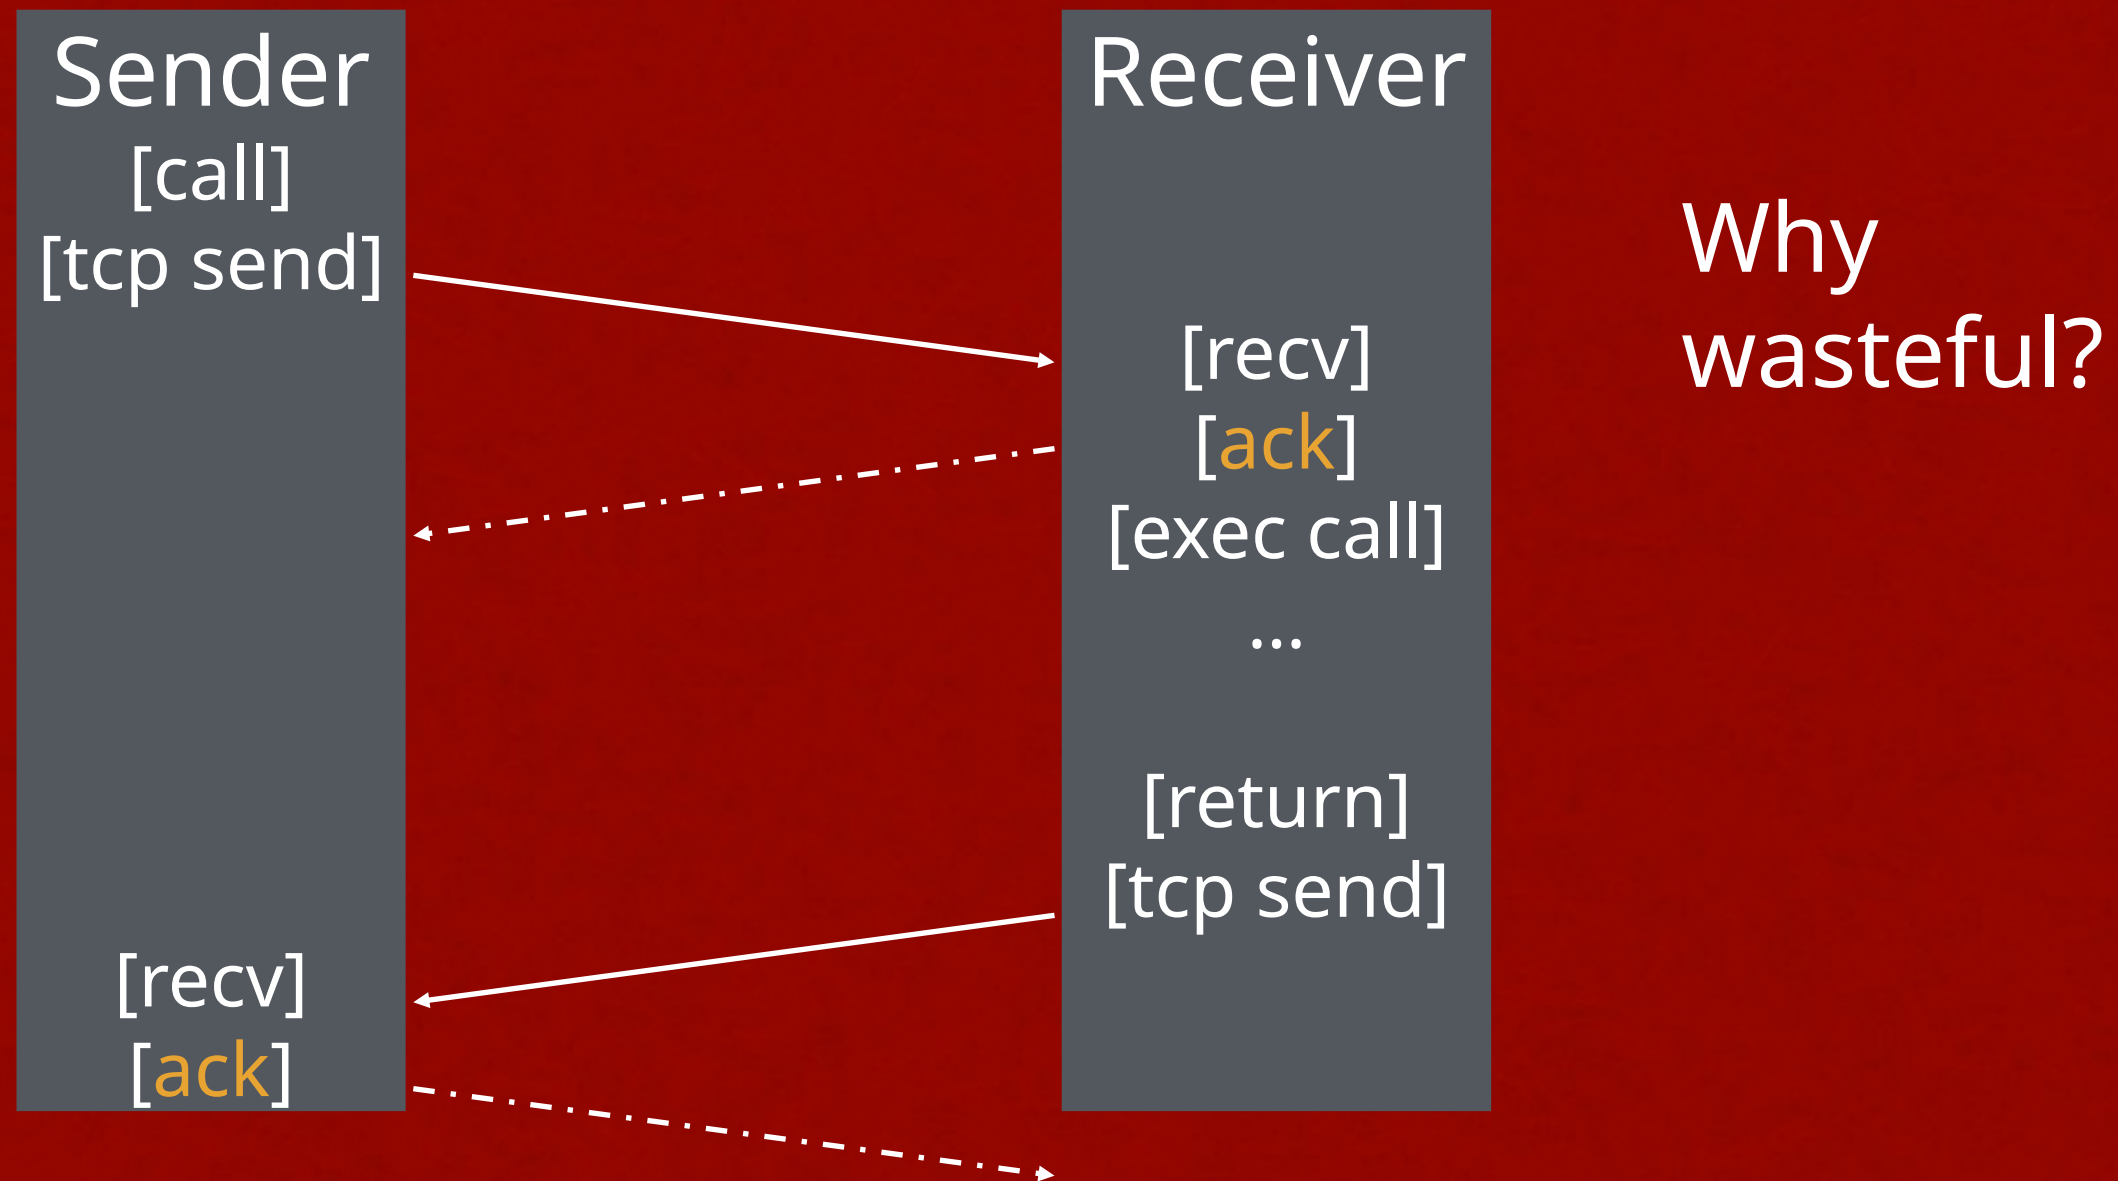
Address passed from client not valid on server

Solutions?

- smart RPC package: follow pointers and copy data



# RPC over TCP?





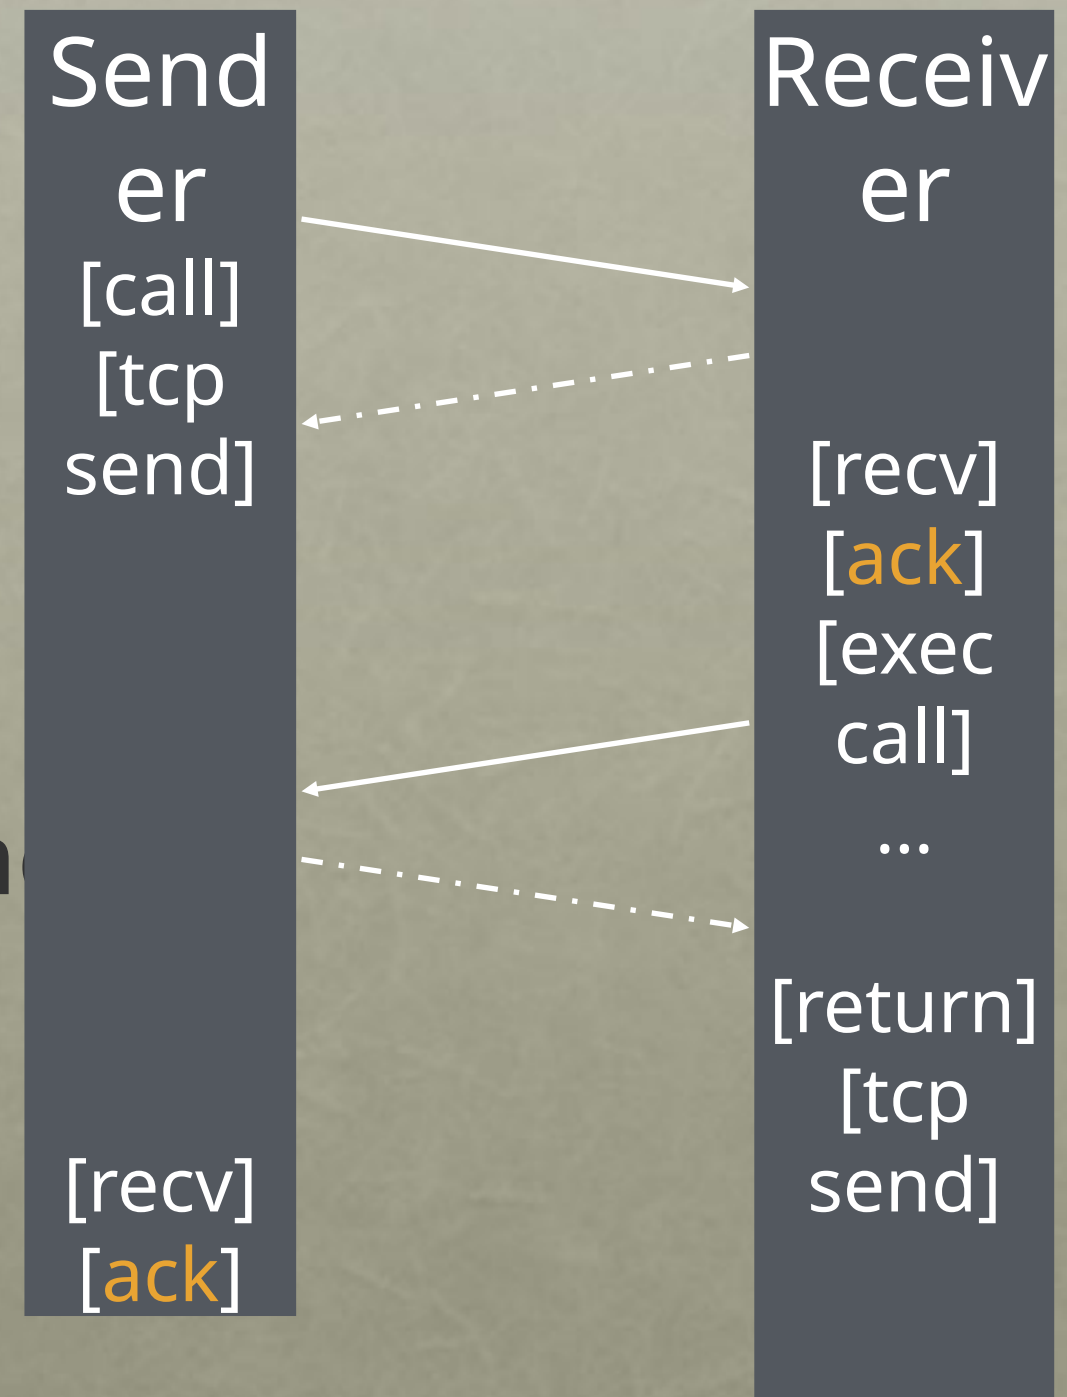
# RPC over UDP

Strategy: use function return as implicit ACK

Piggybacking technique

What if function takes a long time

- then send a separate ACK



# Distributed File Systems

**File systems are great use case for distributed systems**

## **Local FS:**

processes on same machine access shared files

## **Network FS:**

processes on different machines access shared files in same way

# Goals for distributed file systems

Fast + simple crash recovery

- both clients and file server may crash

Transparent access

- can't tell accesses are over the network
- normal UNIX semantics

Reasonable performance

# NFS

Think of NFS as more of a protocol than a particular file system

Many companies have implemented NFS:  
Oracle/Sun, NetApp, EMC, IBM

We're looking at NFSv2

- NFSv4 has many changes

Why look at an older protocol?

- Simpler, focused goals
- To compare and contrast NFS with AFS (next lecture)



# Overview

Architecture

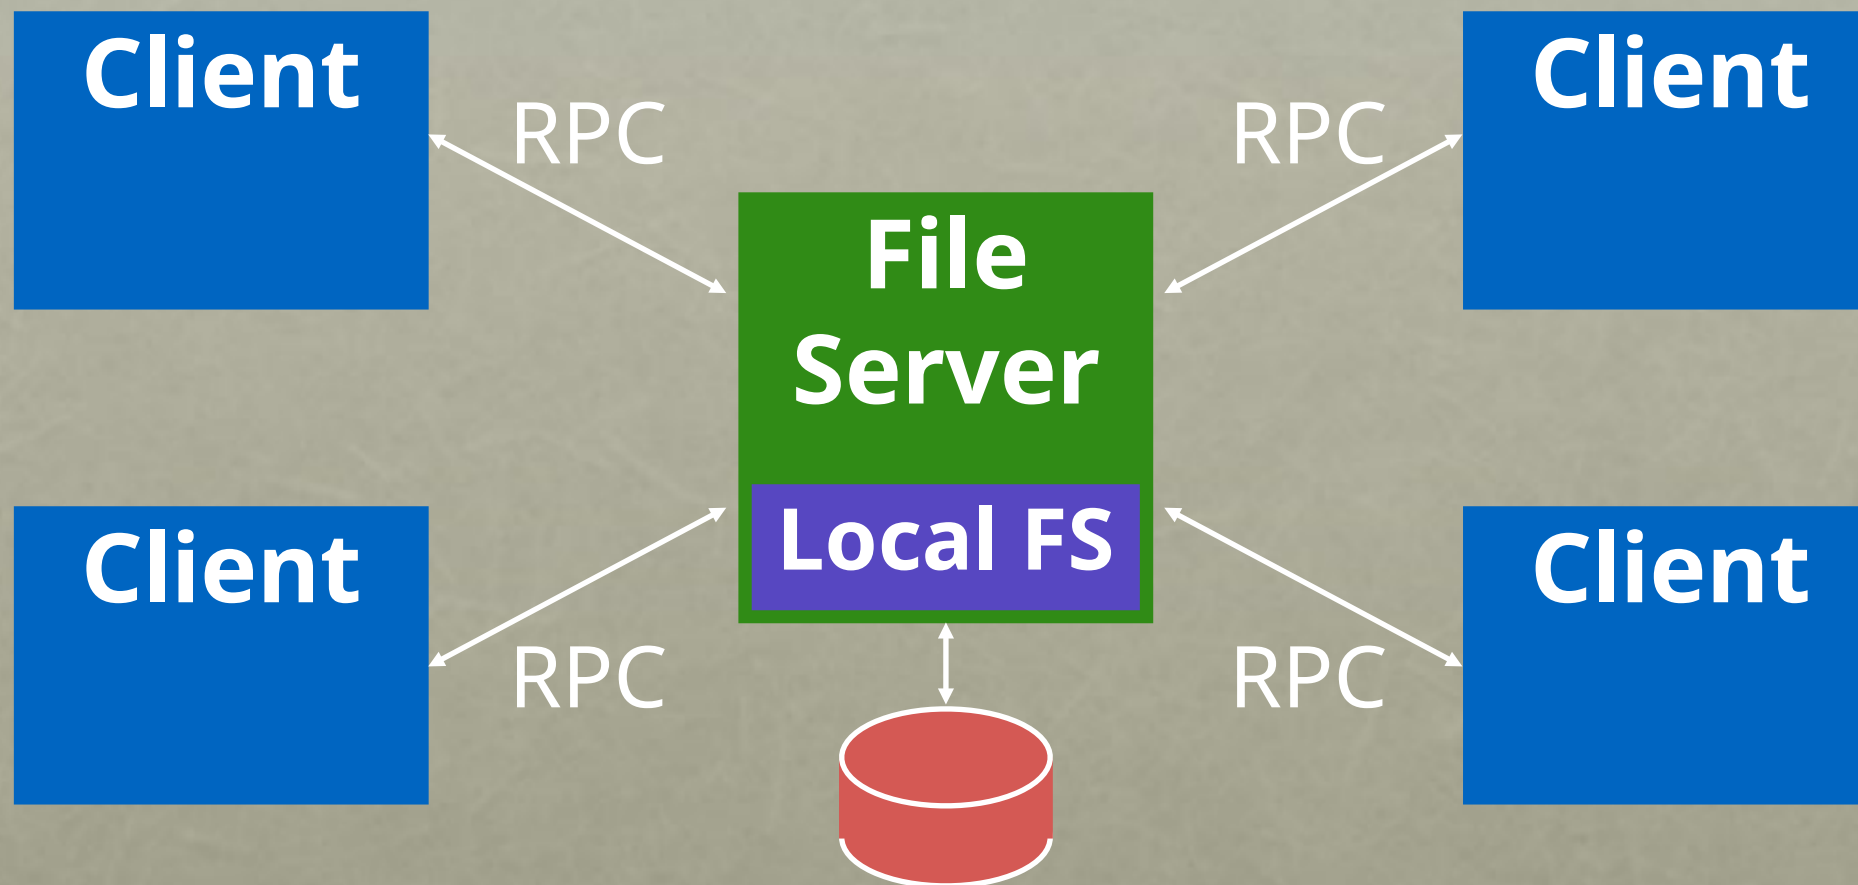
Network API

Write Buffering

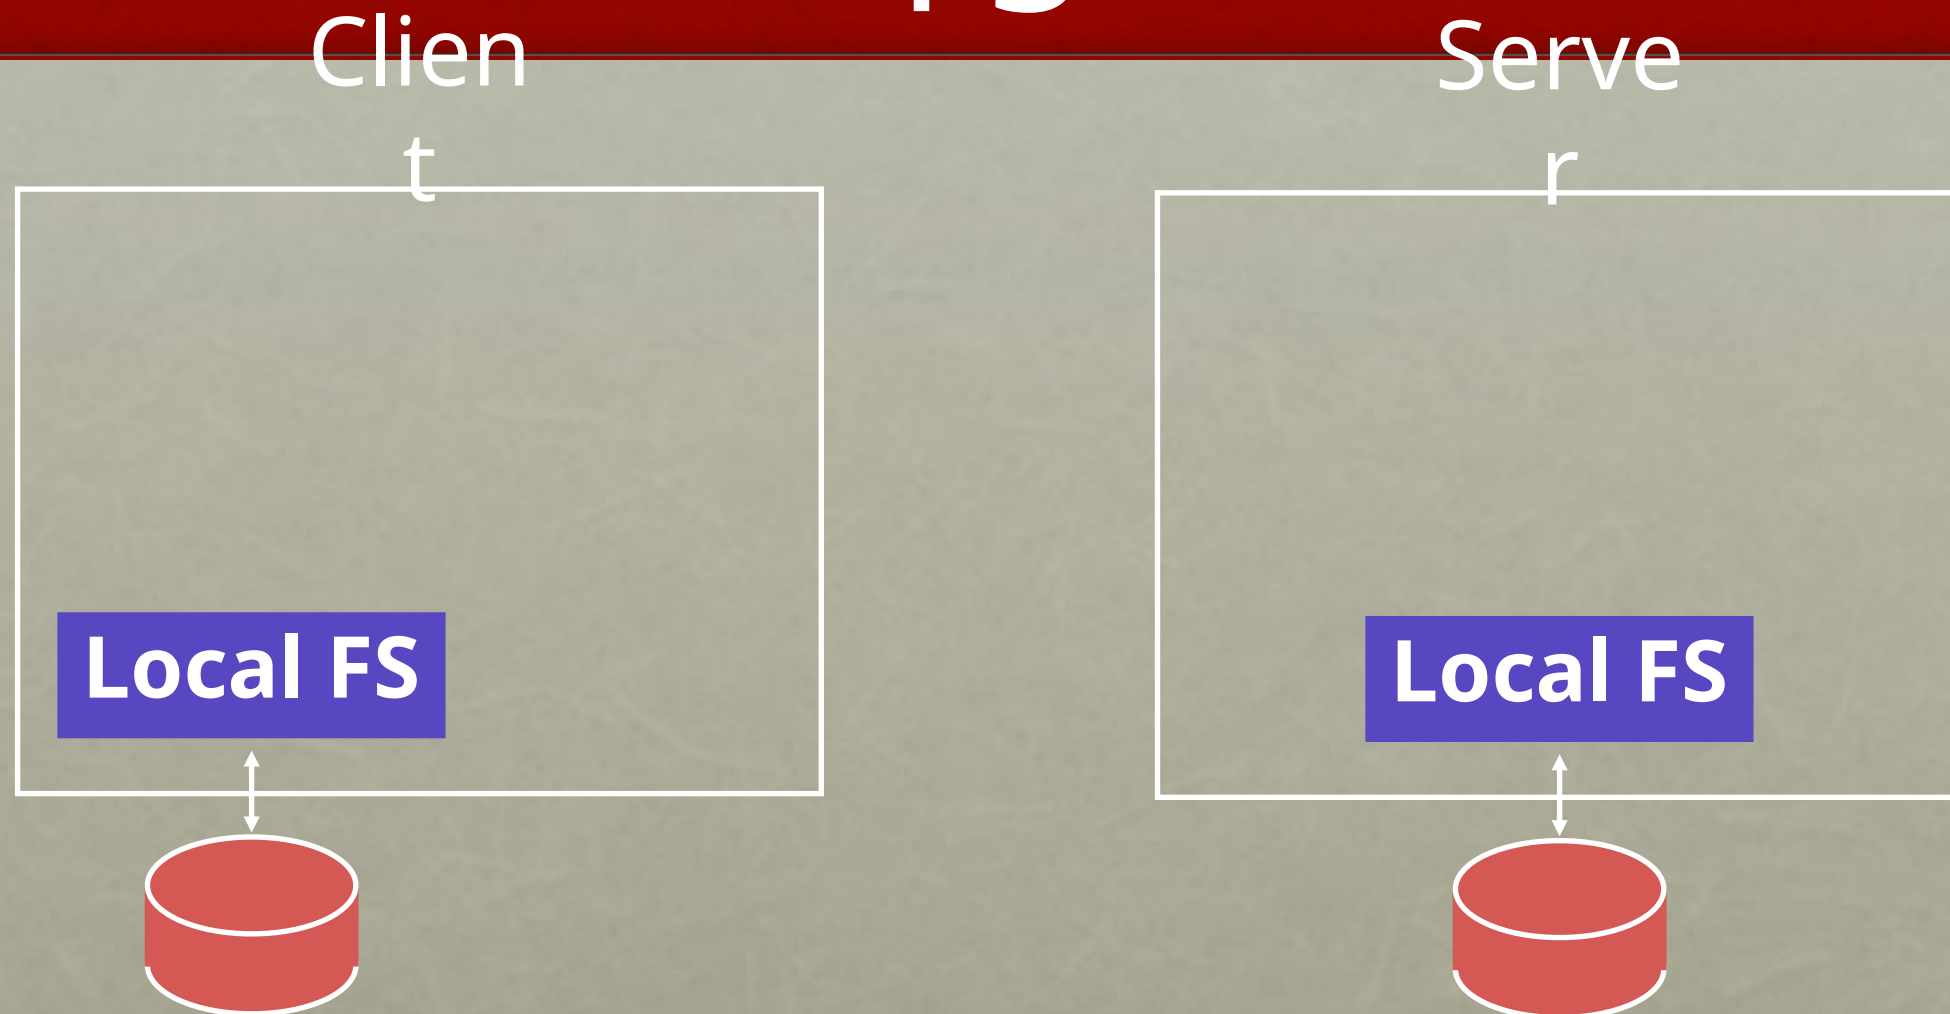
Cache




# NFS Architecture



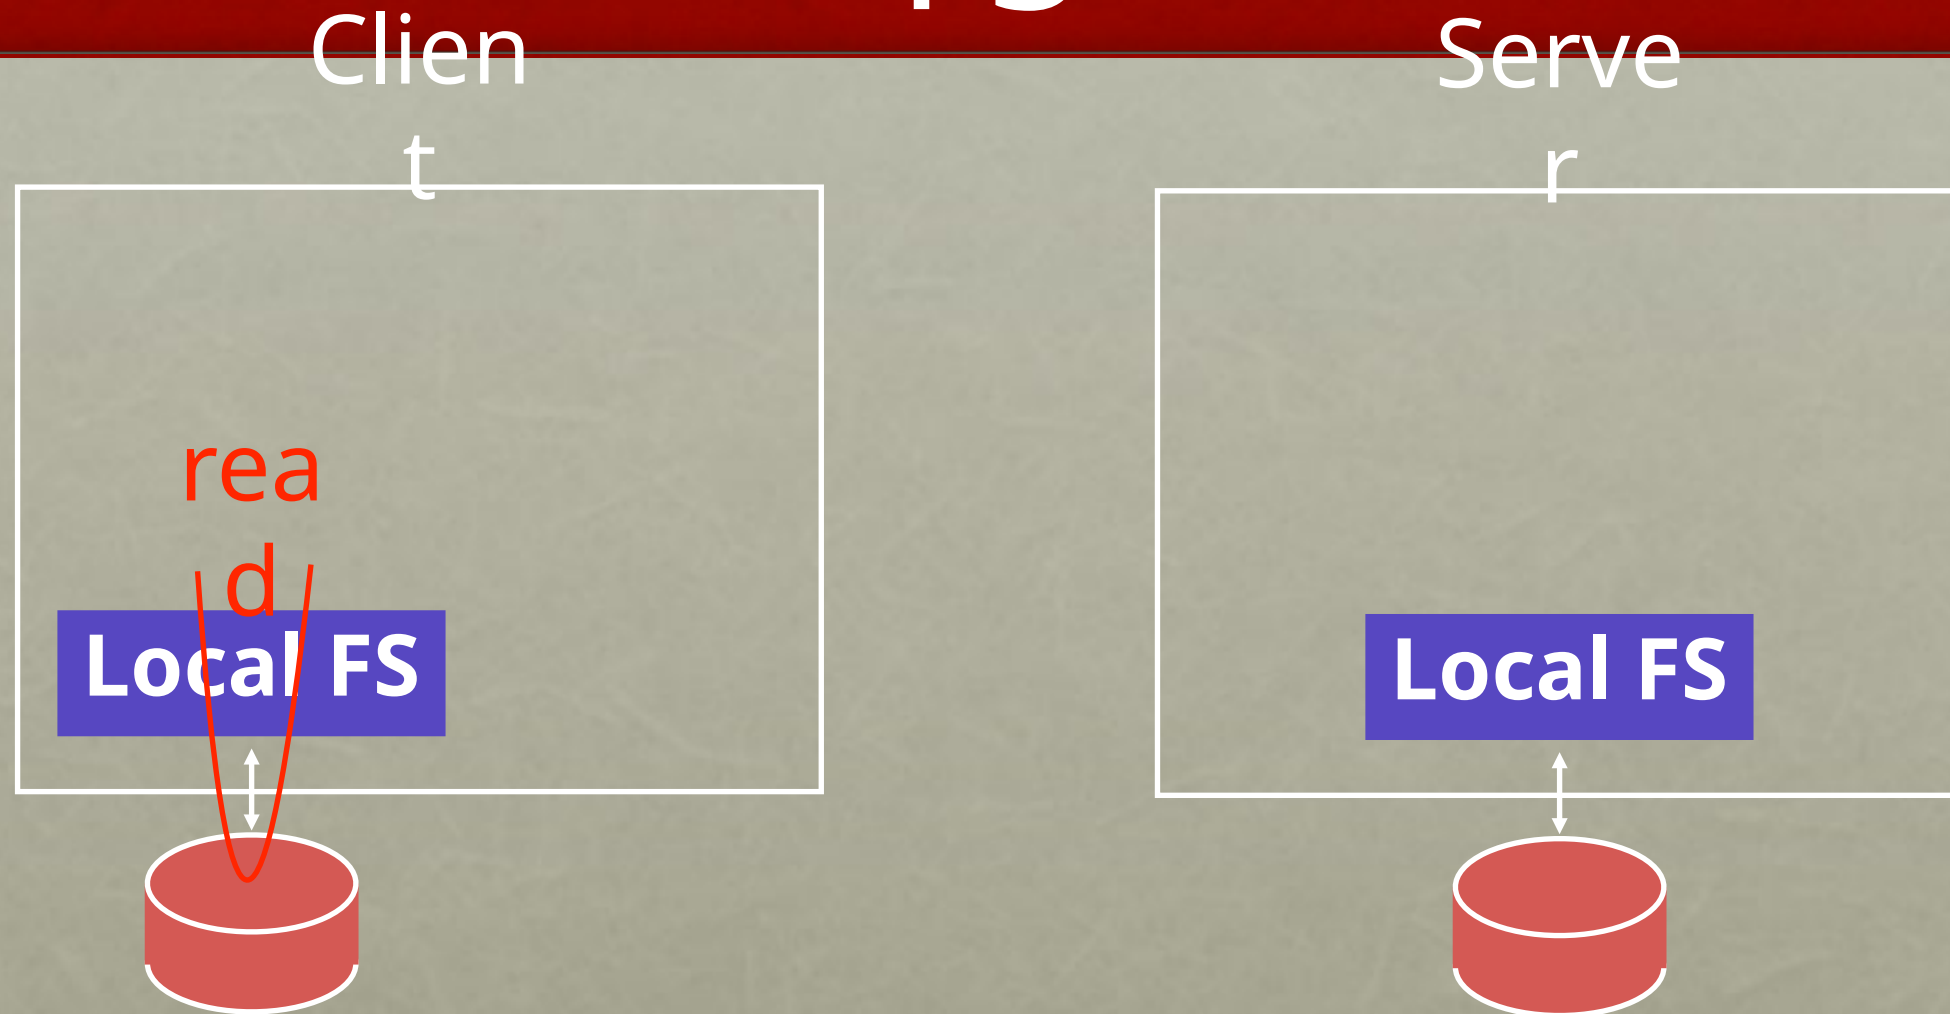
# General Strategy: Export FS




# General Strategy: Export FS



# General Strategy: Export FS




# General Strategy: Export FS





# General Strategy: Export FS






- /dev/sda1 **on** /
- /dev/sdb1 **on** /backups
- AFS **on** /home/tyler

# General Strategy: Export FS



# General Strategy: Export FS





# Goals for NFS

Fast + simple crash recovery

- both clients and file server may crash

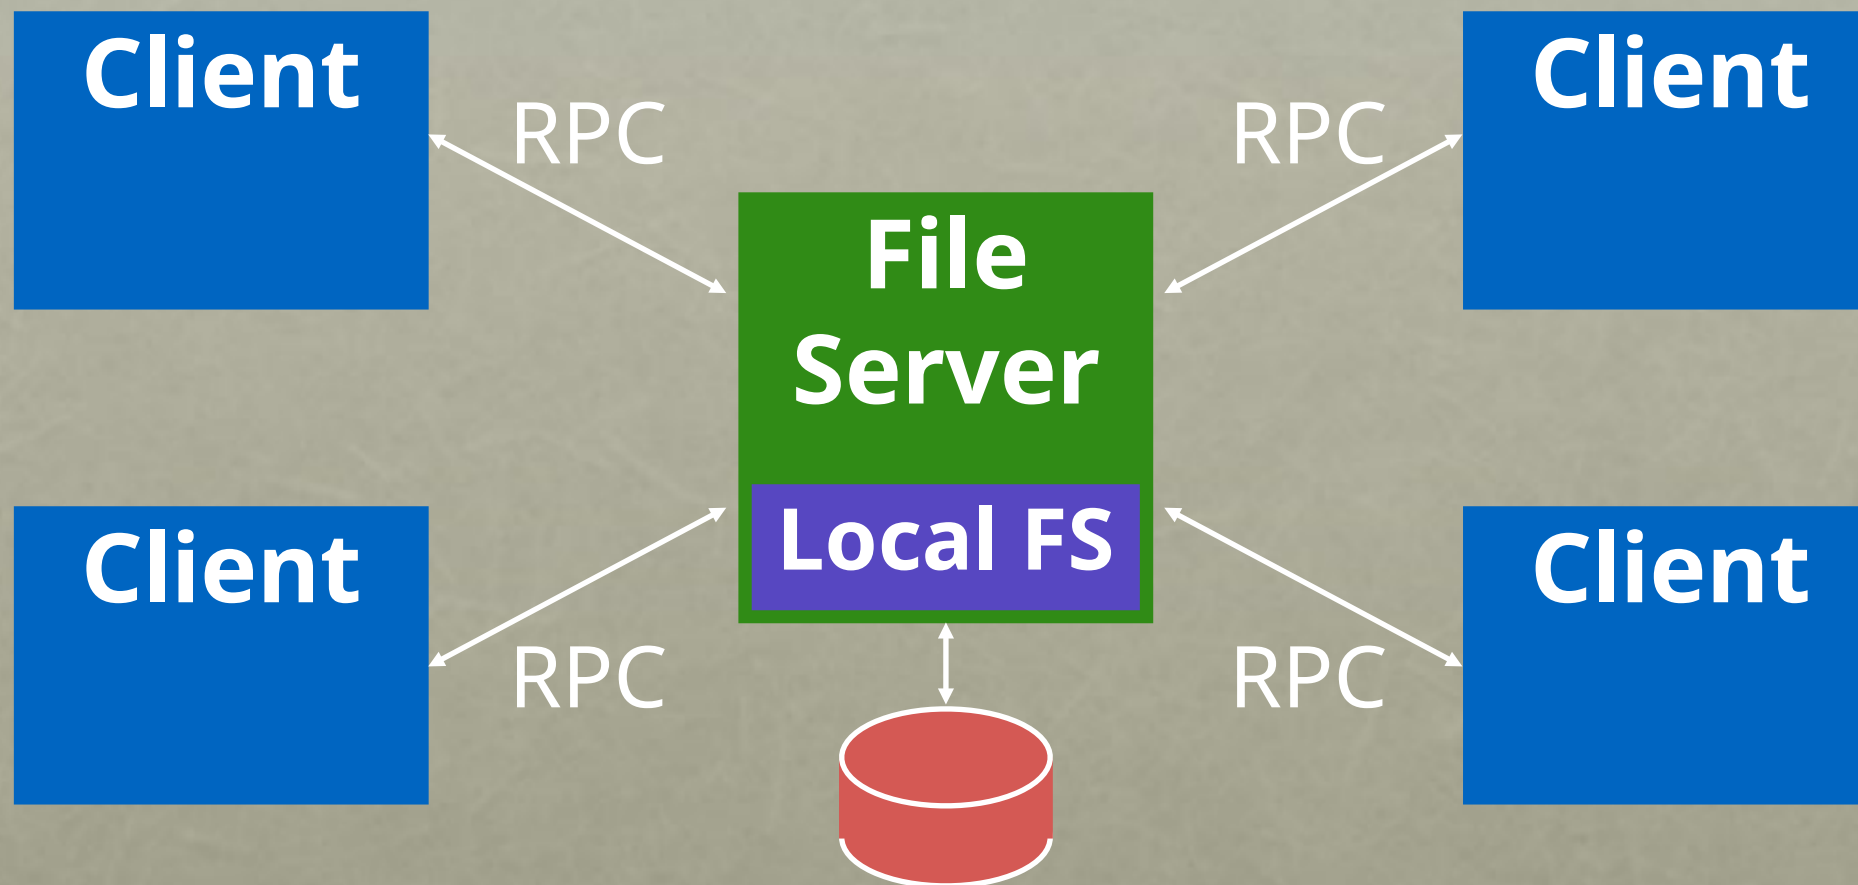
Transparent access

- can't tell accesses are over the network
- normal UNIX semantics

Reasonable performance



# NFS Architecture



# Overview

~~Architecture~~

Network API

Write Buffering

Cache

# Strategy 1

Attempt: Wrap regular UNIX system calls using RPC

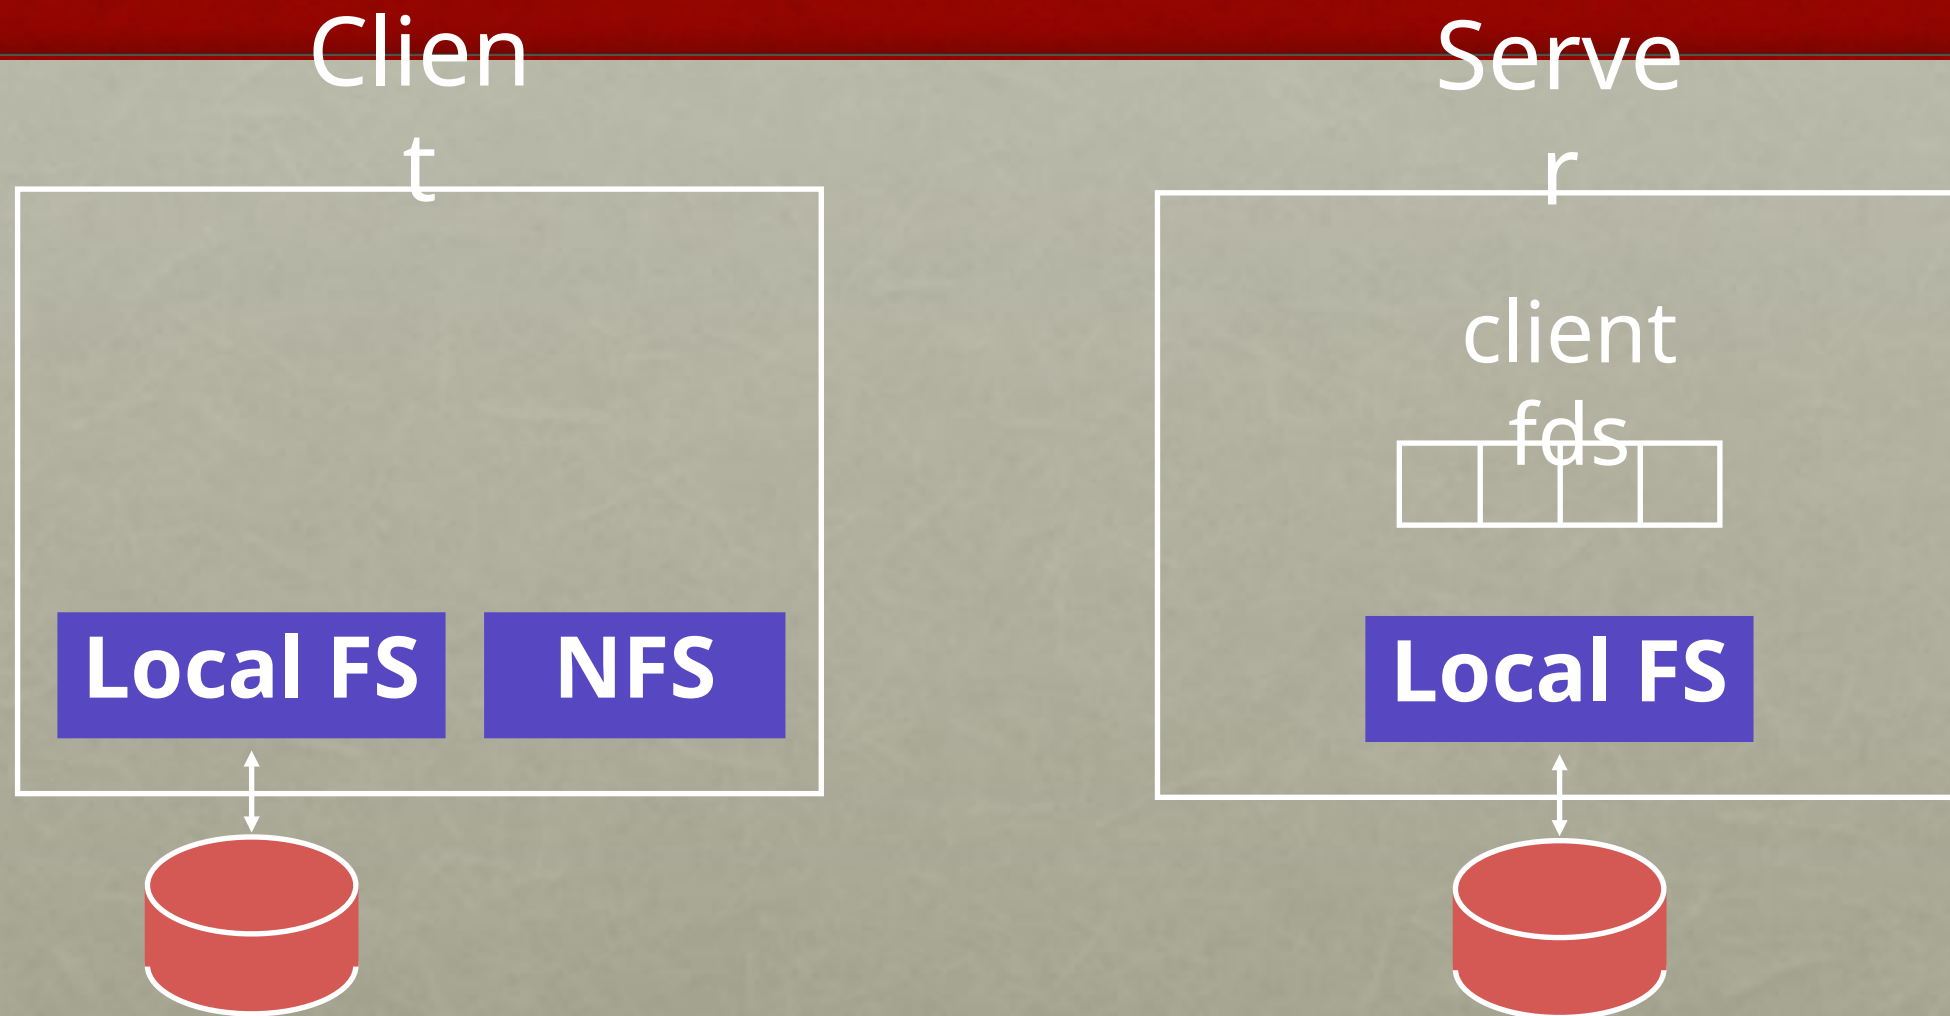
`open()` on client calls `open()` on server

`open()` on server returns fd back to client

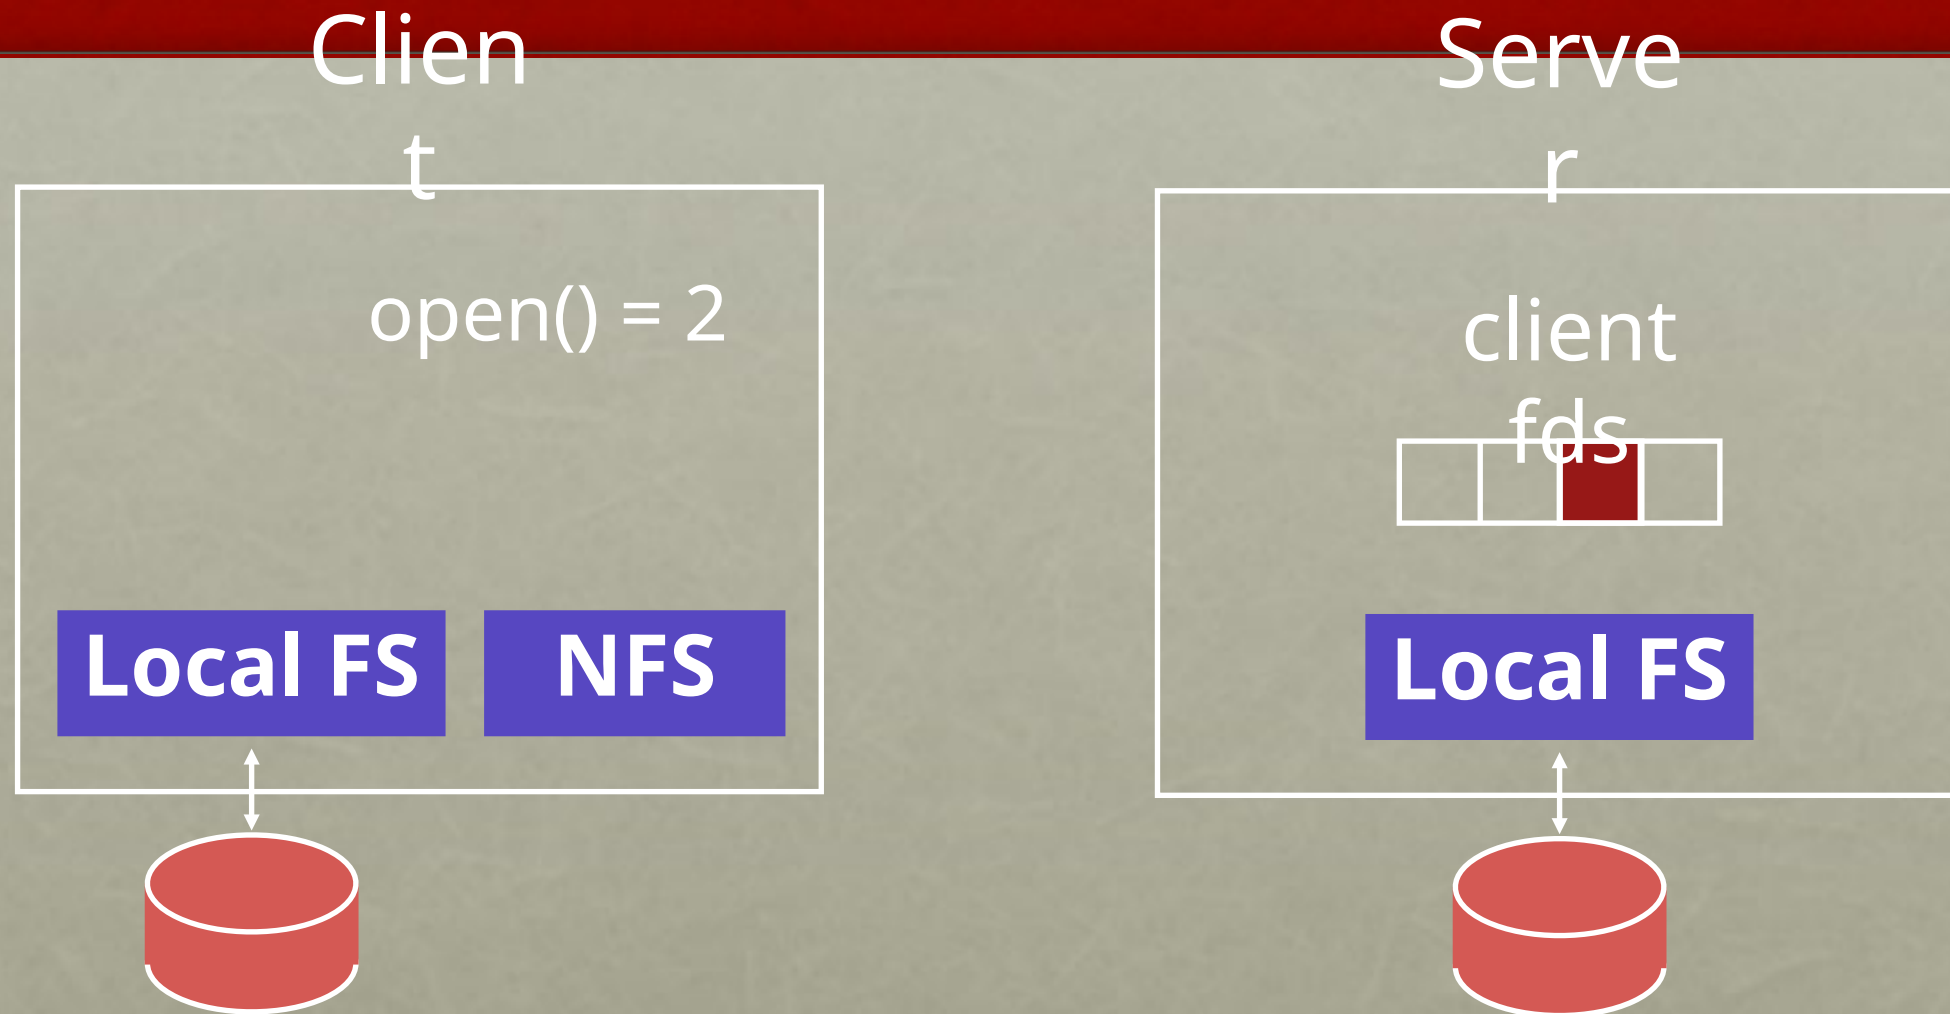
`read(fd)` on client calls `read(fd)` on server

`read(fd)` on server returns data back to client

# File Descriptors

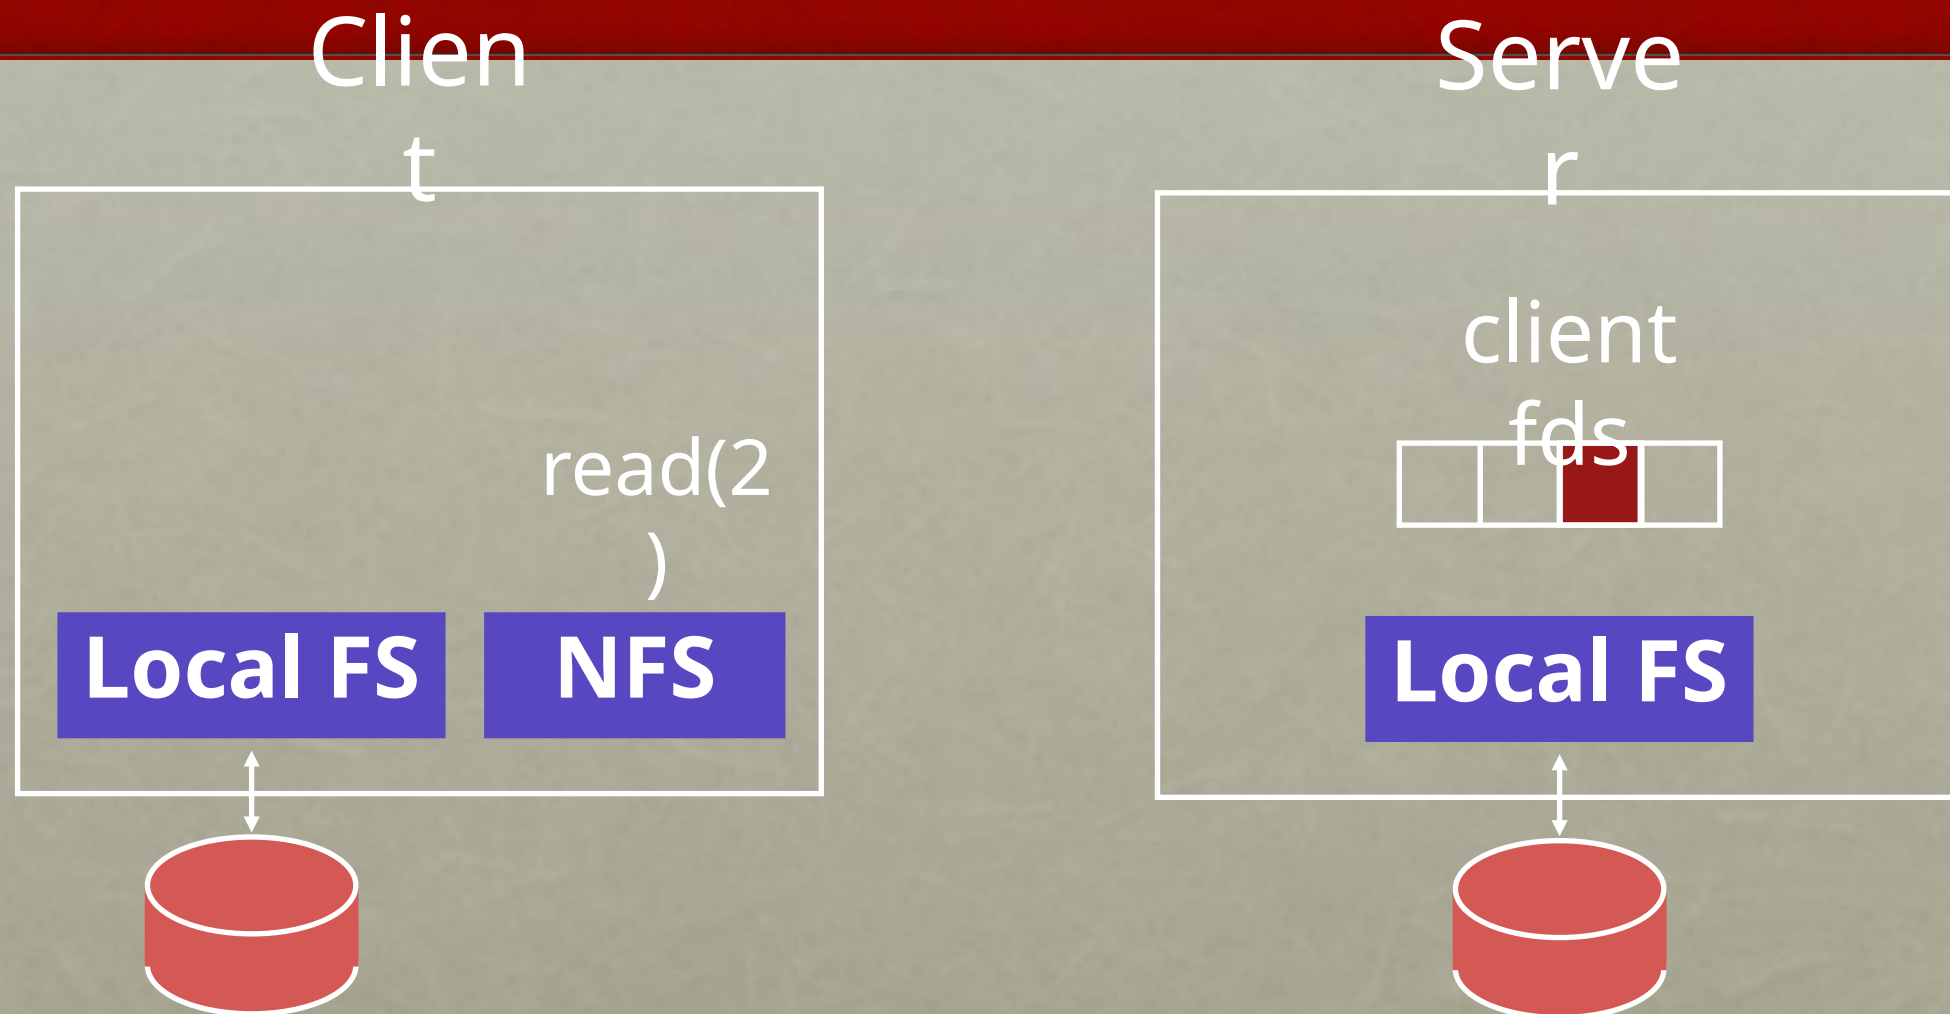


# File Descriptors

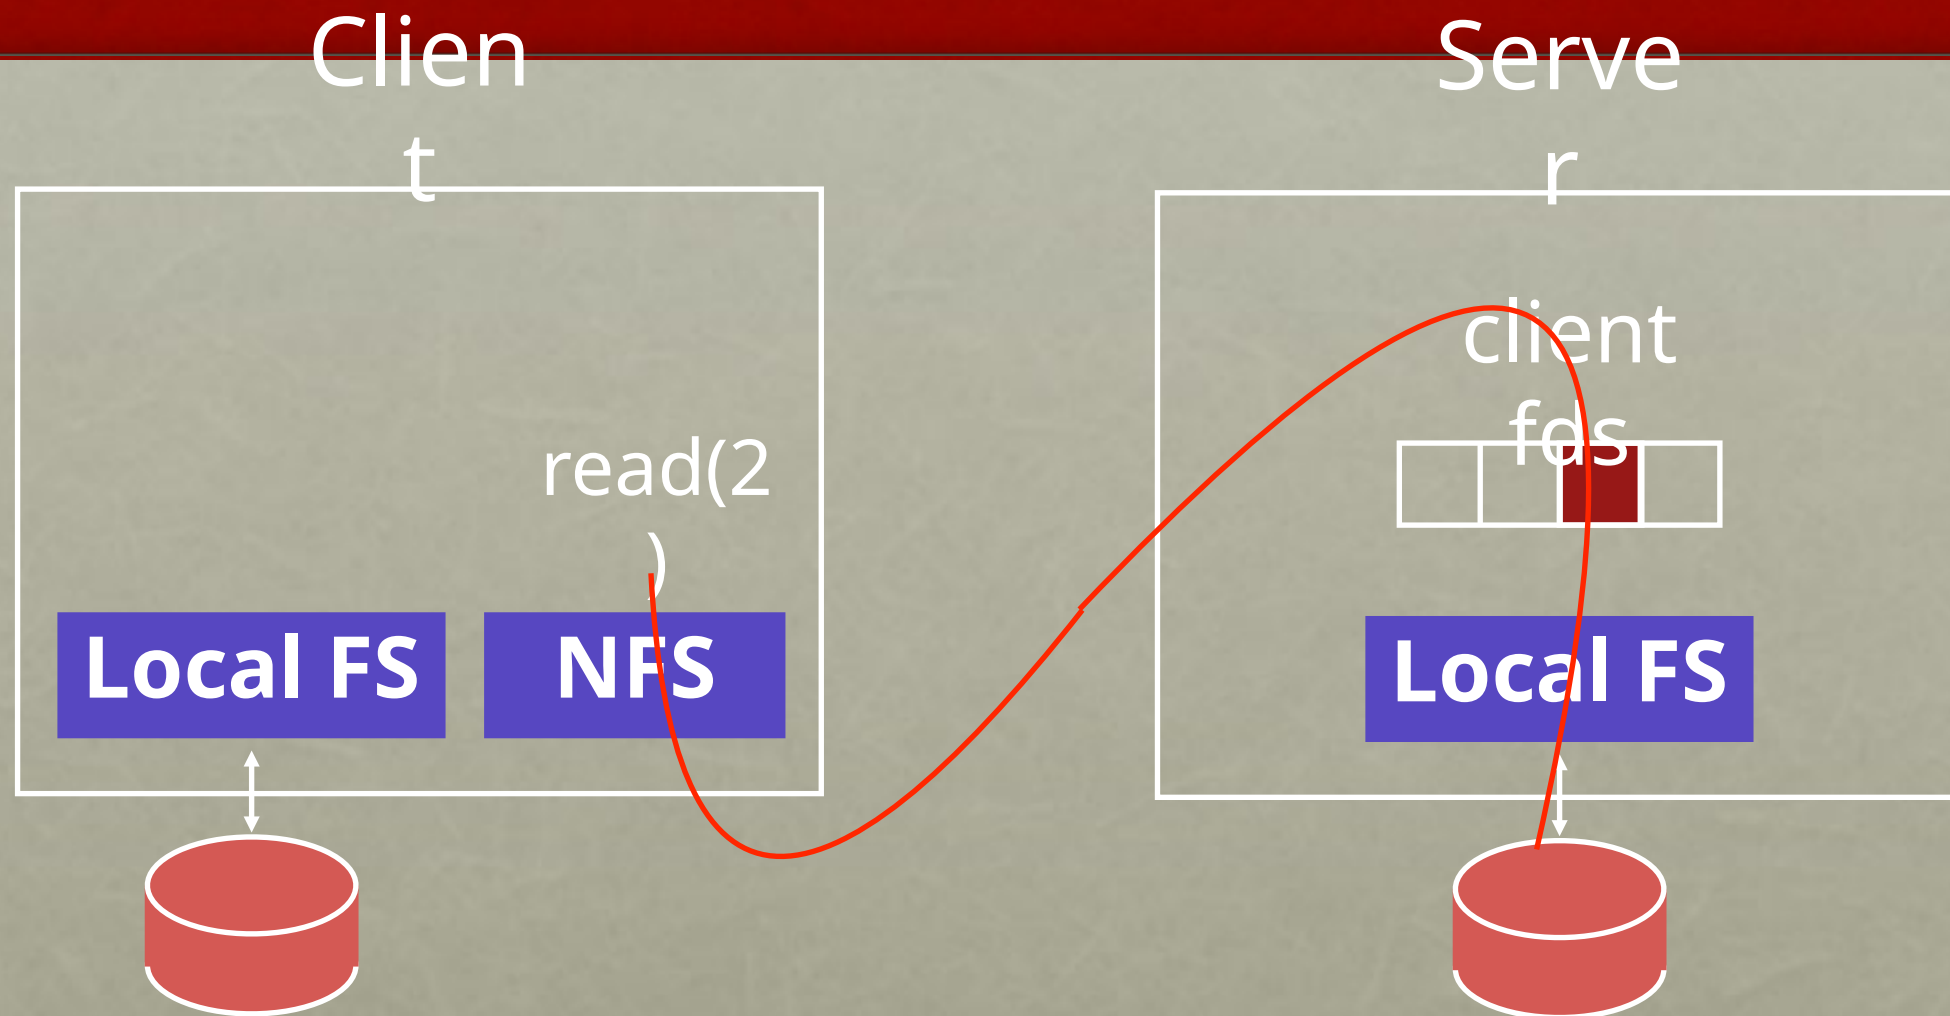




# File Descriptors



# File Descriptors



# Strategy 1 Problems

What about crashes?


```
int fd = open("foo", O_RDONLY);
```

```
read(fd, buf, MAX);
```

```
read(fd, buf, MAX);
```

```
...
```

```
read(fd, buf, MAX);
```



Server  
crash!

nice if acts like a slow read

Imagine server crashes and reboots during reads...

# Potential Solutions

1. Run some crash recovery protocol upon reboot
  - Complex
2. Persist fds on server disk.
  - Slow
  - What if client crashes? When can fds be garbage collected?


# Strategy 2: put all info in requests

Use “stateless” protocol!

- server maintains no state about clients
- server still keeps other state, of course



# Eliminate File Descriptors



# Strategy 2: put all info in requests

Use “stateless” protocol!

- server maintains no state about clients

Need API change. One possibility:

**pread**(char \*path, buf, size, offset);

**pwrite**(char \*path, buf, size, offset);

Specify path and offset each time. Server need not remember anything from clients.

Pros?

Server can crash and reboot transparently to clients.

Cons?

Too many path lookups.

# Strategy 3: inode requests

```
inode = open(char *path);
```

```
pread(inode, buf, size, offset);
```

```
pwrite(inode, buf, size, offset);
```

This is pretty good! Any correctness problems?

If file is deleted, the inode could be reused

- Inode not guaranteed to be unique over time

# Strategy 4: file handles

```
fh = open(char *path);
```

```
pread(fh, buf, size, offset);
```

```
pwrite(fh, buf, size, offset);
```

File Handle = <volume ID, inode #, **generation** #>

Opaque to client (client should not interpret internals)



# Can NFS Protocol include Append?

```
fh = open(char *path);
```

```
pread(fh, buf, size, offset);
```

```
pwrite(fh, buf, size, offset);
```

```
append(fh, buf, size);
```

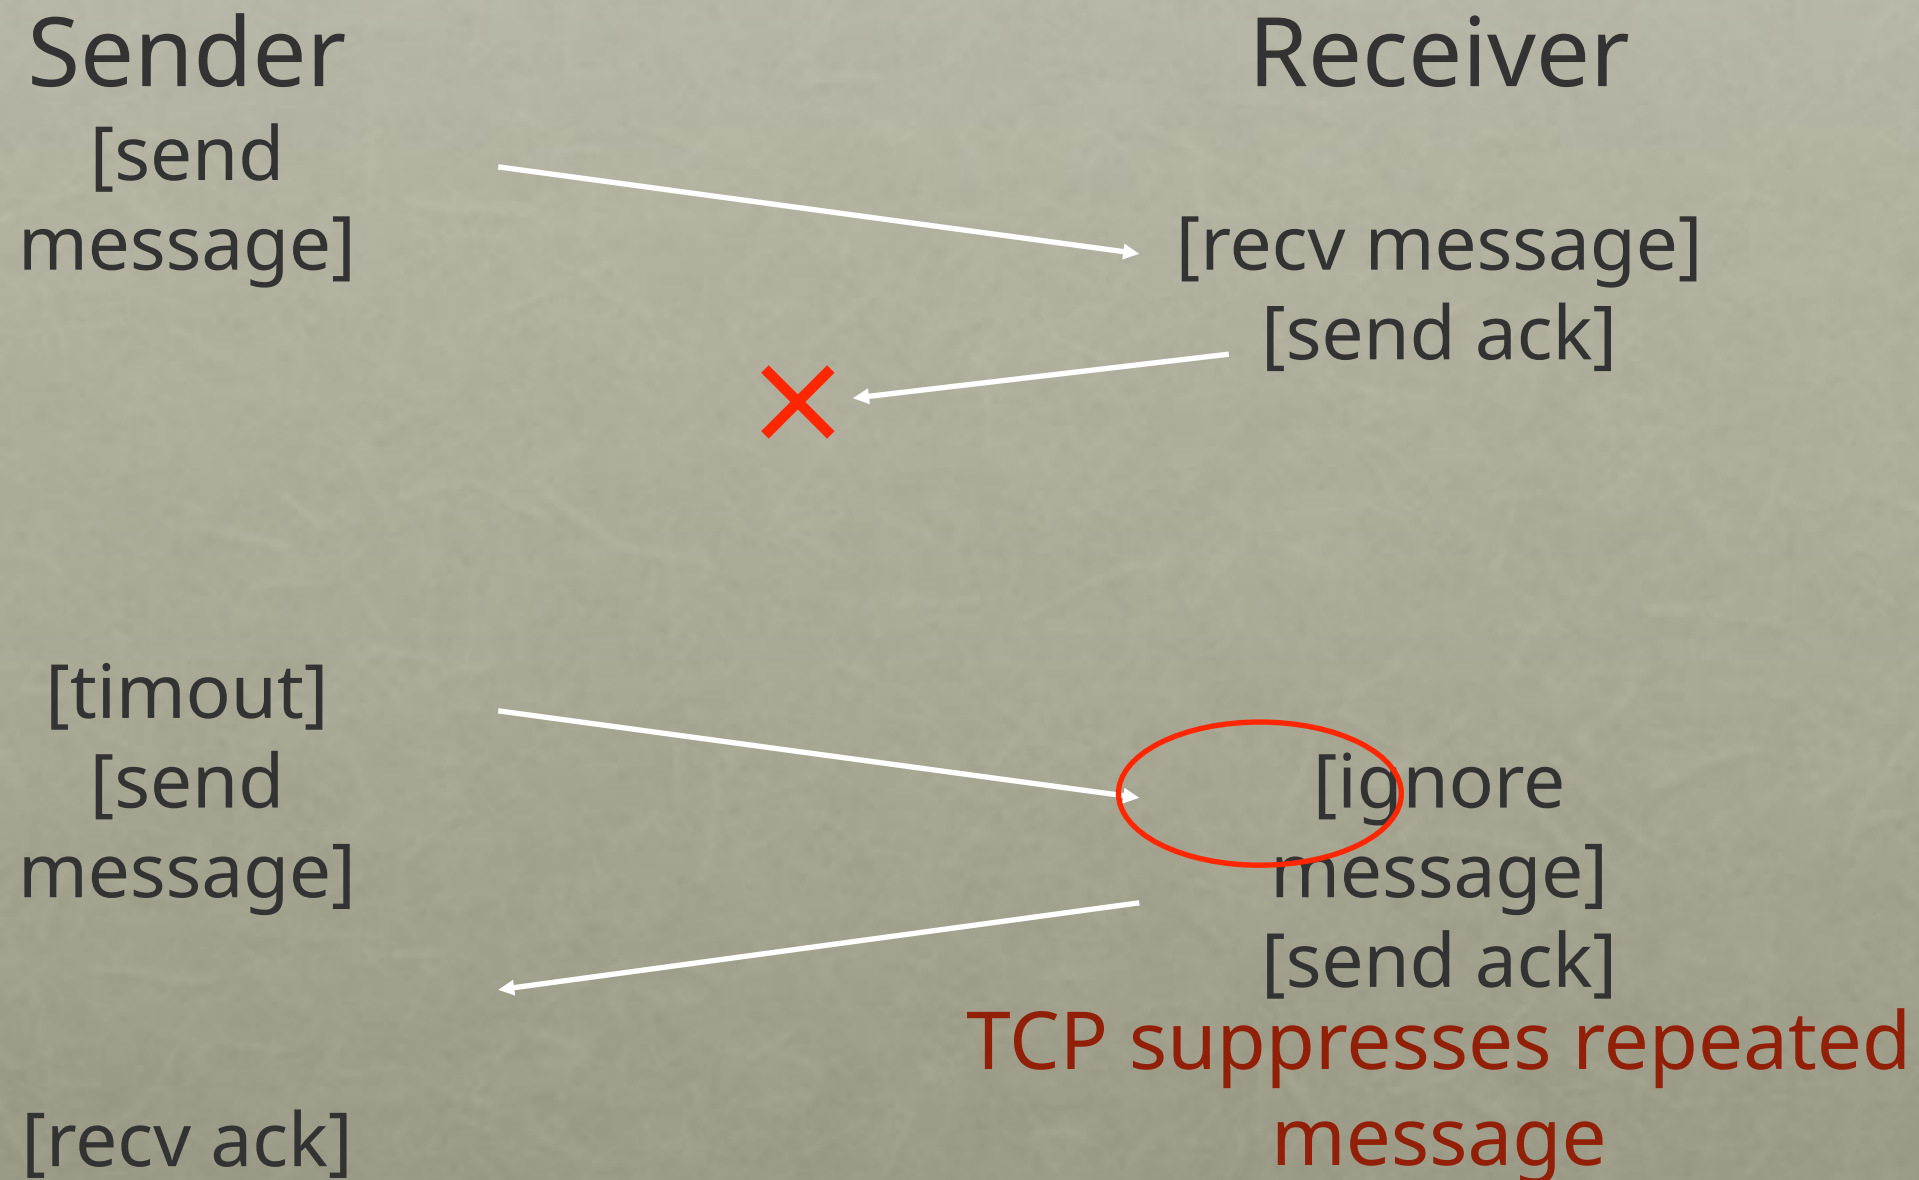
Problem with append()?

If RPC library retries, what happens when append() is retried?

Problem: Why is it difficult to not replay append()?



# Replica Suppression is Stateful



Problem: TCP is stateful

If server crashes, forgets which RPC's have been executed!

# Idempotent Operations


Solution:

Design API so no harm to executing function more than once

If  $f()$  is idempotent, then:

$f()$  has the same effect as  $f(); f(); \dots f(); f()$

# fwrite is idempotent



# append is NOT idempotent



# What operations are Idempotent?

## Idempotent

- any sort of read that doesn't change anything
- pwrite

## Not idempotent

- append

## What about these?

- mkdir
- creat



# Strategy 4: file handles

**fh = open(char \*path);**

**pread(fh, buf, size, offset);**

**pwrite(fh, buf, size, offset);**

**~~append~~(fh, buf, size);**

File Handle = <volume ID, inode #, generation #>

# Strategy 5: client logic

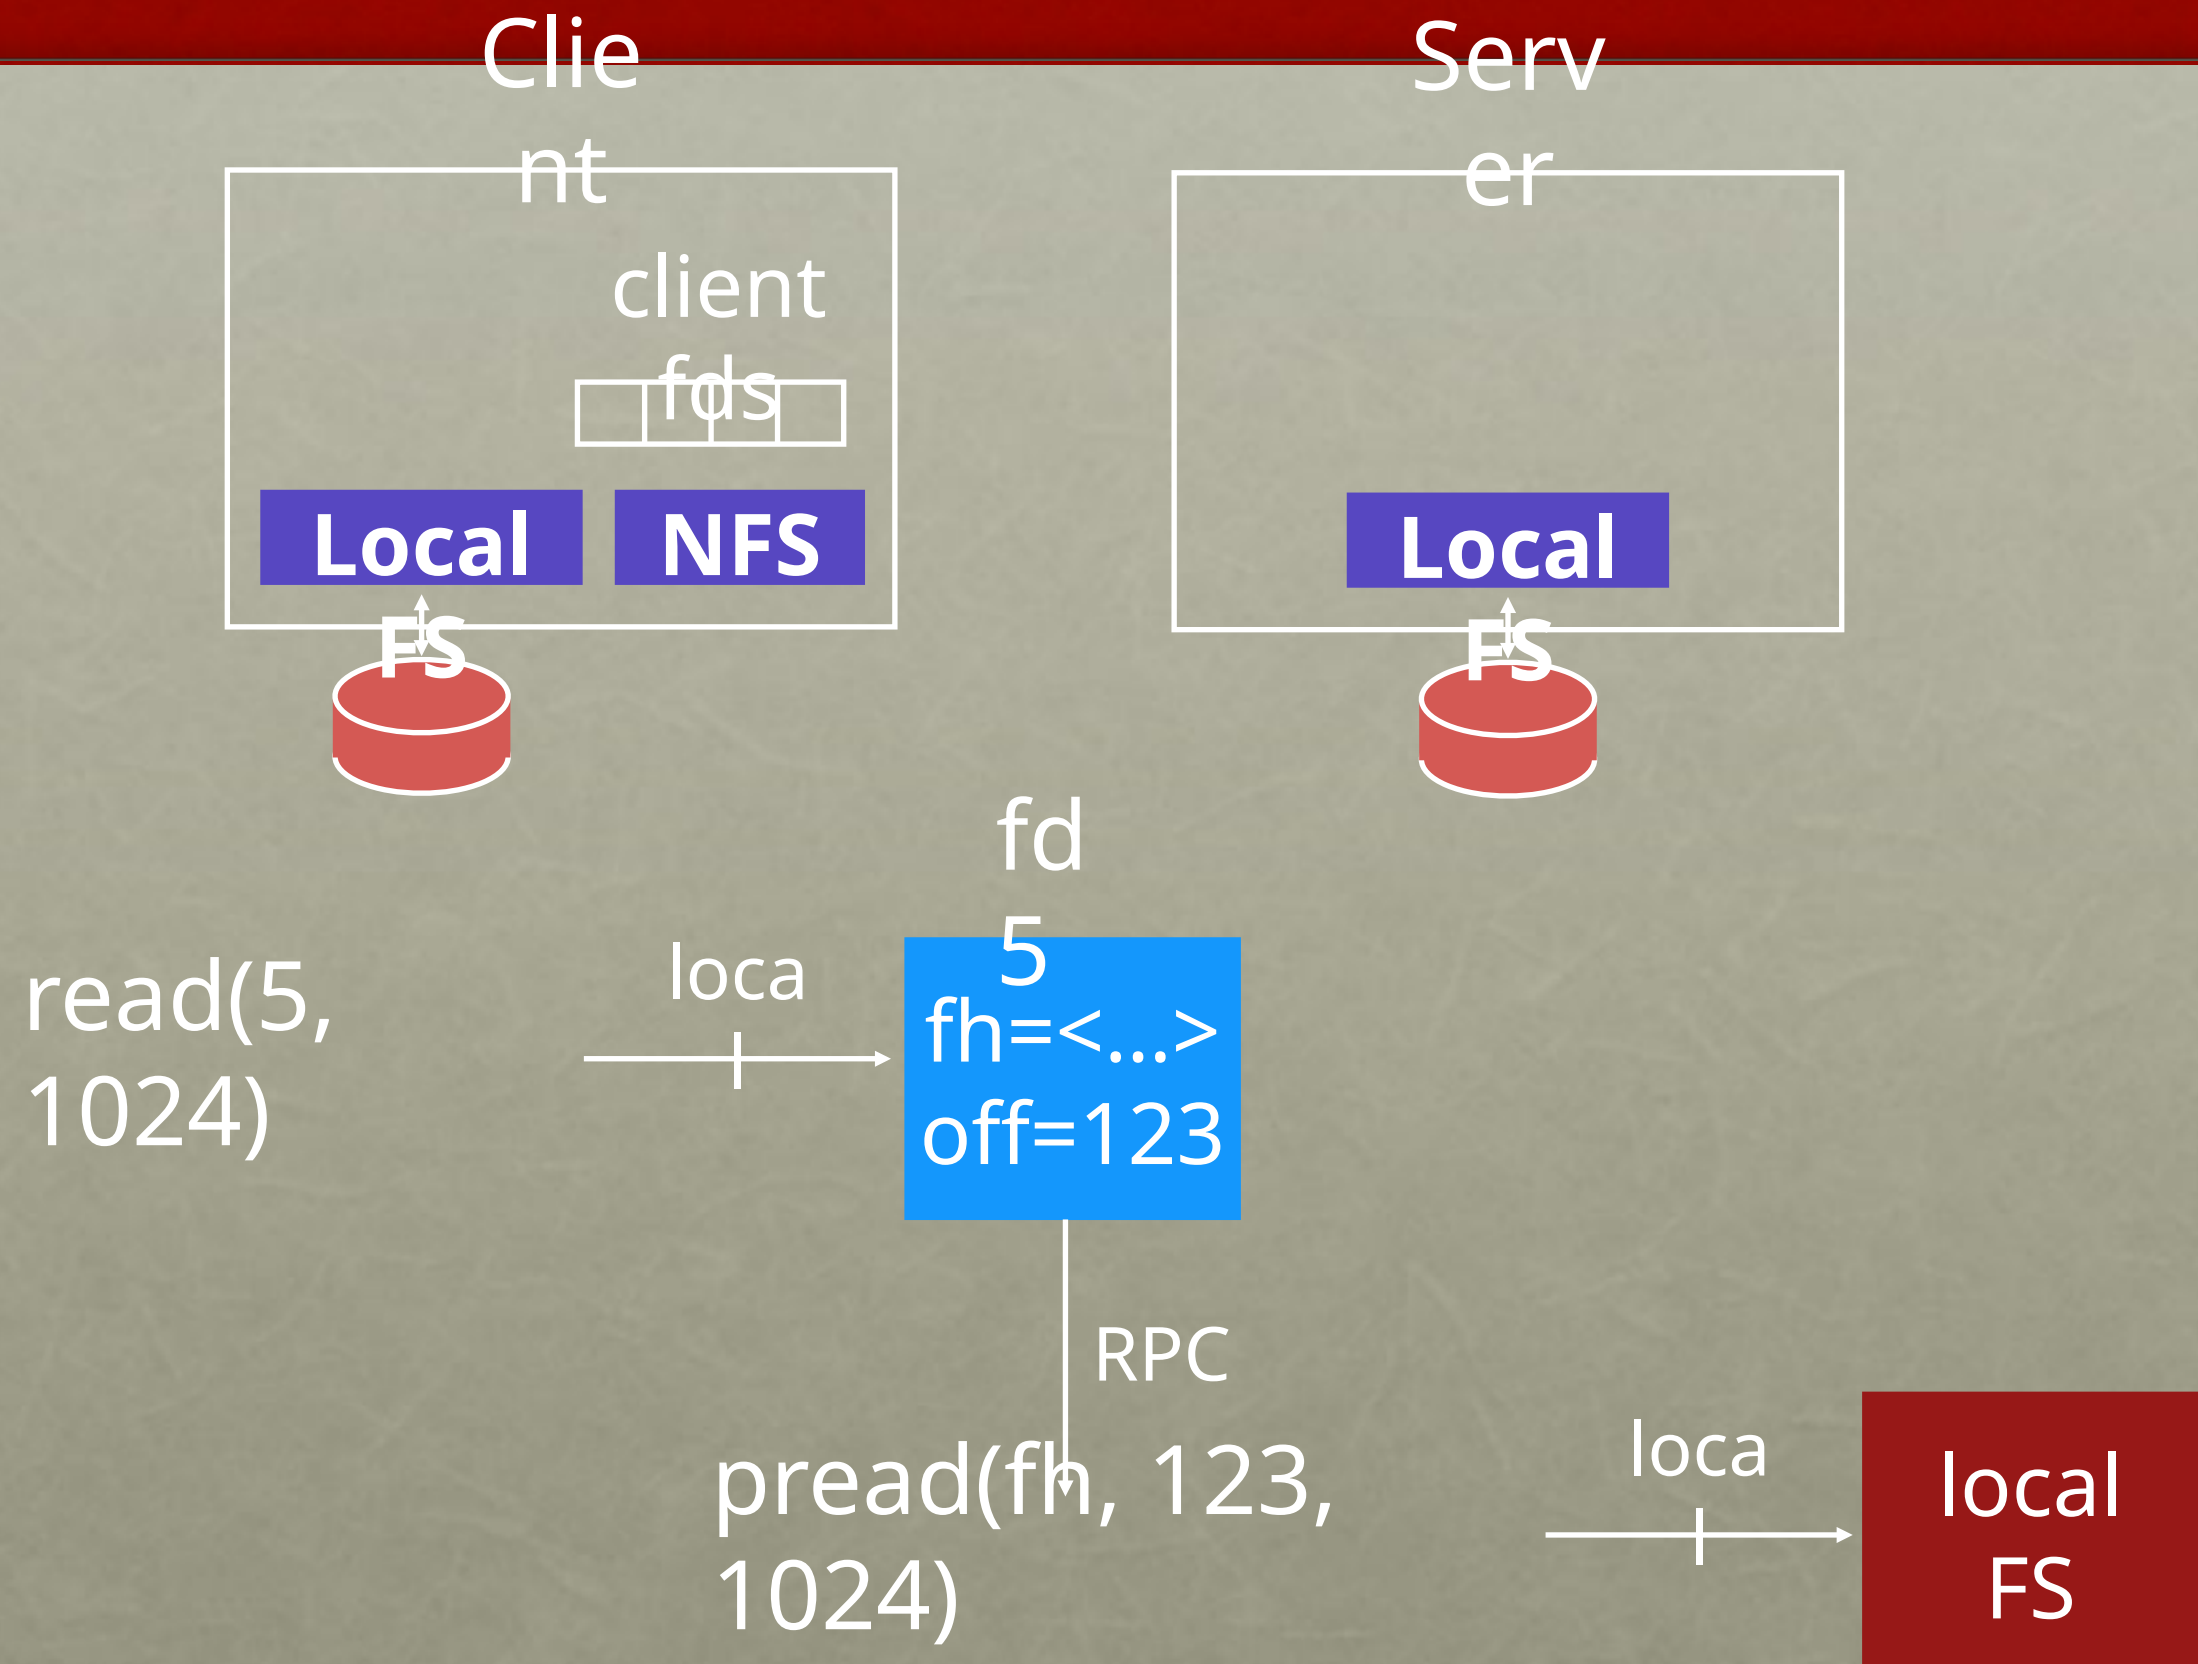
Build normal UNIX API on client side on top of idempotent, RPC-based API

Client `open()` creates a local fd object

It contains:

- file handle
- offset

# File Descriptors



# Overview


~~Architecture~~

~~Network API~~

Write Buffering

Cache

# Write Buffers



server acknowledges write before write is pushed  
to disk;  
what happens if server crashes?



# Server Write Buffer Lost

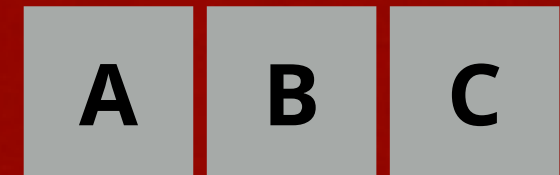
client:

write A to 0

write B to 1

write C to 2

server  
mem:



server disk:



server acknowledges write before write is pushed  
to disk

# Server Write Buffer Lost

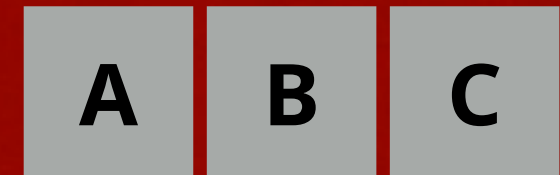
client:

write A to 0

write B to 1

write C to 2

server  
mem:



server disk:



server acknowledges write before write is pushed  
to disk

# Server Write Buffer Lost

client:

write A to 0

write B to 1

write C to 2

write X to 0

server  
mem:

**X**

**B**

**C**

server disk:

**A**

**B**

**C**

server acknowledges write before write is pushed  
to disk

# Server Write Buffer Lost

client:

write A to 0

write B to 1

write C to 2

write X to 0

server  
mem:

**X**

**B**

**C**

server disk:

**X**

**B**

**C**

server acknowledges write before write is pushed  
to disk



# Server Write Buffer Lost

client:

write A to 0

write B to 1

write C to 2

write X to 0

write Y to 1

server  
mem:

**X**

**Y**

**C**

server disk:

**X**

**B**

**C**

server acknowledges write before write is pushed  
to disk



# Server Write Buffer Lost

client:

write A to 0

write B to 1

write C to 2

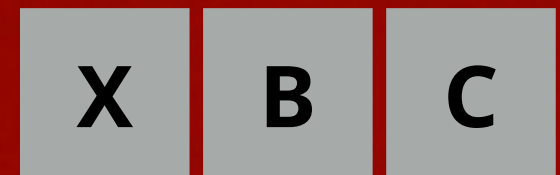
write X to 0

write Y to 1

server  
mem:



server disk:



crash  
!

server acknowledges write before write is pushed  
to disk

# Server Write Buffer Lost

client:

write A to 0

write B to 1

write C to 2


write X to 0

write Y to 1

server  
mem:



server disk:



server acknowledges write before write is pushed  
to disk

# Server Write Buffer Lost

client:

write A to 0

write B to 1


write C to 2

write X to 0

write Y to 1

write Z to 2

server  
mem:



server disk:



server acknowledges write before write is pushed  
to disk

# Server Write Buffer Lost

client:

write A to 0


write B to 1

write C to 2

write X to 0

write Y to 1

write Z to 2

server  
mem: 

server disk: 


Problem:

No write failed, but disk state  
doesn't match any point in time

Solutions????



# Write Buffers

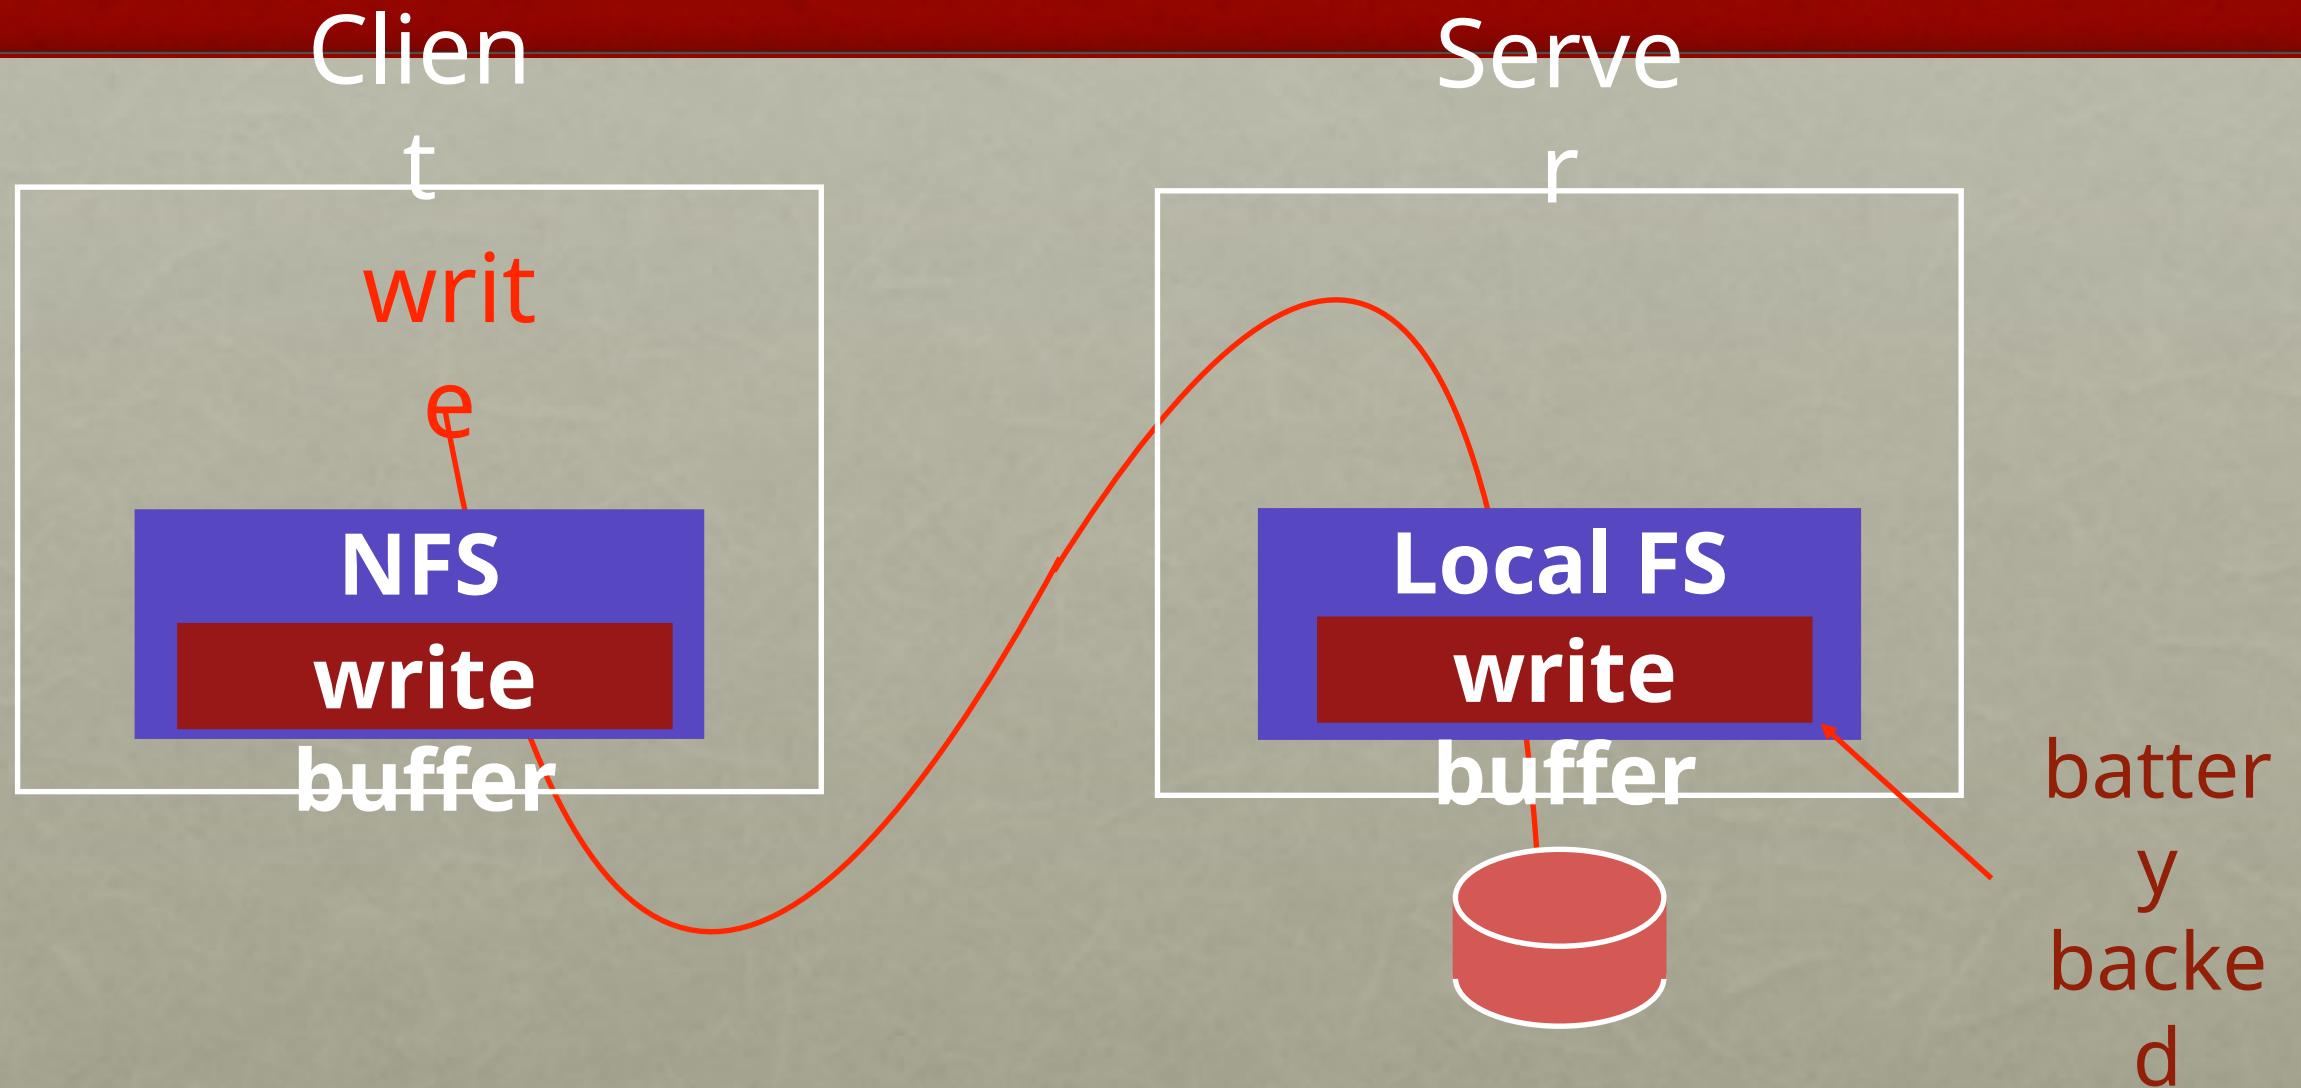


1. Don't use server write buffer  
(persist data to disk before acknowledging  
write)

Problem: Slow!



# Write Buffers



2. use persistent write buffer (more expensive)

# Overview

~~Architecture~~

~~Network API~~

~~Write Buffering~~

Cache


# Cache Consistency

NFS can cache data in three places:


- server memory
- client disk
- client memory

How to make sure all versions are in sync?

# Distributed Cache




# Cache






# Cache



# Cache

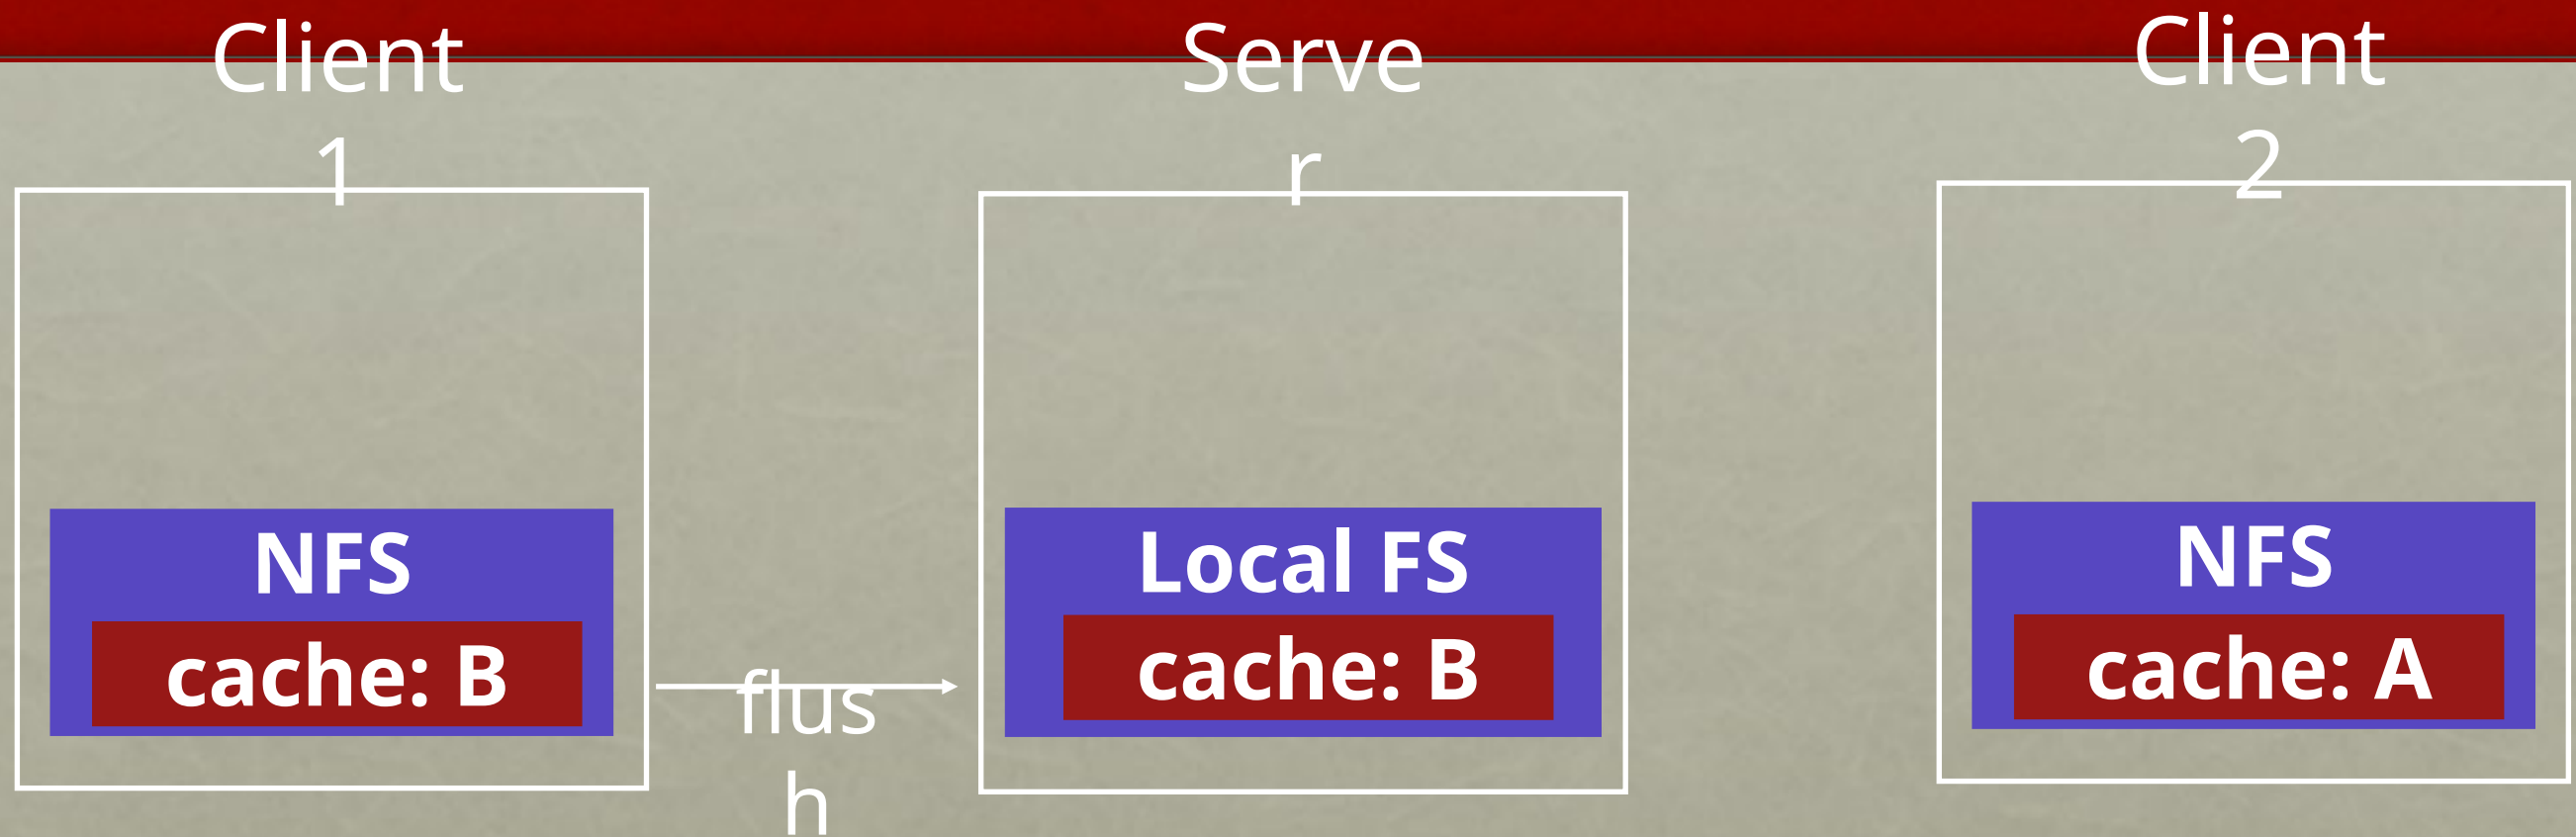


“Update Visibility” problem:  
server doesn't have latest version

What happens if Client 2 (or any other client) reads data?

Sees old version (different semantics than local FS)


# Cache



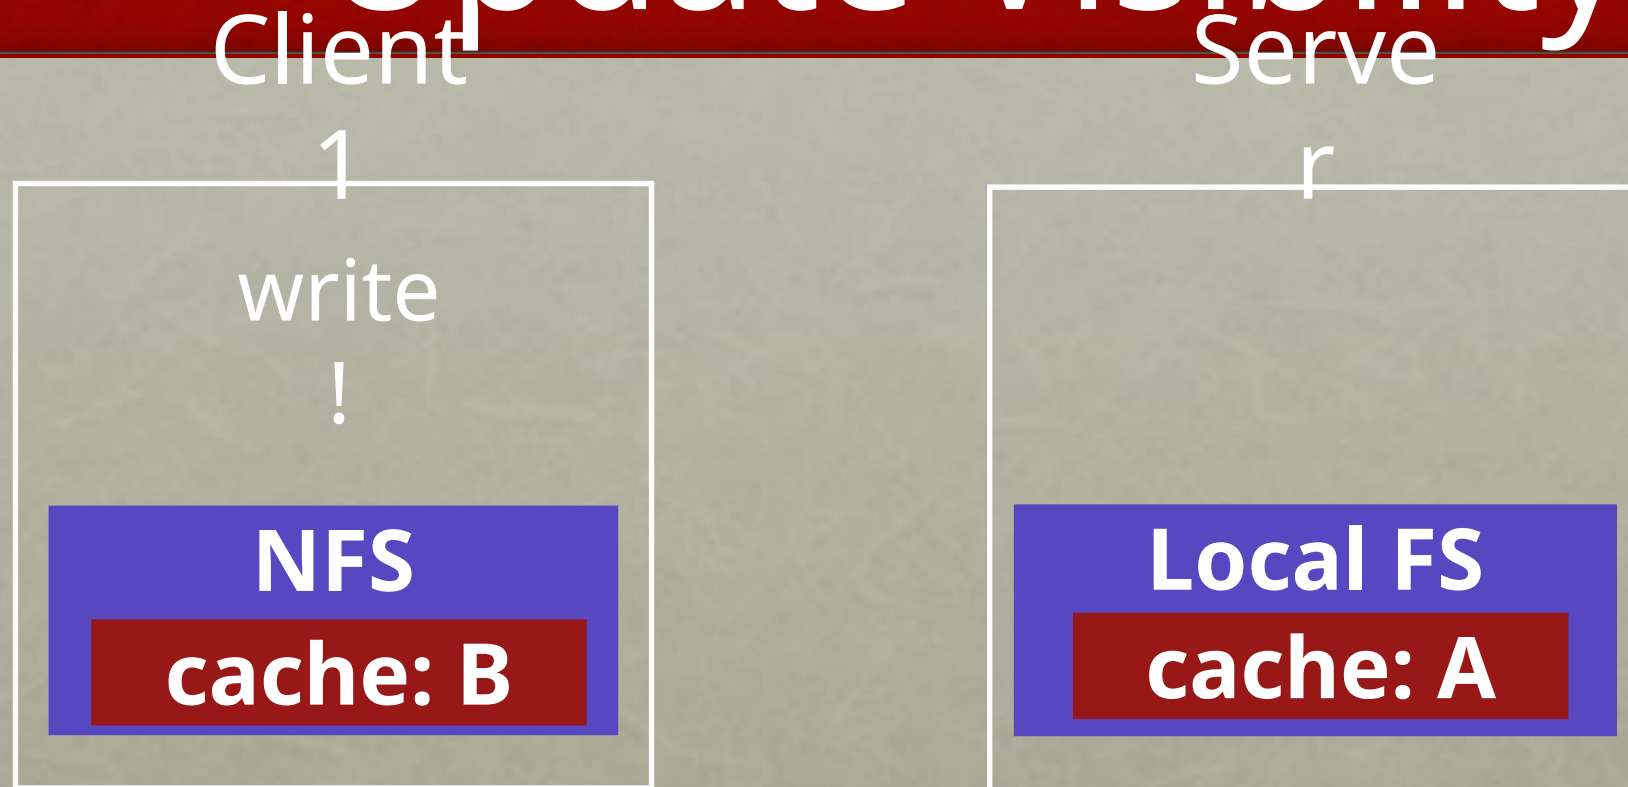
“Stale Cache” problem:  
client 2 doesn't have latest version

What happens if Client 2 reads data?  
Sees old version (different semantics than  
local FS)

# Cache



# Problem 1: Update Visibility



When client buffers a write, how can server (and other clients) see update?

- Client flushes cache entry to server

**When** should client perform flush????? (3 reasonable options??)

NFS solution: flush on fd close



# Problem 2: Stale Cache



Problem: Client 2 has stale copy of data; how can it get the latest?

One possible solution:

- If NFS had state, server could push out update to relevant clients

NFS solution:

- Clients recheck if cached copy is current before using data

# Stale Cache Solution



Client cache records time when data block was fetched ( $t_1$ )

Before using data block, client does a STAT request to server

- get's last modified timestamp for this file ( $t_2$ ) (not block...)
- compare to cache timestamp
- refetch data block if changed since timestamp ( $t_2 > t_1$ )

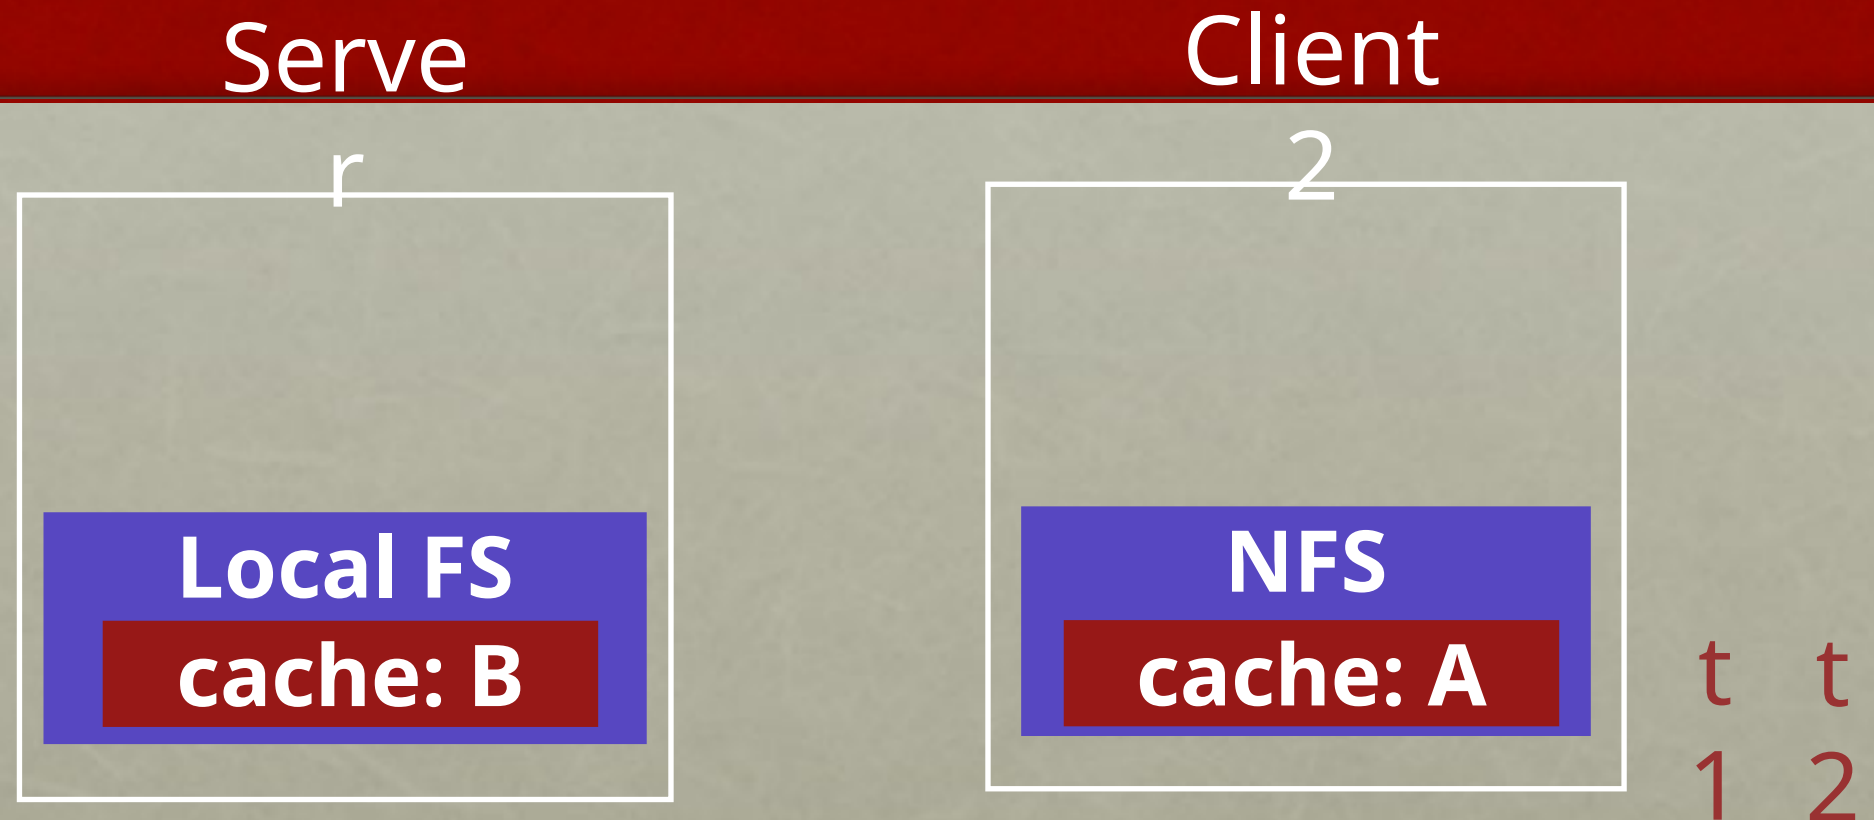
# Measure then Build

NFS developers found stat accounted for 90% of server requests

Why?

Because clients frequently recheck cache

# Reducing Stat Calls



Solution: cache results of stat calls

What is the result? **Never see updates on server!**

Partial Solution: Make stat cache entries expire after a given time (e.g., 3 seconds) (discard t2 at client 2)

What is the result? **Could read data that is up to 3 seconds old**



# NFS Summary

NFS handles client and server crashes very well; robust APIs are often:

- **stateless**: servers don't remember clients
- **idempotent**: doing things twice never hurts

Caching and write buffering is harder in distributed systems, especially with crashes

Problems:

- Consistency model is odd (client may not see updates until 3 seconds after file is closed)
- Scalability limitations as more clients call `stat()` on server



# AFS Goals


Primary goal: scalability! (many clients per server)

More reasonable semantics for concurrent file access

# AFS Design


NFS: Server exports local FS

AFS: Directory tree stored across many server machines  
(helps scalability!)



Break directory tree into  
“volumes”  
I.e., partial sub trees


# Volume Architecture




collection of servers store different volumes that together form directory tree

# Volume Architecture

Server



Server






Server



volumes may be moved by  
an administrator.




# Volume Architecture



volumes may be moved by an administrator.



# Volume Architecture



Client library gives seamless view of directory tree by automatically finding volumes


Communication via RPC  
Servers store data in local file systems

# AFS Cache Consistency


Update visibility

Stale cache

# Update Visibility



# Update Visibility



“Update Visibility” problem: server doesn’t have latest.

# Update Visibility

NFS solution is to flush blocks

- on close()
- other times too – e.g., when low on memory

Problems

- flushes not atomic (one block at a time)
- two clients flush at once: mixed data



# Update Visibility


AFS solution:

- also flush on close
- buffer **whole files** on local disk; update file on server atomically

Concurrent writes?

- Last writer (i.e., last file closer) wins
- Never get mixed data on server

# Cache Consistency



“Stale Cache” problem: client 2 doesn’t have latest

# Stale Cache

NFS rechecks cache entries compared to server before using them, assuming check hasn't been done "recently"

How to determine how recent? (about 3 seconds)

"Recent" is too long?

"Recent" is too short?

client reads old  
data  
server overloaded with  
stats

# Stale Cache



AFS solution: Tell clients when data is overwritten

- Server must remember which clients have this file open right now

When clients cache data, ask for “callback” from server if changes

- Clients can use data without checking all the time

Server no longer stateless!



# Callbacks: Dealing with STATE

What if client crashes?

What if server runs out of memory?

What if server crashes?



# Client Crash



What should client do after reboot?  
(remember cached data can be on disk too...)

Concern? **may have missed notification that cached copy changed**

Option 1: evict everything from cache

Option 2: ??? **recheck entries before using**

# Low Server Memory



Strategy: tell clients you are dropping their callback

What should client do?

Option 1: Discard entry from cache

Option 2: ???

Mark entry for  
recheck

# Server Crashes

What if server crashes?

Option: tell all clients to recheck all data before next read

Handling server and client crashes without inconsistencies or race conditions is very difficult...

# Prefetching

AFS paper notes: “the study by Ousterhout *et al.* has shown that most files in a 4.2BSD environment are read in their entirety.”

What are the implications for client prefetching policy?

Aggressively prefetch whole files.



# Whole-File Caching

Upon open, AFS client fetches whole file (even if huge), storing in local memory or disk

Upon close, client flushes file to server (if file was written)

Convenient and intuitive semantics:

- AFS needs to do work only for open/close
  - Only check callback on open, not every read
- reads/writes are local
  - Use same version of file entire time between open and close



# AFS Summary

**State** is useful for **scalability**, but makes handling crashes hard

- Server tracks callbacks for clients that have file cached
- Lose callbacks when server crashes...

Workload drives design: **whole-file caching**

- More intuitive semantics (see version of file that existed when file was opened)

# AFS vs nfs Protocols

Can you summarize the consistency semantics provided by NFSv2?

Time	Client A	Client B	Server Action?
0	fd = open("file A");		
10	read(fd, block1);		
20	read(fd, block2);		
30	read(fd, block1);		
31	read(fd, block2);		
40		fd = open("file A");	
50		write(fd, block1);	
60	read(fd, block1);		
70		close(fd);	
80	read(fd, block1);		
81	read(fd, block2);		
90	close(fd);		
100	fd = open("fileA");		
110	read(fd, block1);		
120	close(fd);		

When will server be contacted for NFS? For AFS?

# Nfs Protocol

Time	Client A	Client B	Server Action?
0	fd = open("file A");		lookup()
10	read(fd, block1);		read
20	read(fd, block2);		read
30	read(fd, block1);	check cache: attr expired getattr() → okay, use local	get attr
31	read(fd, block2);	attr not expired, use local	
40		fd = open("file A");	lookup
50		write(fd, block1);	keep local
60	read(fd, block1);	attr. expired use local data	getattr()
70		close(fd); write bl to server!	write to disk
80	read(fd, block1);	attr expired. get attr. CHANGED FILE - kickout	read()
81	read(fd, block2);	not in cache → read	read()
90	close(fd);		
100	fd = open("fileA");		lookup
110	read(fd, block1);	attr expire; get new attr local ok	get attr
120	close(fd);		



# AFS Protocol

Time	Client A	Client B	Server Action?
0	fd = open("file A");		setup callback for A
10	read(fd, block1);		send all of file A
20	read(fd, block2);		
30	read(fd, block1);		
31	read(fd, block2);		
40		fd = open("file A");	setup callback
50		write(fd, block1);	send all of A
60	read(fd, block1); local		
70		close(fd);	send block changes of A break call backs
80	read(fd, block1); local		
81	read(fd, block2); local		
90	close(fd); nothing changed		
100	fd = open("fileA"); no callback!! need to fetch A again		
110	read(fd, block1);		
120	close(fd);		send A