

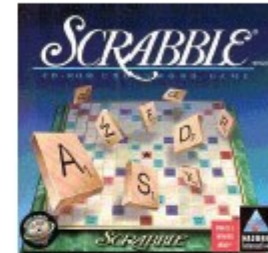
# CSE 401

*Game Playing  
Alpha-Beta Search  
and  
General Issues*

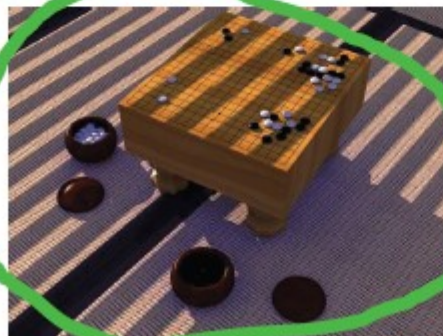
## Which of these are: 2-player zero-sum discrete finite deterministic games of perfect information



- **Two player:** Duh!
- **Zero-sum:** In any outcome of any game, Player A's gains equal player B's losses.
- **Discrete:** All game states and decisions are discrete values.
- **Finite:** Only a finite number of states and decisions.
- **Deterministic:** No chance (no die rolls).
- **Games:** See next page
- **Perfect information:** Both players can see the state, and each decision is made sequentially (no simultaneous moves).



## Which of these are: 2-player zero-sum discrete finite deterministic games of perfect information



- **Two player:** Duh!
- **Zero-sum:** In any outcome of any game, Player A's gains equal player B's losses.
- **Discrete:** All game states and decisions are discrete values.
- **Finite:** Only a finite number of states and decisions.
- **Deterministic:** No chance (no die rolls).
- **Games:** See next page
- **Perfect information:** Both players can see the state, and each decision is made sequentially (no simultaneous moves).

# Definition

A Two-player zero-sum discrete finite deterministic game of perfect information is a quintuplet:  $(S, I, Succs, T, V)$  where

$S$	=	a finite set of states (note: state includes information sufficient to deduce who is due to move)
$I$	=	the initial state
$Succs$	=	a function which takes a state as input and returns a set of possible next states available to whoever is due to move
$T$	=	a subset of $S$ . It is the terminal states: the set of states at which the game is over
$V$	=	a mapping from terminal states to real numbers. It is the amount that A wins from B. (If it's negative A loses money to B).

Convention: assume Player A moves first.

For convenience: assume turns alternate.

# Utility Function

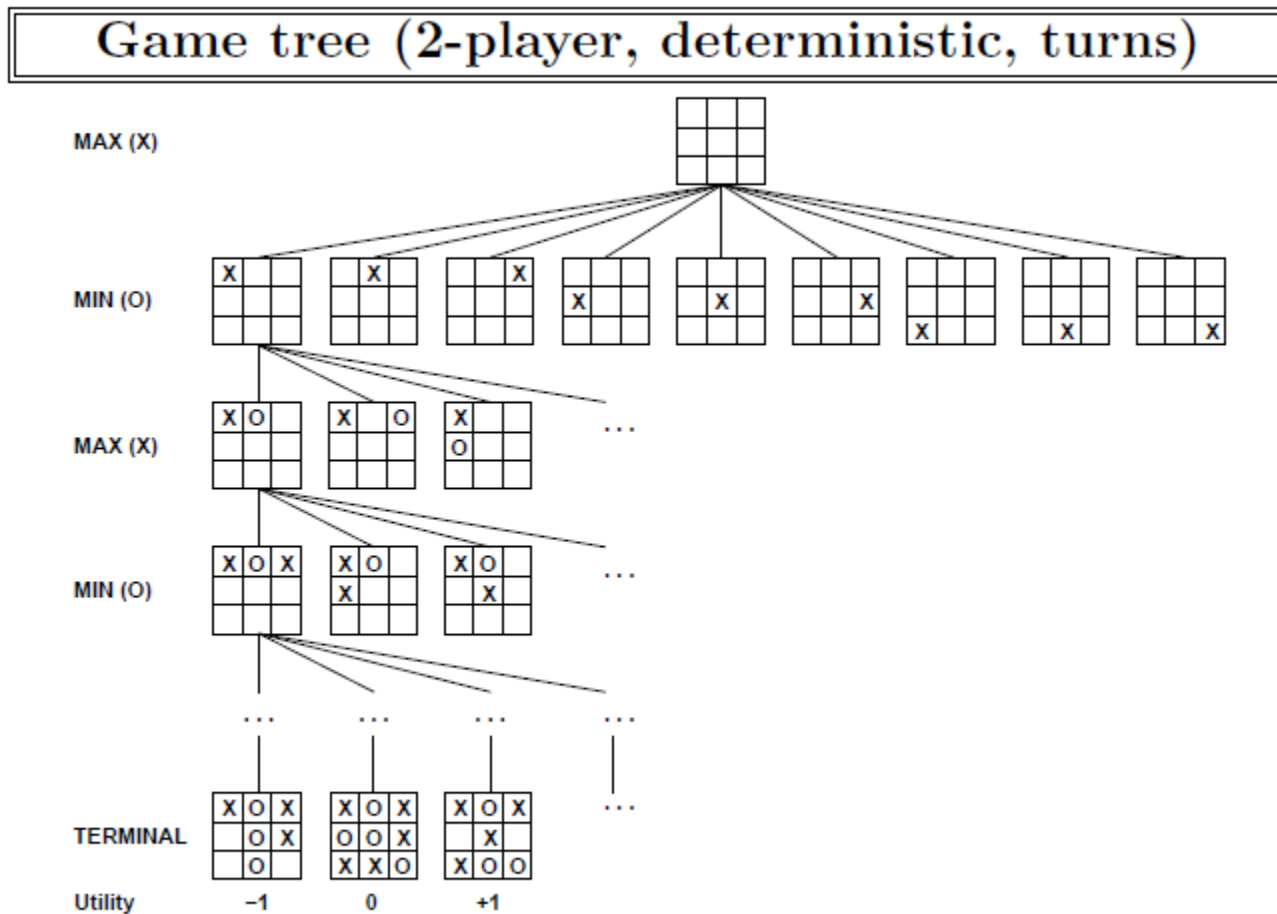
- Gives the utility of a game state
  - $\text{utility}(\text{State})$
- Examples
  - -1, 0, and +1, for Player 1 loses, draw, Player 1 wins, respectively
  - Difference between the point totals for the two players
  - Weighted sum of factors (e.g. Chess)
    - $\text{utility}(S) = w_1 f_1(S) + w_2 f_2(S) + \dots + w_n f_n(S)$ 
      - $f_1(S) = (\text{Number of white queens}) - (\text{Number of black queens}), \quad w_1 = 9$
      - $f_2(S) = (\text{Number of white rooks}) - (\text{Number of black rooks}), \quad w_2 = 5$



# *Game Playing*

- **Game tree**
  - describes the possible sequences of play
  - is a graph if we merge together identical states
- **Minimax:**
  - utility values assigned to the leaves
- **Values “backed up” the tree by**
  - MAX node takes max value of children
  - MIN node takes min value of children
  - Can read off best lines of play
- **Depth Bounded Minimax**
  - utility of terminal states estimated using an “evaluation function”

# Game Tree for Tic-tac-toe



# Minimax

MINIMAX-VALUE( $n$ ) =

$$\begin{cases} \text{UTILITY}(n) & \text{if } n \text{ is a terminal state} \\ \max_{s \in \text{Successors}(n)} \text{MINIMAX-VALUE}(s) & \text{if } n \text{ is a MAX node} \\ \min_{s \in \text{Successors}(n)} \text{MINIMAX-VALUE}(s) & \text{if } n \text{ is a MIN node.} \end{cases}$$



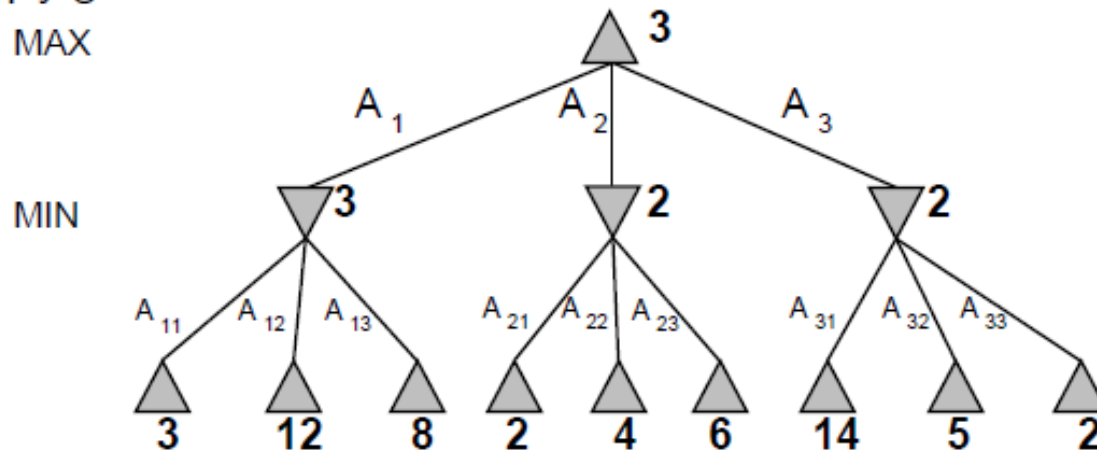
# MiniMax

## Minimax

Perfect play for deterministic, perfect-information games

Idea: choose move to position with highest **minimax value**  
= best achievable payoff against best play

E.g., 2-ply game:



# Basic Algorithm

function **MINIMAX-DECISION**(*state*) returns *an action*

inputs: *state*, current state in game

return the *a* in **ACTIONS**(*state*) maximizing **MIN-VALUE**(**RESULT**(*a*, *state*))

---

function **MAX-VALUE**(*state*) returns *a utility value*

if **TERMINAL-TEST**(*state*) then return **UTILITY**(*state*)

$v \leftarrow -\infty$

for *a*, *s* in **SUCCESSORS**(*state*) do  $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return *v*

---

function **MIN-VALUE**(*state*) returns *a utility value*

if **TERMINAL-TEST**(*state*) then return **UTILITY**(*state*)

$v \leftarrow \infty$

for *a*, *s* in **SUCCESSORS**(*state*) do  $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return *v*

MAX

MIN

MAX

1

A

1

B

-3

C

4

D

1

E

2

F

-3

G

4

-5

-5

1

-7

2

-3

-8



= terminal position



= agent



= opponent

# *Game Playing – Beyond Minimax*

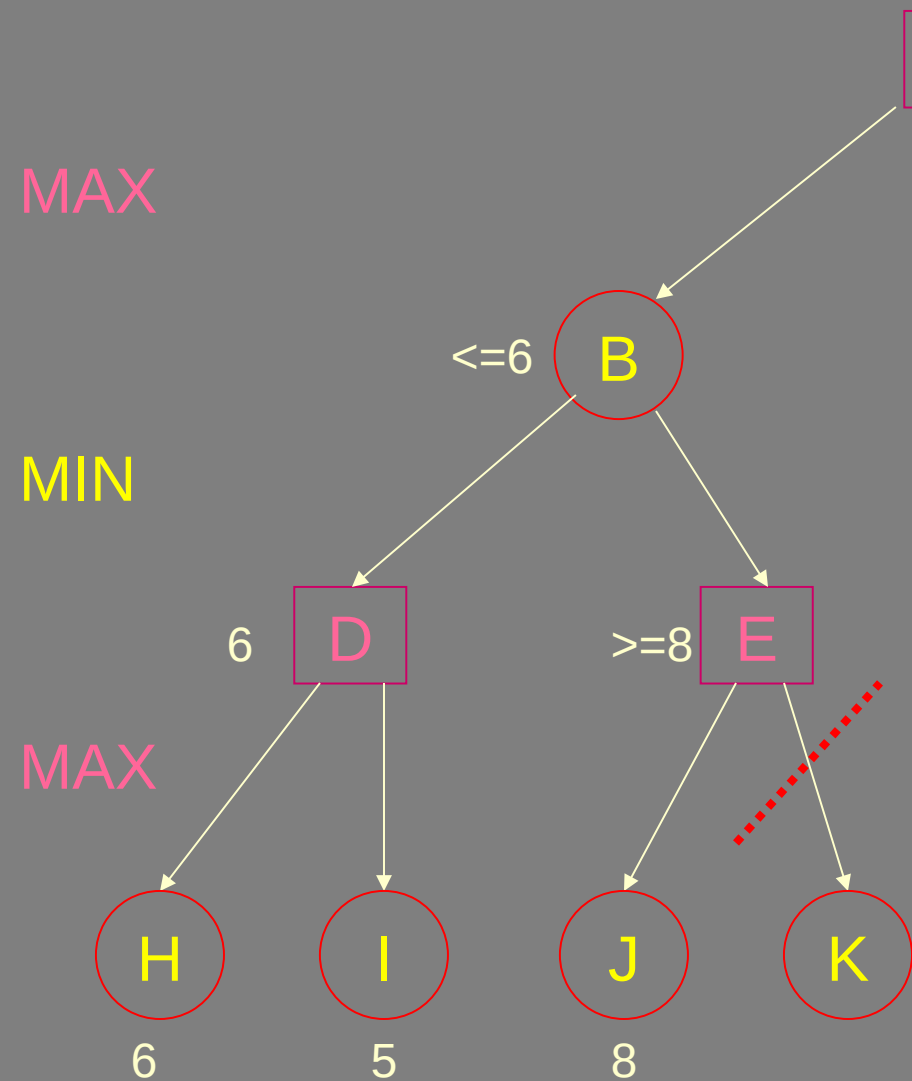
- **Efficiency of the search**
  - Game trees are very big
  - Evaluation of positions is time-consuming
- **How can we reduce the number of nodes to be evaluated?**
  - “alpha-beta search”
- **Bounding the depth of minimax has deficiencies**
  - Why?
  - How can we mitigate these deficiencies?


## *Game Playing – Improving Efficiency*


- **Suppose that we are doing depth-bounded minimax**
- **We have a game tree to create and then insert the minimax values in order to find the values for our possible moves from the current position**

## *Game Playing – Minimax using DFS*

- The presentation of minimax was done by “backing up from the leaves” – a “bottom-up” breadth-first search.
- This has the disadvantage of taking a lot of space
  - Compare this with the space usage issues for DFS vs. BFS in earlier lectures
- If we can do minimax using DFS then it is likely to take a lot less space
- Minimax can be implemented using DFS
- But reduced space is not the only advantage:



 = agent

 = opponent

**STOP! What else can  
you deduce now!?**

On discovering  $\text{util}(D) = 6$   
we know that  $\text{util}(B) \leq 6$

On discovering  $\text{util}(J) = 8$   
we know that  $\text{util}(E) \geq 8$

Can stop expansion of E as best  
play will not go via E

Value of K is irrelevant – prune it!



## *Game Playing – Pruning nodes*

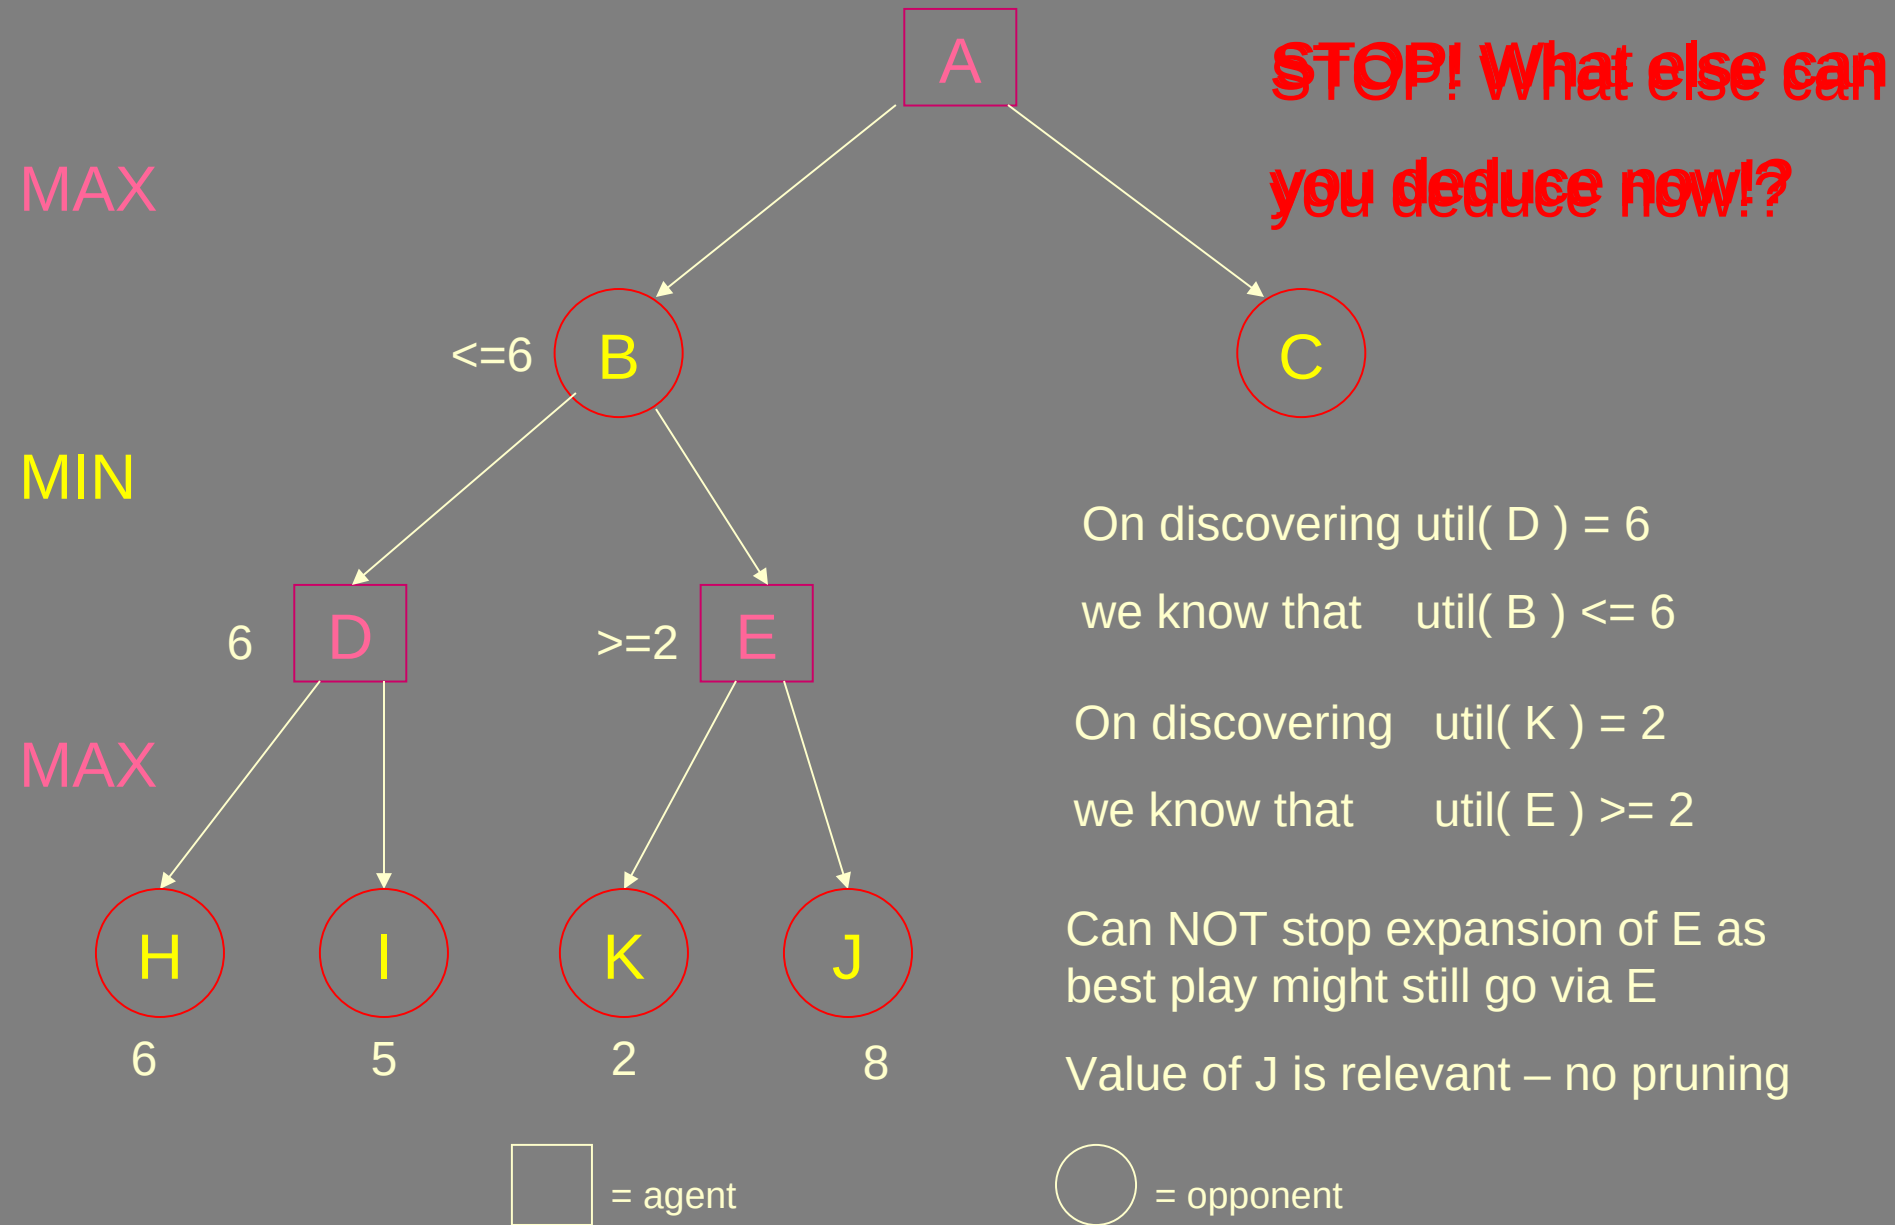
- If we are scanning the tree using DFS then there was no point in evaluating node K
- Whatever the value of K there cannot be any rational sequence of play that would go through it
  - **Node K can be pruned from the search: i.e. just not selected for further expansion**
- “At node B then MIN will never select E; because J is better than D for MAX and so MIN must not allow MAX to have that opportunity”
- Q. So what! It's just one node?
- A. Suppose that the depth limit were such that K was far from the depth bound. Then evaluating K corresponds to a large sub-tree. Such prunings can save an exponential amount of work

## *Game Playing – Improving Efficiency*

- Suppose that we were doing Breadth-First Search, would you still be able to prune nodes in this fashion?
- NO! Because the pruning relied on the fact that we had already evaluated node D by evaluating the tree underneath D
- This form of pruning is an example of “alpha-beta pruning” and relies on doing a DEPTH-FIRST search of the depth bounded tree

## *Game Playing – Node-ordering*

- **Suppose that**
  - nodes K and J were evaluated in the opposite order
  - can we expect that we would be able to do a similar pruning?
- **The answer depends on the value of K**
- **Suppose that K had a value of 2 and is expanded first:**

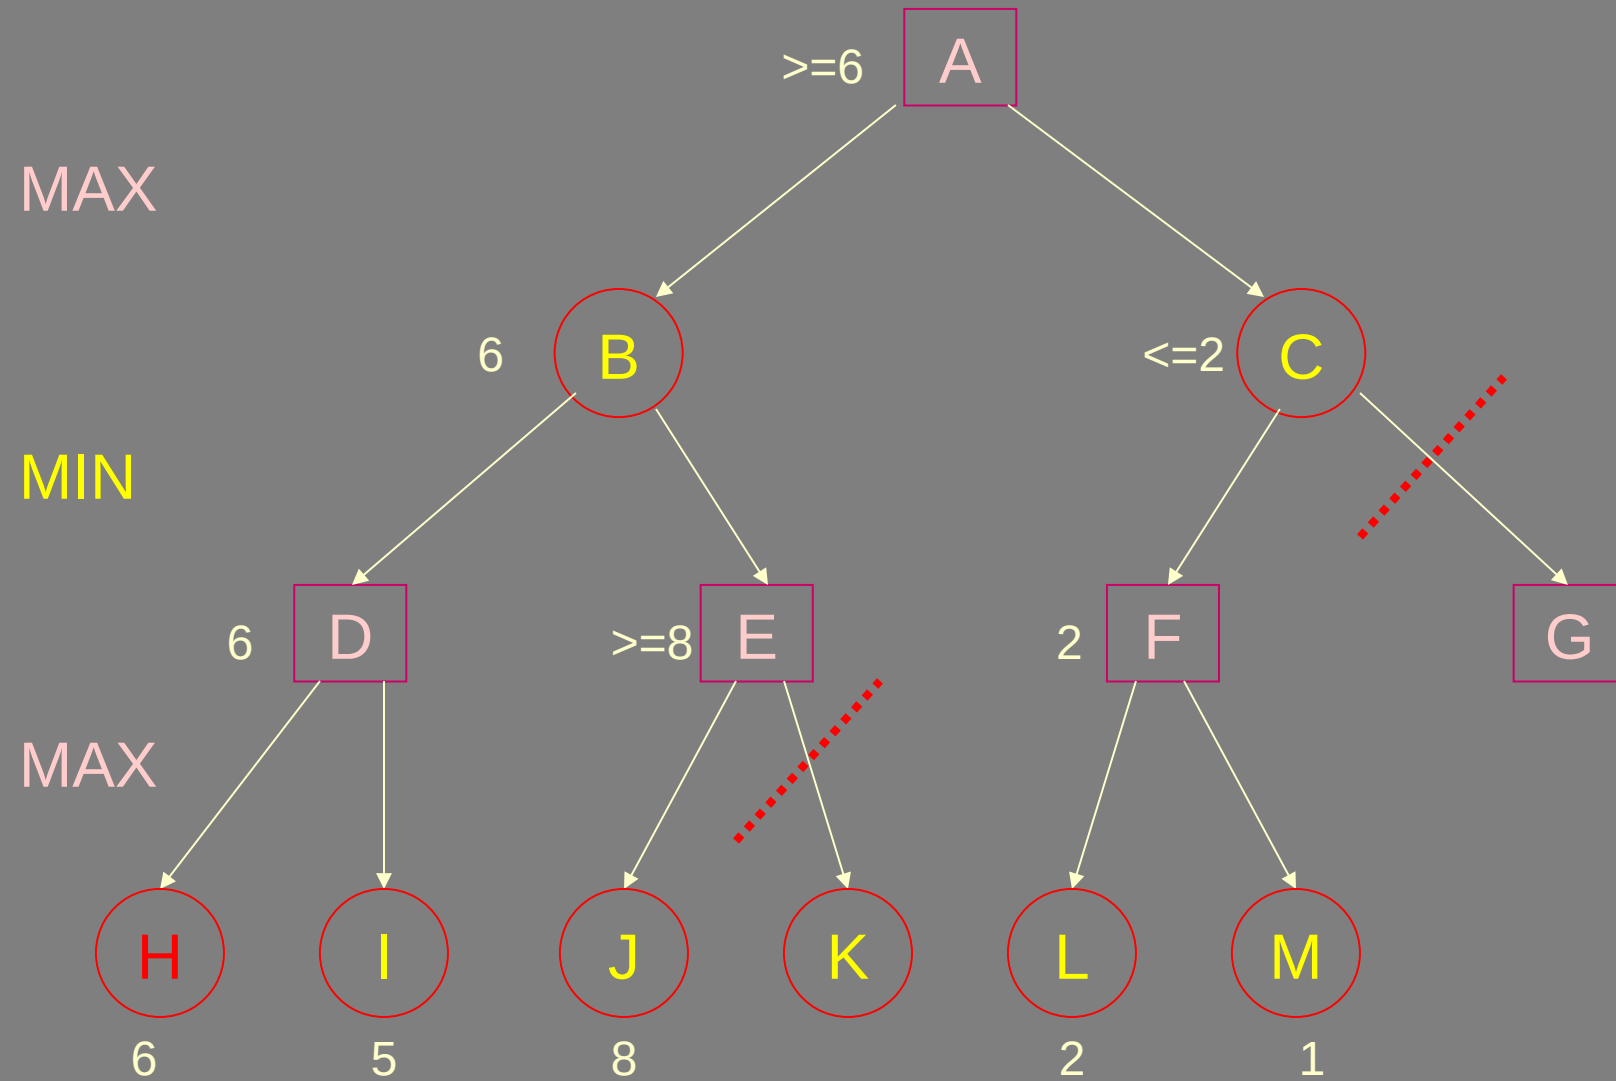



## *Game Playing – Node-ordering*


- When K had a value of 2 and was expanded first then we did not get to prune a child of E
- To maximise pruning we want to first expand those children that are best for the parent
  - cannot know which ones are really best
  - use heuristics for the “best-first” ordering
- If this is done well then alpha-beta search can effectively double the depth of search tree that is searchable in a given time
  - Effectively reduces the branching factor in chess from about 30 to about 8
  - This is an enormous improvement!

## *Game Playing – Improving Efficiency*

- **The games are symmetric so is natural that we can also do a similar pruning with the MIN and MAX roles reversed**
- **The reasoning is identical other than for the reversal of roles**
- **Can deduce that some other nodes can not be involved in the line of best play**



 = agent

 = opponent



## *Game Playing – Alpha-Beta Implementation*

- The pruning was based on using the results of the “DFS so far” to deduce upper and lower bounds on the values of nodes
- Conventionally these bounds are stored in terms of two parameters
  - alpha  $\alpha$
  - beta  $\beta$

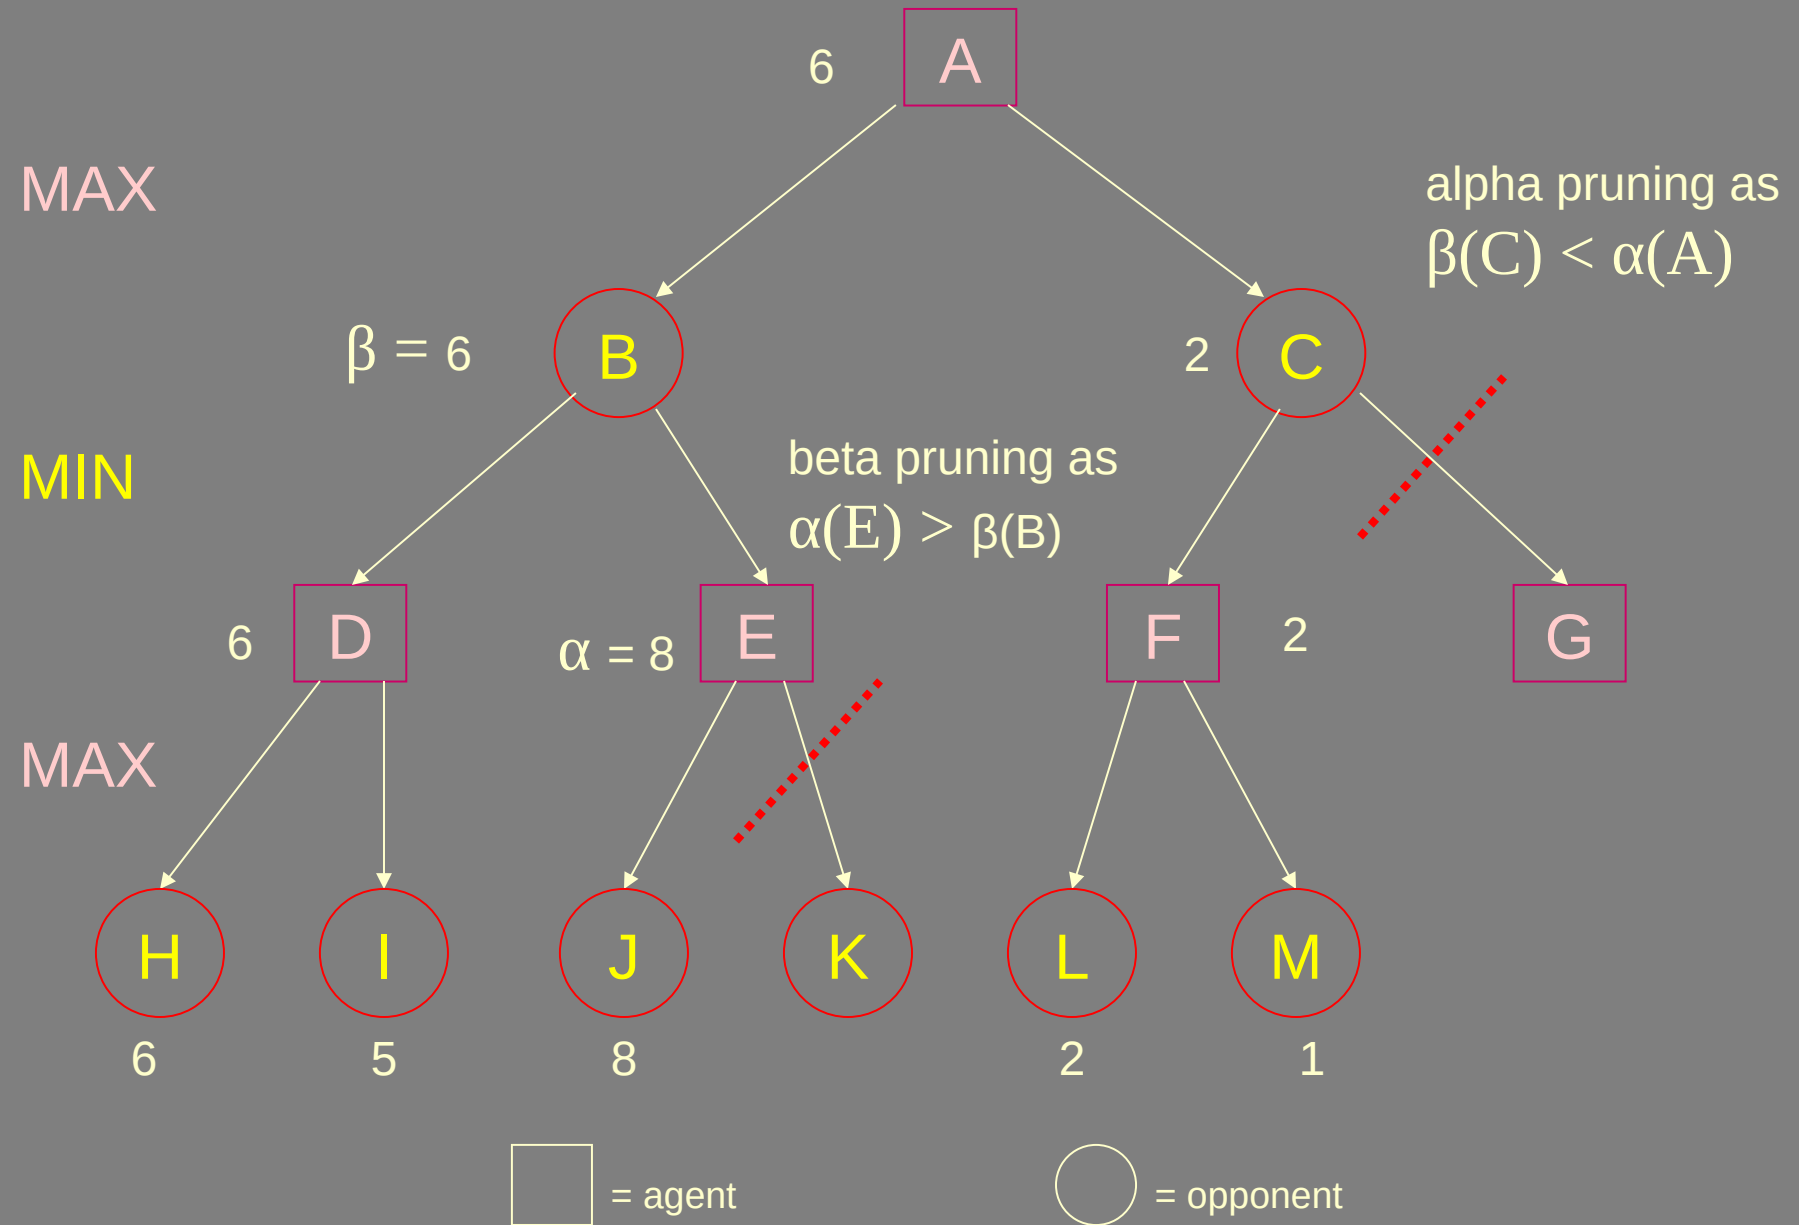
## *Game Playing – Alpha-Beta Implementation*

- **$\alpha$  values are stored with each MAX node**
- **each MAX node is given a value of alpha that is the current best lower-bound on its final value**
  - initially is  $-\infty$  to represent that nothing is known
  - as we do the search then  $\alpha$  at a node can increase, but it can never decrease – it always gets better for MAX

## *Game Playing – Alpha-Beta Implementation*

- **$\beta$  values are stored with each MIN node**
- **each MIN node is given a value of beta that is the current best upper-bound on its final value**
  - initially is  $+\infty$  to represent that nothing is known
  - as we do the search then  $\beta$  at a node can decrease, but it can never increase – it always gets better for MIN

# Alpha-beta Pruning



## Properties of $\alpha$ - $\beta$

Pruning **does not** affect final result

Good move ordering improves effectiveness of pruning

With “perfect ordering,” time complexity =  $O(b^{m/2})$   
 $\Rightarrow$  **doubles** solvable depth

A simple example of the value of reasoning about which computations are relevant (a form of **metareasoning**)

Unfortunately,  $35^{50}$  is still impossible!

## Resource limits

Standard approach:

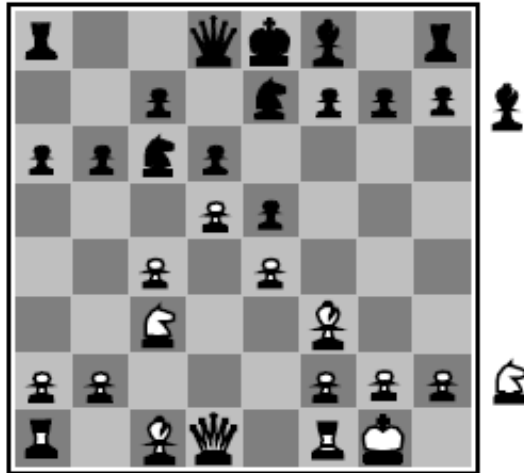
- Use CUTOFF-TEST instead of TERMINAL-TEST  
e.g., depth limit (perhaps add quiescence search)
- Use EVAL instead of UTILITY  
i.e., evaluation function that estimates desirability of position

Suppose we have 100 seconds, explore  $10^4$  nodes/second

$\Rightarrow 10^6$  nodes per move  $\approx 35^{8/2}$

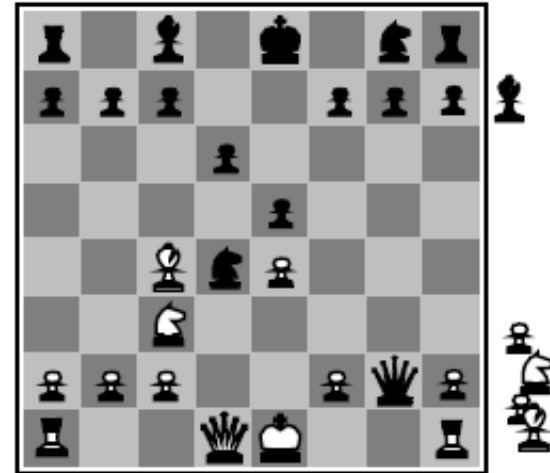
$\Rightarrow \alpha\text{-}\beta$  reaches depth 8  $\Rightarrow$  pretty good chess program

## Evaluation functions



Black to move

White slightly better



White to move

Black winning

For chess, typically linear weighted sum of features

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

e.g.,  $w_1 = 9$  with

$f_1(s) = (\text{number of white queens}) - (\text{number of black queens}), \text{ etc.}$



## *Game Playing – Deficiencies of Minimax*

- **The bound on the depth of search is artificial and can lead to many anomalies.**
- **We only consider two:**
  1. **Non-quiescence**  
“quiescent” = inactive, quiet, calm, ...
  2. **Horizon Effect**
- **(These deficiencies also apply to alpha-beta as it is just a more efficient way to do the same calculation as minimax)**

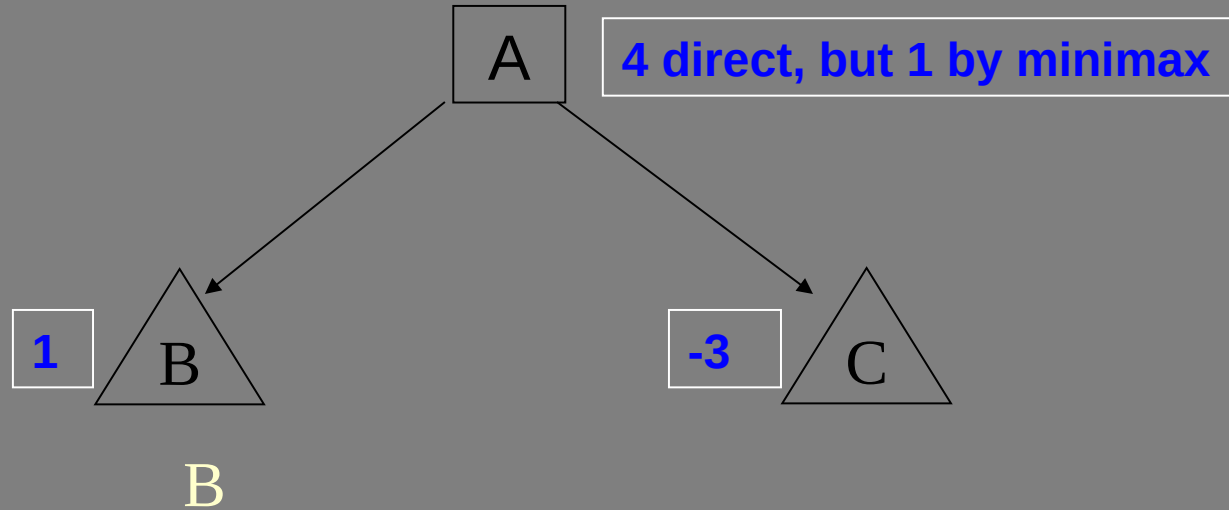
## *Game Playing – Non-Quiescence*

- **Suppose that change depth bound from  $k$  to  $k+1$  – i.e. expand one more move**
- **The values given to a node might change wildly**

# Example of non-quiescence

MAX

MIN



Utility values of “terminal” positions obtained by an evaluation function

Direct evaluation does not agree with one more expansion and then using of minimax



= terminal position



= agent



= opponent

## *Game Playing – Quiescence Search*

- Suppose that change depth bound from  $k$  to  $k+1$  – i.e. expand one more move
- The values given to a node might change wildly
- Keep on increasing the depth bound in that region of the game tree until the values become “quiescent” (“quiet”, i.e. stop being “noisy”)

## *Game Playing – Quiescence Search*

- **In quiescence search the depth bound is not applied uniformly but adapted to the local situation**
  - in this case so that the values are not wildly changing
- **Many other improvements to minimax also work by adapting to depth-bound to get better results and/or do less work**

## *Game Playing – Horizon Effect*

- Sometimes there is a bad thing, “X”, such that
  1. X cannot be avoided
  2. X can be delayed by some pointless moves
  3. X is not detected by the evaluation function
- In this case depth-limited minimax can be fooled
- It will use the pointless moves to push X beyond the depth limit, “horizon”, in which case it will “not see X”, and ignore it.
- This can lead the search to take bad moves because it ignores the inevitability of X

## *Game Playing – Beyond alpha-beta*

- We looked briefly at two problems
  - “non-quiescence”, and the “horizon effect”
- and one solution “quiescence search”
- To seriously implement a game
  - Deep-blue, chinook, etc
- it is necessary to solve many such problems!
- Good programs implement many techniques and get them to work together effectively



## *Game Playing – Game Classification*

- So far have only considered games such as chess, checkers, and nim.

These games are:

### 1. Fully observable

- Both players have full and perfect information about the current state of the game

### 2. Deterministic

- There is no element of chance
- The outcome of making a sequence of moves is entirely determined by the sequence itself

## *Game Playing – Game Classification*

- **Fully vs. Partially Observable**
- **Some games are only partially observable**
- **Players do not have access to the full “state of the game”**
- **E.g. card games – you typically cannot see all of your opponents cards**

## *Game Playing – Game Classification*

- **Deterministic vs. Stochastic**
- **In many games there is some element of chance**
- **E.g. Backgammon – throw dice in order to move**
- **(You are expected to be aware of these simple classifications)**

# *Game Playing – Summary*

- **Game Trees**
- **Minimax**
  - utility values propagate back to the root
- **Bounded Depth Minimax**
- **Alpha-Beta Search**
  - uses DFS
  - with depth bound
  - ordering of nodes is important in order to maximise pruning
- **Deficiencies of Bounded Depth Search**
  - **Non-quiescence**
    - Combat using quiescence search
  - **Horizon Problem**
    - Combat with ?? (look it up!)

## *End of Game Playing*



**Garry Kasparov and Deep  
Blue. © 1997, GM Gabriel  
Schwartzman's Chess  
Camera, courtesy IBM.**