

Chapter 3

The Data Link Layer

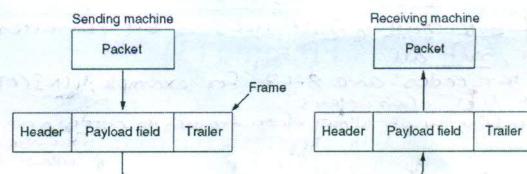
Data Link Layer Design Issues

- Services Provided to the Network Layer
- Framing
- Error Control
- Flow Control

Functions of the Data Link Layer

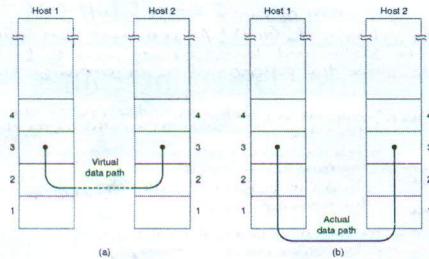
- Provide service interface to the network layer
- Dealing with transmission errors
- Regulating data flow
 - Slow receivers not swamped by fast senders

Functions of the Data Link Layer (2)



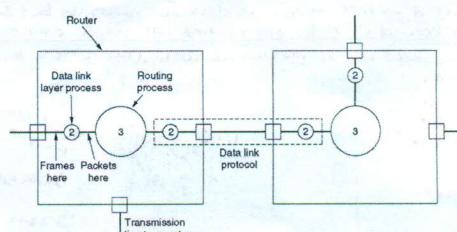
Relationship between packets and frames.

Services Provided to Network Layer



(a) Virtual communication.
(b) Actual communication.

Services Provided to Network Layer (2)



Placement of the data link protocol.

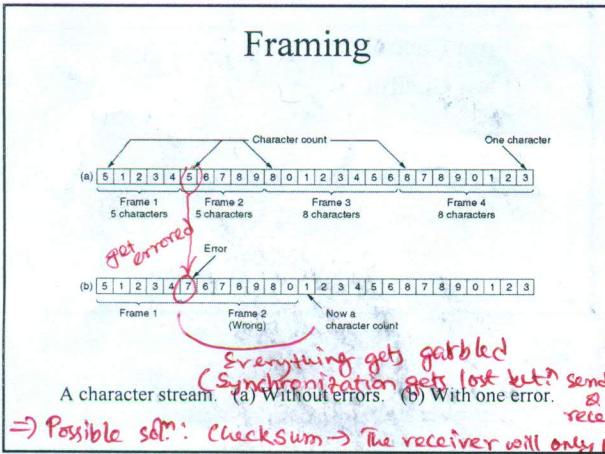
Framing \Rightarrow Breaking up bit streams into frames

intuitive way like our ordinary text writing \Rightarrow insert time gaps between frames
Problem: Networks rarely make any guarantees about timing, so it is possible these gaps might be squeezed out or other gaps might be inserted during transmission.

Remedy \Rightarrow four methods of framing

1. Character count
2. Flag bytes with byte stuffing
3. Starting and ending flags, with bit stuffing
4. Physical layer coding violations

8/26,

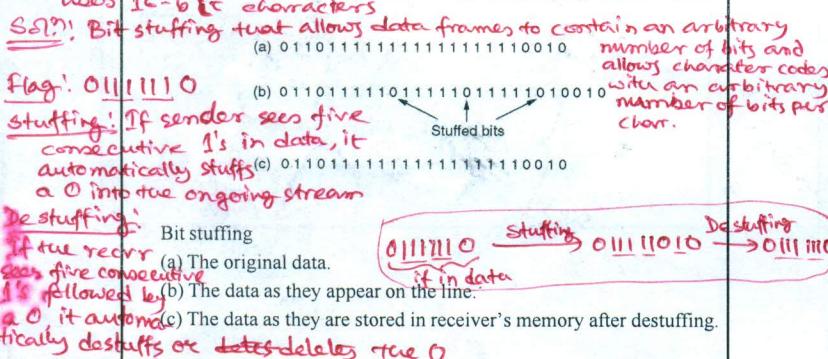


\Rightarrow Possible soln: Checksum \rightarrow The receiver will only know that the frame is bad, but, no way to know where the next frame starts

\Rightarrow Another possible soln: Sending a frame back to the src asking for a retrans \rightarrow still no way to know how many bits to skip

Problem of byte stuffing \Rightarrow It uses 8-bit characters.

However, not all character codes are 8-bit. For example, UNICODE uses 16-bit characters



\Rightarrow Error correction: n redundant or check bits $\xrightarrow{\text{for}} m+r$ bits data. Total n bits.

\Rightarrow Hamming distance: # of bit positions at which two code words differ.

\rightarrow To detect d errors, Hamming distance needs to be $(d+1)$

\rightarrow To correct d errors, Hamming distance needs to be $(2d+1)$

\Rightarrow Calculation of lower bound on r for correcting all single errors

each of 2^m legal messages has r illegal code words (can't be inverted each other)
distancce 1 from it.
So, each legal message (total 2^m) has $(n+r)$ bit patterns dedicated to it.

So, $(n+r)^2 \leq 2^m$

$\therefore (m+r)^2 \leq 2^m$

$\therefore (m+r) \leq 2^{\frac{m}{2}}$
gives the lower limit on r .

Error-Correcting Codes

Char.	ASCII (7bit)	Check bits	at power of 2, i.e., at 1, 2, 4, 8
H	1001000	00110010000	data bits at 3, 5, 6, 7, 9, 10, 11
a	1100001	1011001001	checkbit parity of data bits
m	1101101	1110101010	1 \rightarrow 2+1
m	1101011	1110101010	2 \rightarrow 4+1
i	1101001	0110101001	3 \rightarrow 4+2
n	1101110	0110101110	4 \rightarrow 4+2+1
g	0100000	1001100000	5 \rightarrow 8+1
c	1100011	1111000100	6 \rightarrow 8+2
o	1101111	1010101111	7 \rightarrow 8+2+1
d	1100100	1111001100	8 \rightarrow 8+2+1
e	1100101	0011001010	

Use of a Hamming code to correct burst errors.

* Check bits \Rightarrow Give even parity! # of 1's needs to (including itself) be even!

\Rightarrow When a codeword arrives at a receiver, it initializes a counter to zero. It then adds the posn k of an errored check bit to the counter.

If counter = 0, then no error.

If counter $\neq 0$, then the counter gives posn of the single error.

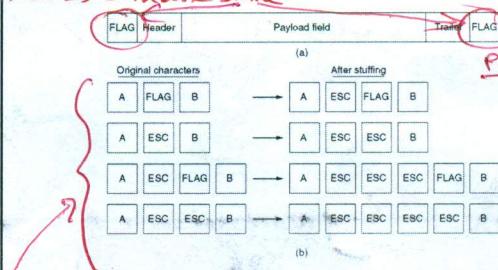
\Rightarrow To handle burst error of at most k length, k consecutive codewords are arranged as a matrix, one codeword per row. Data is transmitted (one column at a time). So, in case of k -bit error (in burst), each column will be affected by only one bit. At the receiver, the matrix is reconstructed, one column at a time. Needs $(k+1)$ check bits to make blocks of $k+m$ data bits immune to a single burst error of length k or less.

If encoding in physical medium contains some redundancy. For example, in some network 1 data bit is encoded by 2 physical bits. Here, high-low transition refers to 1 and low-high refers to 0. High-high and low-low are unlinked \Rightarrow can be a delimiter.

Byte stuffing or character stuffing \Rightarrow omits the synchronization problem

In prot \Rightarrow Starting and ending bytes were different

Now \Rightarrow Both are same



Problem: what happens if a flag appears in original data (may happen in binary files)?

Soln: Insert a special escape character just before the flag when it appears in data.

In case of destuffing, the extra added escape characters are removed.

New problem: what happens if the "escape" character itself appears in the data???

Soln: Again, it is stuffed and destuffed in the same way with additional "escape" character.

Different situations can occur here.

Error Detection and Correction

- Error-Correcting Codes
- Error-Detecting Codes

Telephone system \Rightarrow i) switches, ii) interoffice trunks, iii) local loops

For noisy medium (with noise) \rightarrow for noisy medium (less noise) \rightarrow for noisy medium (copper fiber).

\Rightarrow Errors are more common in analog than digital. We may experience this long, as replacing local loops are highly costly. Also, wireless case is far more error prone. So, we need to know how to deal with them.

\Rightarrow Errors may come in bursts or singly.

\Rightarrow Burst: Gives errors in less # of packets

Ex: if error rate is 0.01% and block size is 1000 bits, then most of the blocks will be errored for singly error. For burst errors of 100 bits, then 1 or 2 out of 100 packets will be in error.

\Rightarrow Problem of burst error: Difficult to correct

Generator polynomial, $G(x)$: Agreed upon by both sender & receiver

Both the high- and low-order bits must be 1

Should divide the polynomial represented by the checksummed frame

Error-Detecting Codes

Basic idea: Append a checksum to the end of the frame in such a way that the polynomial represented by the checksummed frame is divisible by $G(x)$

When a receiver gets the checksummed frame, it divides it by $G(x)$ (receiver)

The remainder whether it is 0 or not

1. Let r be the degree of $G(x)$. Append r zero bits to the low-order end of the frame, so it contains $m+r$ bits and corresponds to the polynomial $x^r M(x)$.

2. Divide the bit string corresponding to $x^r M(x)$ into the bit string corresponding to $x^r M(x)$.

Calculation of the polynomial code checksum using modulo 2 division (Add/Sub \rightarrow BBSXOR)

3. Subtract the remainder from the bit string corresponding to $x^r M(x)$ using modulo 2 subtraction. Result $T(x)$ is the checksummed frame to be transmitted.

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Sent to receiver

$T(x) + E(x) \rightarrow$ Received frame

* Each bit in $E(x)$ corresponds to a bit that has been inverted

Receiver

$T(x) + E(x) \rightarrow G(x)$ to error

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Transmitter frame: 110101101110

Received frame: 110101101110

Generator polynomial, $G(x)$: Received frame

$T(x) = 0$; $E(x)$ reflects $G(x)$ to error

Trans

→ proof of no polynomial with an odd # of terms has $(x+1)$ as its factor in modulo 2 system.
 → let the polynomial can have it. So, $E(x) = (x+1) \& G(x)$; Now, if $x=1$, then $E(1) = (1+1)G(1) = 0 \cdot G(1) = 0$
 On the other hand, if $E(x)$ contains an odd # of terms, putting $x=1$ will give $E(1) = 1 \neq contradiction!$

Con't: A polynomial code with r check bits detects all burst errors of length $\leq r$.
 → A burst error of length k can be represented as $x^i(x^{k-1} + \dots + 1)$, where i determines how far from the right-hand end of the received frame the burst is located. Now, if $G(x)$ contains an x^i term, then it will not have x^i as a factor.
 Besides, if the degree of the polynomial in parenthesized expression is less than the degree of $G(x)$, then the remainder X can never be zero.
 So, if the degree of the parenthesized expression is less than the degree of $G(x)$, i.e., r , then the remainder can never be zero.
 So, it will never divide the error by $G(x)$ with degree $\leq r$.

Elementary Data Link Protocols

- An Unrestricted Simplex Protocol
- A Simplex Stop-and-Wait Protocol
- A Simplex Protocol for a Noisy Channel

Protocol Definitions

```
#define MAX_PKT 1024
/* determines packet size in bytes */

typedef enum {false, true} boolean;
typedef unsigned int seq_nr;
typedef struct {unsigned char data[MAX_PKT];} packet; /* packet definition */
typedef enum {data, ack, nak} frame_kind;
/* frame_kind definition */

typedef struct {
  frame_kind kind;
  seq_nr seq;
  seq_nr ack;
  packet info;
} frame;
```

control fields
frame header

Continued →

Some definitions needed in the protocols to follow.
 These are located in the file protocol.h.

Protocol Definitions (ctd.)

Some definitions needed in the protocols to follow.
 These are located in the file protocol.h.

```
/* Wait for an event to happen; return its type in event. */
void wait_for_event(event_type *event);

/* Fetch a packet from the network layer for transmission on the channel. */
void void from_network_layer(packet *p);

/* Deliver information from an inbound frame to the network layer. */
void void to_network_layer(packet *p);

/* Go get an inbound frame from the physical layer and copy it to r. */
void void from_physical_layer(frame *r);

/* Pass the frame to the physical layer for transmission. */
void void to_physical_layer(frame *s);

/* Start the clock running and enable the timeout event. */
void start_timer(seq_nr k);

/* Stop the clock and disable the timeout event. */
void stop_timer(seq_nr k);

/* Start an auxiliary timer and enable the ack_timeout event. */
void start_ack_timer(void);

/* Stop the auxiliary timer and disable the ack_timeout event. */
void stop_ack_timer(void);

/* Allow the network layer to cause a network_layer_ready event. */
void enable_network_layer(void);

/* Forbid the network layer from causing a network_layer_ready event. */
void disable_network_layer(void);

/* Macro inc is expanded in-line. Increment k circularly. */
#define inc(k) if (k < MAX_SEQ) k = k + 1, else k = 0
```

Unrestricted Simplex Protocol

NL → PPL
PL → NL

```
/* Protocol 1 (utopia) provides for data transmission in one direction only, from sender to receiver. The communication channel is assumed to be error free, and the receiver is assumed to be able to process all the input infinitely quickly. Consequently, the sender just sits in a loop pumping data out onto the line as fast as it can. */

typedef enum {frame_arrival} event_type;
#include "protocol.h"

void sender1(void)
{
  frame s;
  packet buffer;
  while (true) {
    from_network_layer(&buffer); /* go get something to send */
    s.info = buffer; /* copy it into s for transmission */
    to_physical_layer(&s); /* bye bye little frame */
  }
}

void receiver1(void)
{
  frame r;
  event_type event;
  while (true) {
    wait_for_event(&event);
    from_physical_layer(&r);
    to_network_layer(&r.info); /* only possibility is frame_arrival */
  }
}
```

No flow control,
No error control

Simplex Stop-and-Wait Protocol

```
/* Protocol 2 (stop-and-wait) also provides for a one-directional flow of data from sender to receiver. The communication channel is once again assumed to be error free, as in protocol 1. However, this time, the receiver has only a finite buffer capacity and a finite processing speed, so the protocol must explicitly prevent the sender from flooding the receiver with data faster than it can be handled. */

typedef enum {frame_arrival} event_type;
#include "protocol.h"

void sender2(void)
{
  frame s;
  packet buffer;
  event_type event;
  while (true) {
    from_network_layer(&buffer); /* go get something to send */
    s.info = buffer; /* copy it into s for transmission */
    to_physical_layer(&s); /* bye bye little frame */
    wait_for_event(&event); /* do not proceed until given the go ahead */
  }
}

void receiver2(void)
{
  frame r;
  event_type event;
  while (true) {
    wait_for_event(&event);
    from_physical_layer(&r);
    to_network_layer(&r.info);
    to_physical_layer(&r);
  }
}
```

flow control

A Simplex Protocol for a Noisy Channel

```
/* Protocol 3 (par) allows unidirectional data flow over an unreliable channel. */
#define MAX_SEQ 1 /* must be 1 for protocol 3 */
typedef enum {frame_arrival, cksum_err, timeout} event_type;
#include "protocol.h"

void sender3(void)
{
  seq_nr next_frame_to_send;
  frame s;
  packet buffer;
  event_type event;
  next_frame_to_send = 0;
  from_network_layer(&buffer);
  while (true) {
    if (next_frame_to_send == 0) {
      s.seq = next_frame_to_send;
      to_physical_layer(&s);
      start_timer(s.seq);
    }
    wait_for_event(&event);
    if (event == frame_arrival) {
      from_physical_layer(&s);
      if (s.seq == next_frame_to_send) {
        stop_timer(s.seq);
        from_network_layer(&buffer);
        inc(next_frame_to_send);
      }
    }
  }
}
```

A positive acknowledgement with retransmission protocol.

Continued →

PAR (Positive Acknowledgement with Retransmission)
 or ARQ (Automatic Repeat Request)

A Simplex Protocol for a Noisy Channel (ctd.)

```

void receiver3(void)
{
    seq_nr frame_expected;
    frame r;
    event_type event;
    frame_expected = 0;
    while (true) {
        wait_for_event(&event);
        if (event == frame_arrival) {
            from_physical_layer(&r);
            if (r.seq == frame_expected) {
                to_network_layer(&r.info);
                inc(frame_expected);
            }
            r.sack = 1 - frame_expected;
            to_physical_layer(&r);
        }
    }
}

```

MAX_SEQ is set to 1, it always gives the A positive acknowledgement with retransmission protocol. Last acked seq # (either 0 or 1)

Sender's sending window \Rightarrow Corresponds to seq#s of frames it is permitted to send

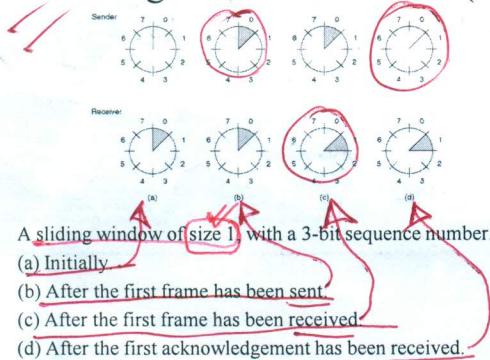
- Lower edge advances when an ACK comes in

- Upper edge advances when a new packet arrives from NL

Receiving window \Rightarrow Corresponds to a set of frames it is permitted to receive

- Rotates by one when a frame whose seq# is equal to the lower edge of the window is received

Sliding Window Protocols (2)



Sliding Window Protocols

- A One-Bit Sliding Window Protocol
- A Protocol Using Go Back N
- A Protocol Using Selective Repeat

A One-Bit Sliding Window Protocol (ctd.)

```

while (true) {
    wait_for_event(&event);
    if (event == frame_arrival) {
        from_physical_layer(&r);
        if (r.seq == frame_expected) {
            to_network_layer(&r.info);
            inc(frame_expected);
        }
        if (r.ack == next_frame_to_send) {
            stop_timer(r.ack);
            from_network_layer(&buffer);
            inc(next_frame_to_send);
        }
    }
    s.info = buffer;
    s.seq = next_frame_to_send;
    s.ack = 1 - frame_expected;
    to_physical_layer(&s);
    start_timer(s.seq);
}

```

has been received

A One-Bit Sliding Window Protocol

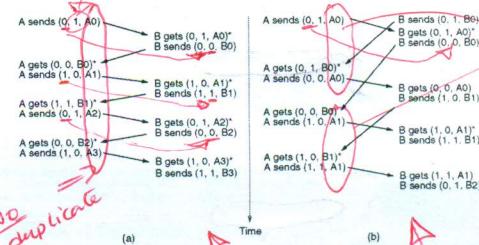
```

/* Protocol 4 (sliding window) is bidirectional. */
#define MAX_SEQ 1 /* must be 1 for protocol 4 */
typedef enum {frame_arrival, cksum_err, timeout} event_type;
#include "protocol.h"
void protocol4(void)
{
    seq_nr next_frame_to_send;
    seq_nr frame_expected;
    frame r, s;
    packet buffer;
    event_type event;
    next_frame_to_send = 0;
    frame_expected = 0;
    from_network_layer(&buffer);
    s.info = buffer;
    s.seq = next_frame_to_send;
    s.ack = 1 - frame_expected;
    to_physical_layer(&s);
    start_timer(s.seq);
}

```

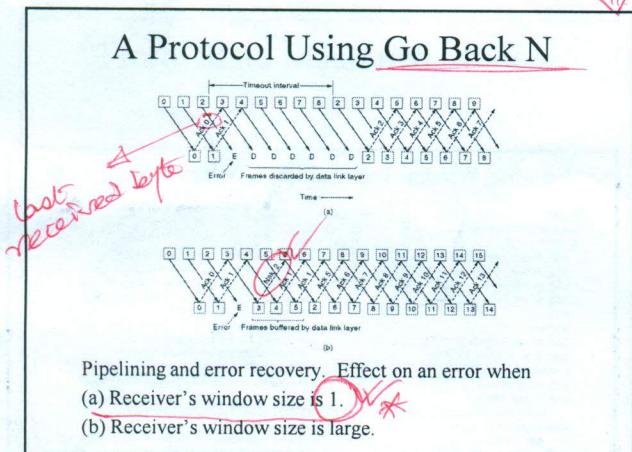
Continued \Rightarrow

A One-Bit Sliding Window Protocol (2)

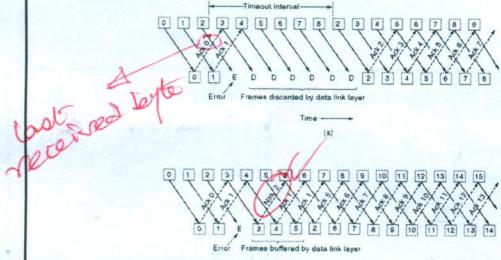


Two scenarios for protocol 4. (a) Normal case, (b) Abnormal case. The notation is (seq, ack, packet number). An asterisk indicates where a network layer accepts a packet.

out of the packets get duplicated



A Protocol Using Go Back N



Pipelining and error recovery. Effect on an error when
~~(a) Receiver's window size is 1.~~
(b) Receiver's window size is large.

Sliding Window Protocol Using Go Back N

```

/* Protocol 5 (pipelining) allows multiple outstanding frames. The sender may transmit up
to MAX_SEQ frames without waiting for an ack. In addition, unlike the previous protocols
the network layer is not assumed to have a new packet all the time. Instead, the
network layer causes a network_layer_ready event when there is a packet to send. */

#define MAX_SEQ 7           /* should be 2^n - 1 */
typedef enum {frame_arrival, cksum_err, timeout, network_layer_ready} event_type;
#include "protocol.h"

static boolean between(seq_nr_a, seq_nr_b, seq_nr_c)
{
    /* Return true if a <= b < c circularly; false otherwise. */
    if (((a <= b) && (b < c)) || ((c < a) && (a <= b)) || (b < c) && (c < a)))
        return(true);
    else
        return(false);
}

static void send_data(seq_nr_frame_nr, seq_nr_frame_expected, packet buffer[])
{
    /* Construct and send a data frame. */
    frame s;                      /* scratch variable +*/
    s.info = buffer[frame_nr];     /* insert packet into frame */
    s.seq = frame_nr;              /* insert sequence number into frame */
    s.ack = (frame_expected + MAX_SEQ) % (MAX_SEQ + 1); /* piggyback ack */
    to_physical_layer(s);          /* transmit the frame */
    start_timer(frame_nr);         /* start the timer running */
}

```

Continued →

Sliding Window Protocol Using Go Back N

```

void protocol5(void)
{
    seq_nr next_frame_to_send;
    seq_nr ack_expected;
    seq_nr frame_expected;
    frame_r;
    packet buffer[MAX_SEQ + 1];
    seq_nr nbuffered;
    seq_nr i;
    event_type event;

    enable_network_layer();
    ack_expected = 0;
    next_frame_to_send = 0;
    frame_expected = 0;
    nbuffered = 0;

    /* MAX_SEQ > 1; used for outbound stream */
    /* oldest frame as yet unacknowledged */
    /* next frame expected on inbound stream */
    /* scratch variable */
    /* buffers for the outbound stream */
    /* # output buffers currently in use */
    /* used to index into the buffer array */

    /* allow network_layer_ready events */
    /* next ack expected inbound */
    /* next frame going out */
    /* number of frame expected inbound */
    /* initially no packets are buffered */
}

```

Continued →

Sliding Window Protocol Using Go Back N

```

while (true) {
    wait_for_event(&event); /* four possibilities: see event_type above */

    switch(event) {
        case network_layer_ready: /* the network layer has a packet to send */
            /* Accept, save, and transmit a new frame. */
            from_network_layer(&buffer[next_frame_to_send]); /* fetch new packet */
            nbuffered = nbuffered + 1; /* expand the sender's window */
            send_data(next_frame_to_send, frame_expected, buffer); /* transmit the frame */
            inc(next_frame_to_send); /* advance sender's upper window edge */
            break;

        case frame_arrival: /* a data or control frame has arrived */
            from_physical_layer(&r); /* get incoming frame from physical layer */

            if (r.seq == frame_expected) {
                /* Frames are accepted only in order. */
                to_network_layer(&r.info); /* pass packet to network layer */
                inc(frame_expected); /* advance lower edge of receiver's window */
            }
    }
}

```

Continued →

Sliding Window Protocol Using Go Back N

```

/* Ack N implies n - 1, n - 2, etc. Check for this. */
while((back(ack_expected, r_ack, next_frame_to_send)) {
    /* Handle piggybacked ack. */
    nbuffered = nbuffered - 1; /* one frame fewer buffered */
    stop_timer(ack_expected); /* frame arrived intact; stop timer */
    inc(ack_expected); /* contract sender's window */
}

break;
}

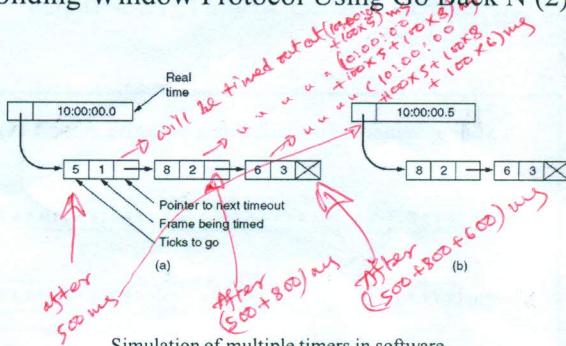
case cksm_err: break; /* just ignore bad frames */

case timeout: /* trouble; retransmit all outstanding frames */
    next_frame_to_send = ack_expected; /* start retransmitting here */
    for (i = 1; i <= nbuffered; i++) {
        send_data(next_frame_to_send, frame_expected, buffer); /* resend 1 frame */
        inc(next_frame_to_send); /* prepare to send the next one */
    }
}

if (nbuffered < MAX_SEQ)
    enable_network_layer();
else
    disable_network_layer();
}

```

Sliding Window Protocol Using Go Back N (2)



⇒ Because Protocol S has multiple outstanding frames, it logically needs multiple timers, one per outstanding frame. However, all the timers can be simulated in software using a single h/w clock that causes interrupt periodically. Pending timeouts form a linked list telling the number of clock ticks (every 500 ms in the example) until the timer expires, the frame being timed, and a pointer to the next node.

A Sliding Window Protocol Using Selective Repeat

```

/* Protocol 6 (nonsequential receive) accepts frames out of order, but passes packets to the
network layer in order. Associated with each outstanding frame is a timer. When the timer
expires, only that frame is retransmitted, not all the outstanding frames, as in protocol 5. */

#define MAX_SEQ 7
#define NR_BUFS (MAX_SEQ + 1)/2
typedef enum {frame_arrival, csum_err, timeout, network_layer_ready, ack_timeout} event_type;
#include "protocol.h"
boolean no_nak = true; /* no nak has been sent yet */
seq_nr oldest_frame = MAX_SEQ + 1;
static boolean between(seq_nr a, seq_nr b, seq_nr c)
{
    /* Same as between in protocol5, but shorter and more obscure. */
    return ((a <= b) && (b <= c)) || ((c < a) && (b <= c)) || ((b < c) && (c < a));
}
static void send_frame(frame_kind lk, seq_nr frame_nr, seq_nr frame_expected, packet buffer[])
{
    /* Construct and send a data, ack, or nak frame. */
    frame s;
    s.kind = lk; /* kind == data, ack, or nak */
    if (lk == data) s.info = buffer[frame_nr % NR_BUFS];
    s.seq = frame_nr; /* only meaningful for data frames */
    s.ack = (frame_expected + MAX_SEQ) % (MAX_SEQ + 1);
    if (lk == nak) no_nak = false; /* one nak per frame, please */
    to_physical_layer(s);
    if (lk == data) start_timer(frame_nr % NR_BUFS);
    stop_ack_timer(); /* no need for separate ack frame */
}

void protocol6(void)
{
    /* lower edge of sender's window */
    seq_nr ack_expected;
    seq_nr next_frame_to_send;
    seq_nr frame_expected;
    seq_nr too_far;
    int i;
    frame r;
    packet out_buf[NR_BUFS];
    packet in_buf[NR_BUFS];
    boolean arrived[NR_BUFS];
    seq_nr nbuffered; /* scratch variable */
    /* buffers for the outbound stream */
    /* buffers for the inbound stream */
    /* inbound bit map */
    /* how many output buffers currently used */
    event_type event;
    enable_network_layer();
    ack_expected = 0; /* initialize */
    next_frame_to_send = 0; /* next ack expected on the inbound stream */
    frame_expected = 0; /* number of next outgoing frame */
    too_far = NR_BUFS;
    nbuffered = 0; /* initially no packets are buffered */
    for (i = 0; i < NR_BUFS; i++) arrived[i] = false;
    event_type event;
}

```

Continued →

A Sliding Window Protocol Using Selective Repeat (3)

```

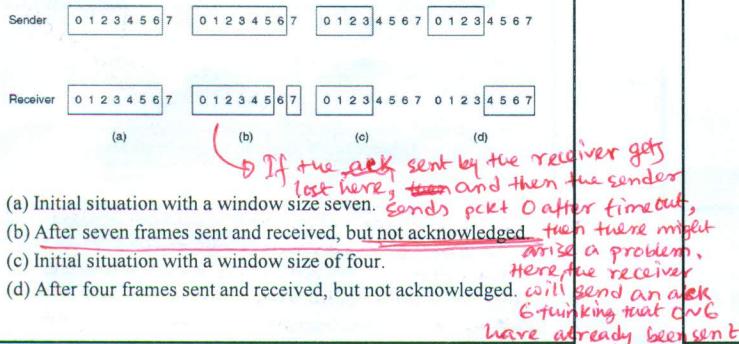
while (true) {
    wait_for_event(&event); /* five possibilities: see event_type above */
    switch(event) {
        case network_layer_ready:
            if (nbuffered > 1): /* accept, save, and transmit a new frame */
                from_network_layer_out.buf(next_frame_to_send % NR_BUFS); /* fetch new packet */
                send_frame(data, next_frame_to_send, frame_expected, out_buf); /* transmit the frame */
                inc(next_frame_to_send); /* advance upper window edge */
            break;

        case frame_arrival:
            /* a data or control frame has arrived */
            from_physical_layer(&r);
            if (r.kind == data) {
                /* An undamaged frame has arrived. */
                if ((r.seq == frame_expected) && no_nak)
                    send_frame(nak, 0, frame_expected, out_buf); /* start ack timer */;
                if (between(frame_expected, r.seq, too_far) && (arrived[r.seq % NR_BUFS] == false)) {
                    /* Frames may be accepted in any order. */
                    arrived[r.seq % NR_BUFS] = true; /* mark buffer as full */
                    in_buf[r.seq % NR_BUFS] = r.info; /* insert data into buffer */
                    while (arrived[frame_expected % NR_BUFS]) {
                        /* Pass frames and advance window. */
                        to_network_layer(&in_buf[frame_expected % NR_BUFS]);
                        no_nak = true;
                        arrived[frame_expected % NR_BUFS] = false;
                        inc(frame_expected); /* advance lower edge of receiver's window */
                        inc(too_far); /* advance upper edge of receiver's window */
                        start_ack_timer(); /* to see if a separate ack is needed */
                    }
                }
            }
    }
}

```

Continued →

A Sliding Window Protocol Using Selective Repeat (5)



(This happens as the newer acks get lost and the sender has no way to understand which pkt 0 (either new or old) has been acked). This problem arises as there is an overlap between earlier window and new window when an window advances. Its remedy is to make sure that there is no such overlap. To do so, the window size needs to be $\text{MAX_SEQ} + 1$.

A Sliding Window Protocol Using Selective Repeat (2)

```

void protocol6(void)
{
    /* lower edge of sender's window */
    seq_nr ack_expected;
    seq_nr next_frame_to_send;
    seq_nr frame_expected;
    seq_nr too_far;
    int i;
    frame r;
    packet out_buf[NR_BUFS];
    packet in_buf[NR_BUFS];
    boolean arrived[NR_BUFS];
    seq_nr nbuffered; /* scratch variable */
    /* buffers for the outbound stream */
    /* buffers for the inbound stream */
    /* inbound bit map */
    /* how many output buffers currently used */
    event_type event;
    enable_network_layer();
    ack_expected = 0; /* initialize */
    next_frame_to_send = 0; /* next ack expected on the inbound stream */
    frame_expected = 0; /* number of next outgoing frame */
    too_far = NR_BUFS;
    nbuffered = 0; /* initially no packets are buffered */
    for (i = 0; i < NR_BUFS; i++) arrived[i] = false;
    event_type event;
}

```

Continued →

A Sliding Window Protocol Using Selective Repeat (4)

```

if((r.kind==nak) && between(ack_expected,(r.ack+1)%(MAX_SEQ+1),next_frame_to_send))
    send_frame(data, (r.ack+1) % (MAX_SEQ + 1), frame_expected, out_buf);

while (between(ack_expected, r.ack, next_frame_to_send)) {
    nbuffered = nbuffered - 1; /* handle piggybacked ack */
    stop_timer(ack_expected % NR_BUFS); /* frame arrived intact */
    inc(ack_expected); /* advance lower edge of sender's window */
}

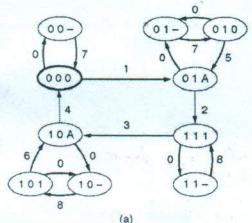
break;
case csum_err:
    if (no_nak) send_frame(nak, 0, frame_expected, out_buf); /* damaged frame */
    break;
case timeout:
    send_frame(data, oldest_frame, frame_expected, out_buf); /* we timed out */
    break;
case ack_timeout:
    send_frame(ack, 0, frame_expected, out_buf); /* ack timer expired; send ack */
}
if (nbuffered < NR_BUFS) enable_network_layer(); else disable_network_layer();
}
}

```

Protocol Verification

- Finite State Machined Models
- Petri Net Models

Finite State Machined Models

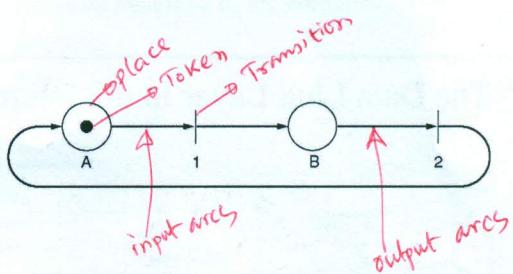


Transition	Who runs?	Frame accepted	Frame emitted	To network layer
0	-	-	-	(frame lost)
1	R	0	A	Yes
2	S	A	1	Yes
3	R	1	A	-
4	S	A	0	-
5	R	0	A	No
6	R	1	A	No
7	S	(timeout)	0	-
8	S	(timeout)	1	-

(b)

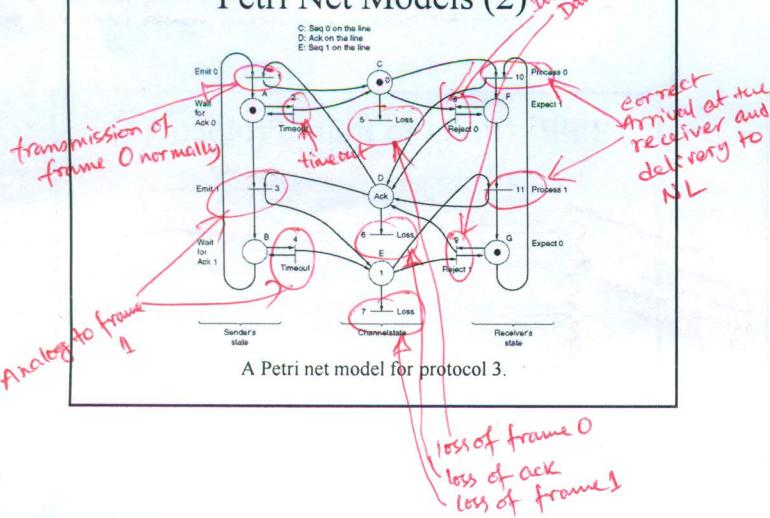
(a) State diagram for protocol 3. (b) Transmissions.

Petri Net Models



A Petri net with two places and two transitions.

Petri Net Models (2)

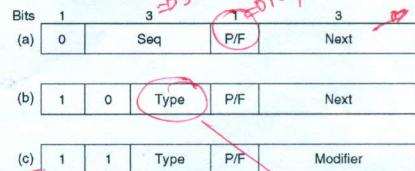


High-Level Data Link Control



Frame format for bit-oriented protocols.

High-Level Data Link Control (2)



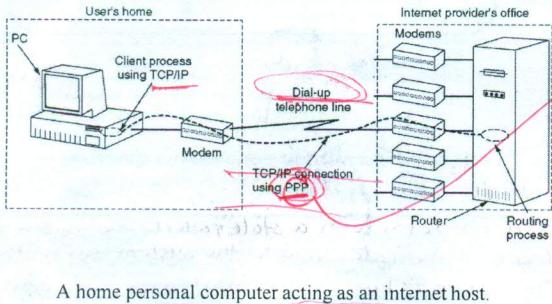
3 kinds of frames

- (a) An information frame.
- (b) A supervisory frame.
- (c) An unnumbered frame.

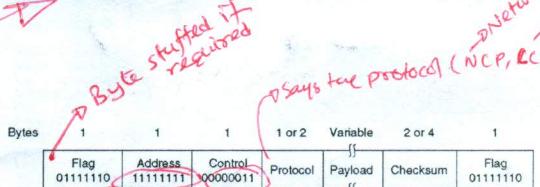
- 0 → Acknowledgment (Similar to Protocol 5)
- 1 → Reject (NAK)
- 2 → Receive NAK ready
- 3 → Selective Reject (Similar to Protocol 6)

D can be used for both carrying data for unreliable connection and for control purpose.

The Data Link Layer in the Internet



PPP – Point to Point Protocol

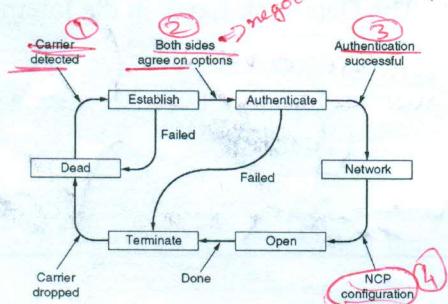


The PPP full frame format for unnumbered mode operation.

NCP → For negotiating network-layer options in a way that is independent of the network layer protocol to be used.

LCP → Bringing lines up, testing them, negotiating options, and bringing them down again when not needed.

PPP – Point to Point Protocol (2)



A simplified phase diagram for bring a line up and down.

PPP – Point to Point Protocol (3)

Name	Direction	Description
Configure-request	I → R	List of proposed options and values
Configure-ack	I ← R	All options are accepted
Configure-nak	I ← R	Some options are not accepted
Configure-reject	I ← R	Some options are not negotiable
Terminate-request	I → R	Request to shut the line down
Terminate-ack	I ← R	OK, line shut down
Code-reject	I ← R	Unknown request received
Protocol-reject	I ← R	Unknown protocol requested
Echo-request	I → R	Please send this frame back
Echo-reply	I ← R	Here is the frame back
Discard-request	I ← R	Just discard this frame (for testing)

The LCP frame types. (11 frames)